





Input combinations and prediction changes

Named modalities use observed features; others receive reference features. Model weights and masks stay fixed.

Train-mean reference Zero reference



All 20 cases included. Full input has 100% agreement and zero deviation by definition, not by accuracy.

Bottom panels replace one modality while retaining both others; these effects are conditional and are not additive.

Sample 13 has all-zero raw vision; samples 07 and 18 have truncated text. Both references mainly differ in text.