


Attachment 4 | Sample 01 | Time-step and feature evidence

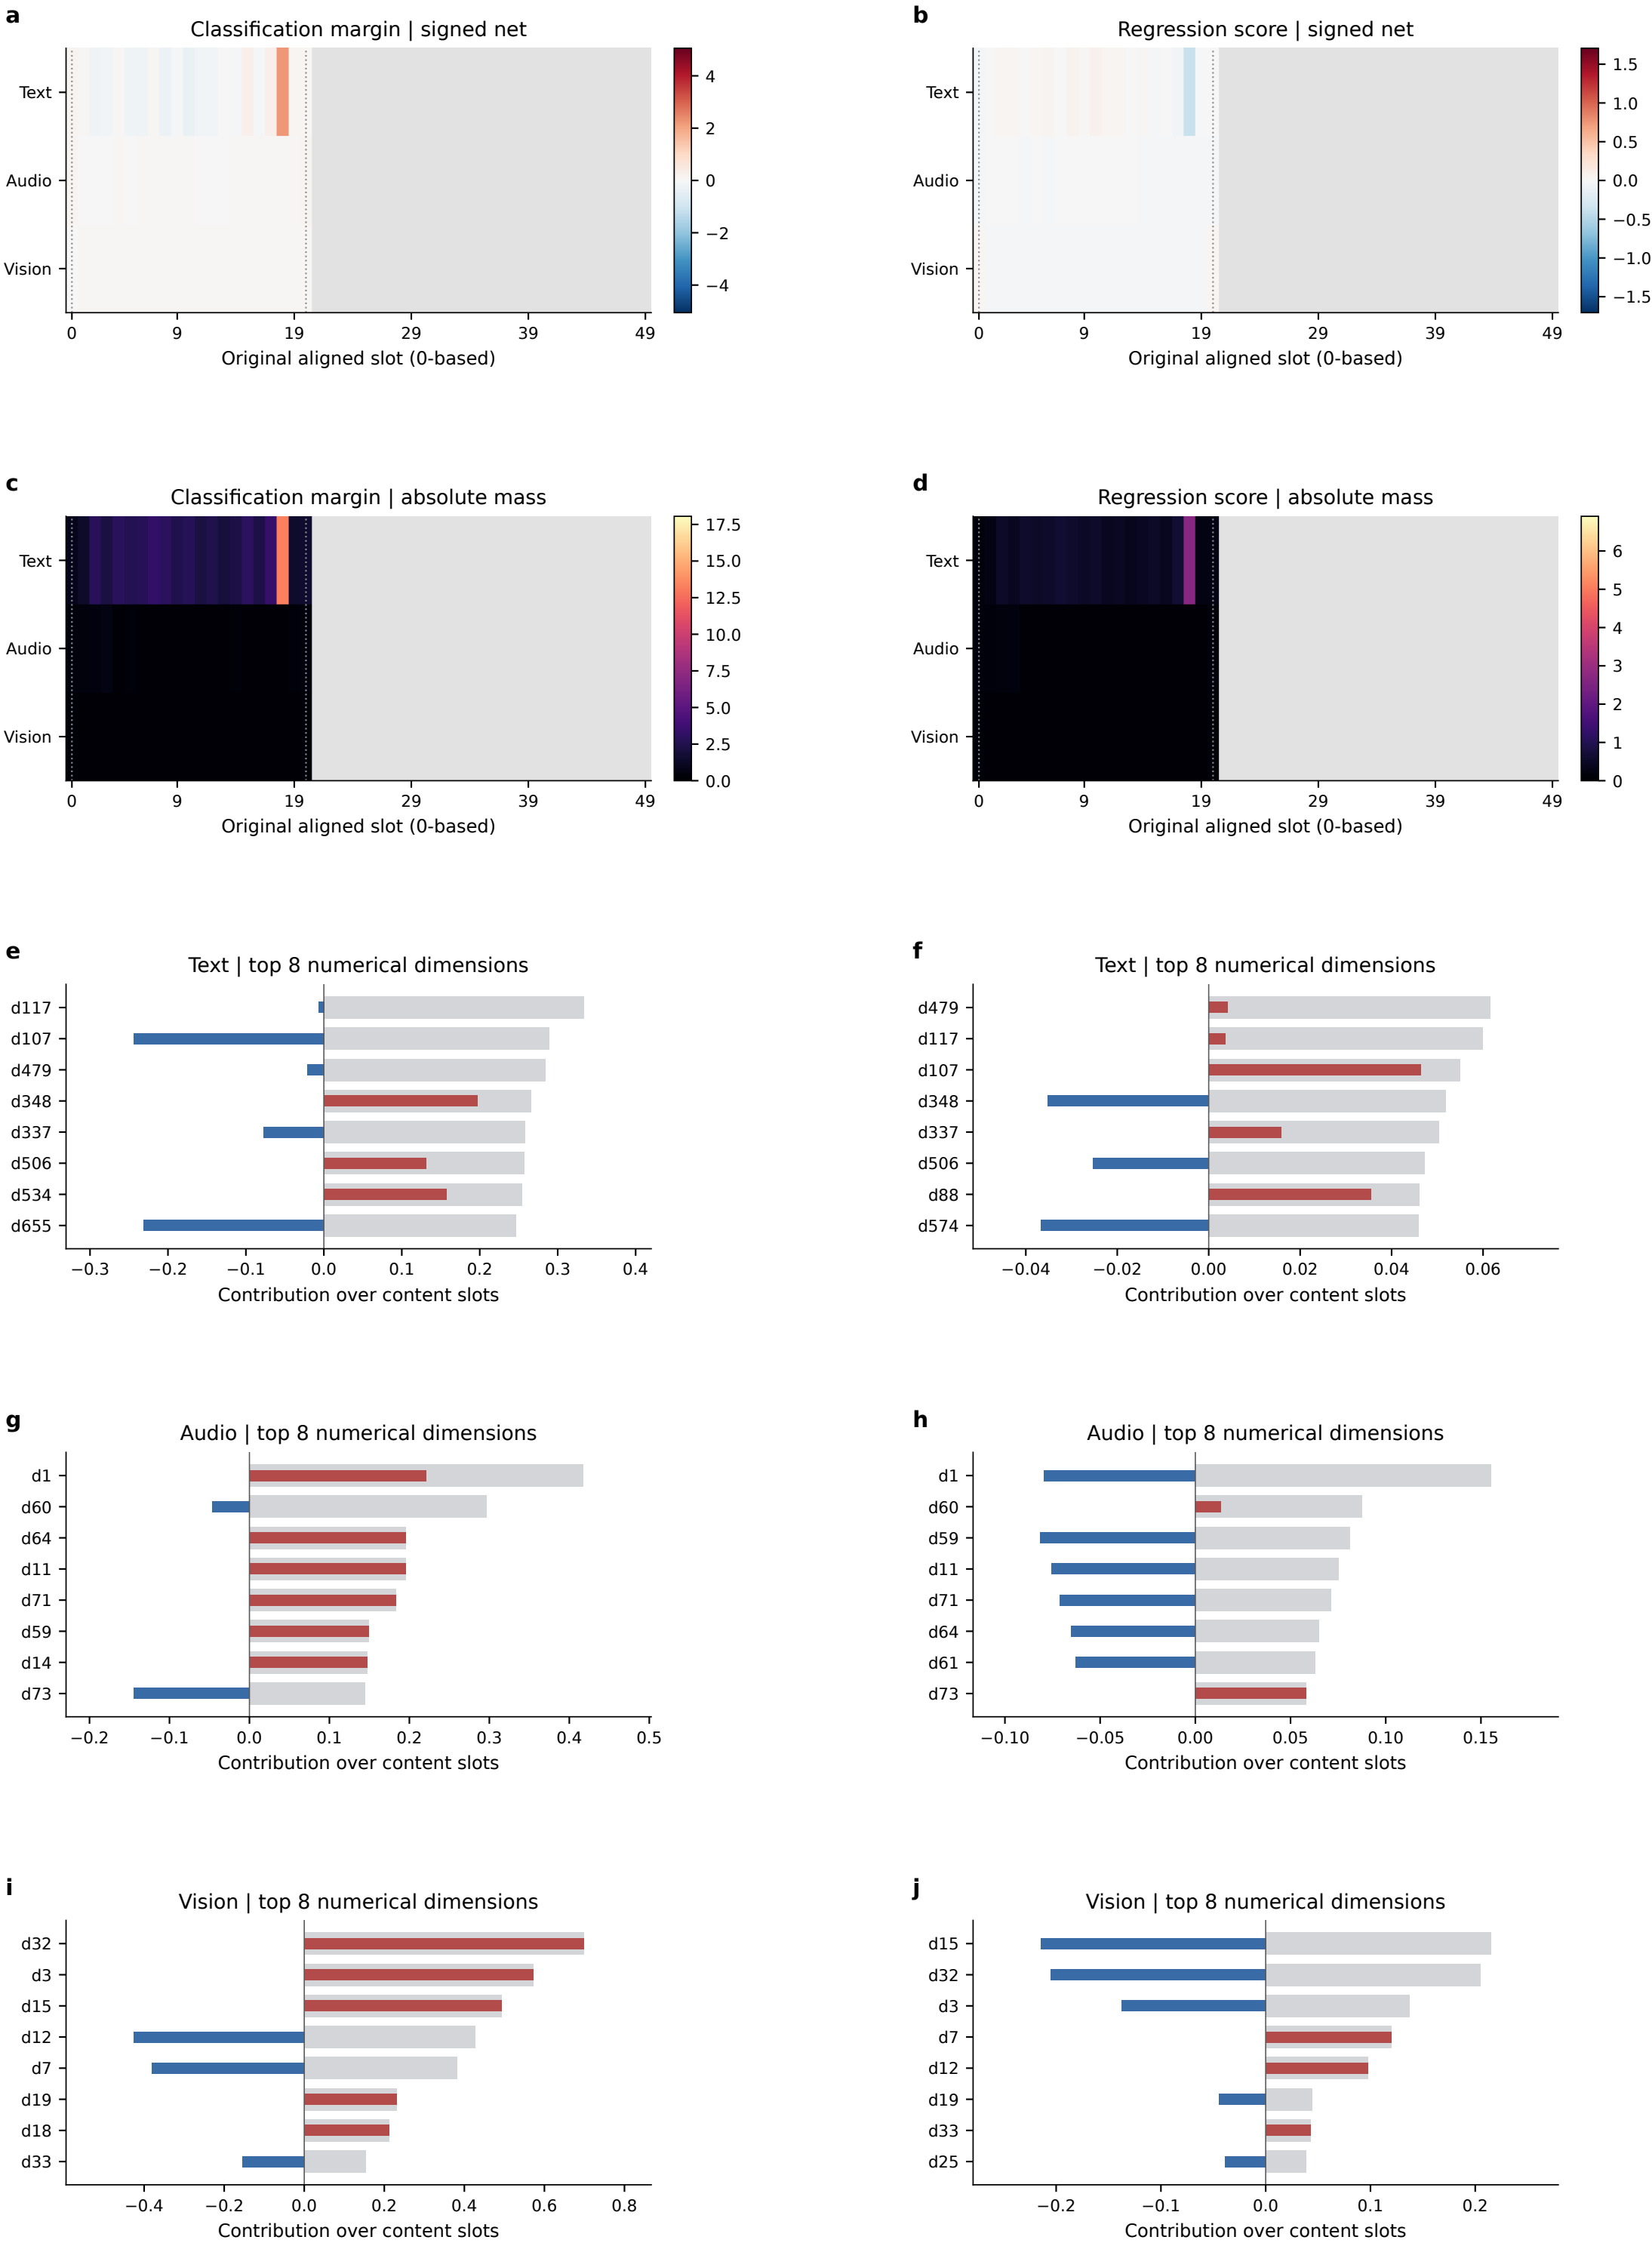
Predicted class 1 | Class probabilities [0: 0.073, 1: 0.490, 2: 0.436] | Regression 0.1541



Train-mean reference. Heatmap scales are shared across all 20 samples for each target. Gray cells: padding; dotted lines: CLS/SEP. Absolute mass is not a modality share. Dimension rankings exclude specials and padding. Dimensions are numerical indices, not semantic labels or video timestamps.

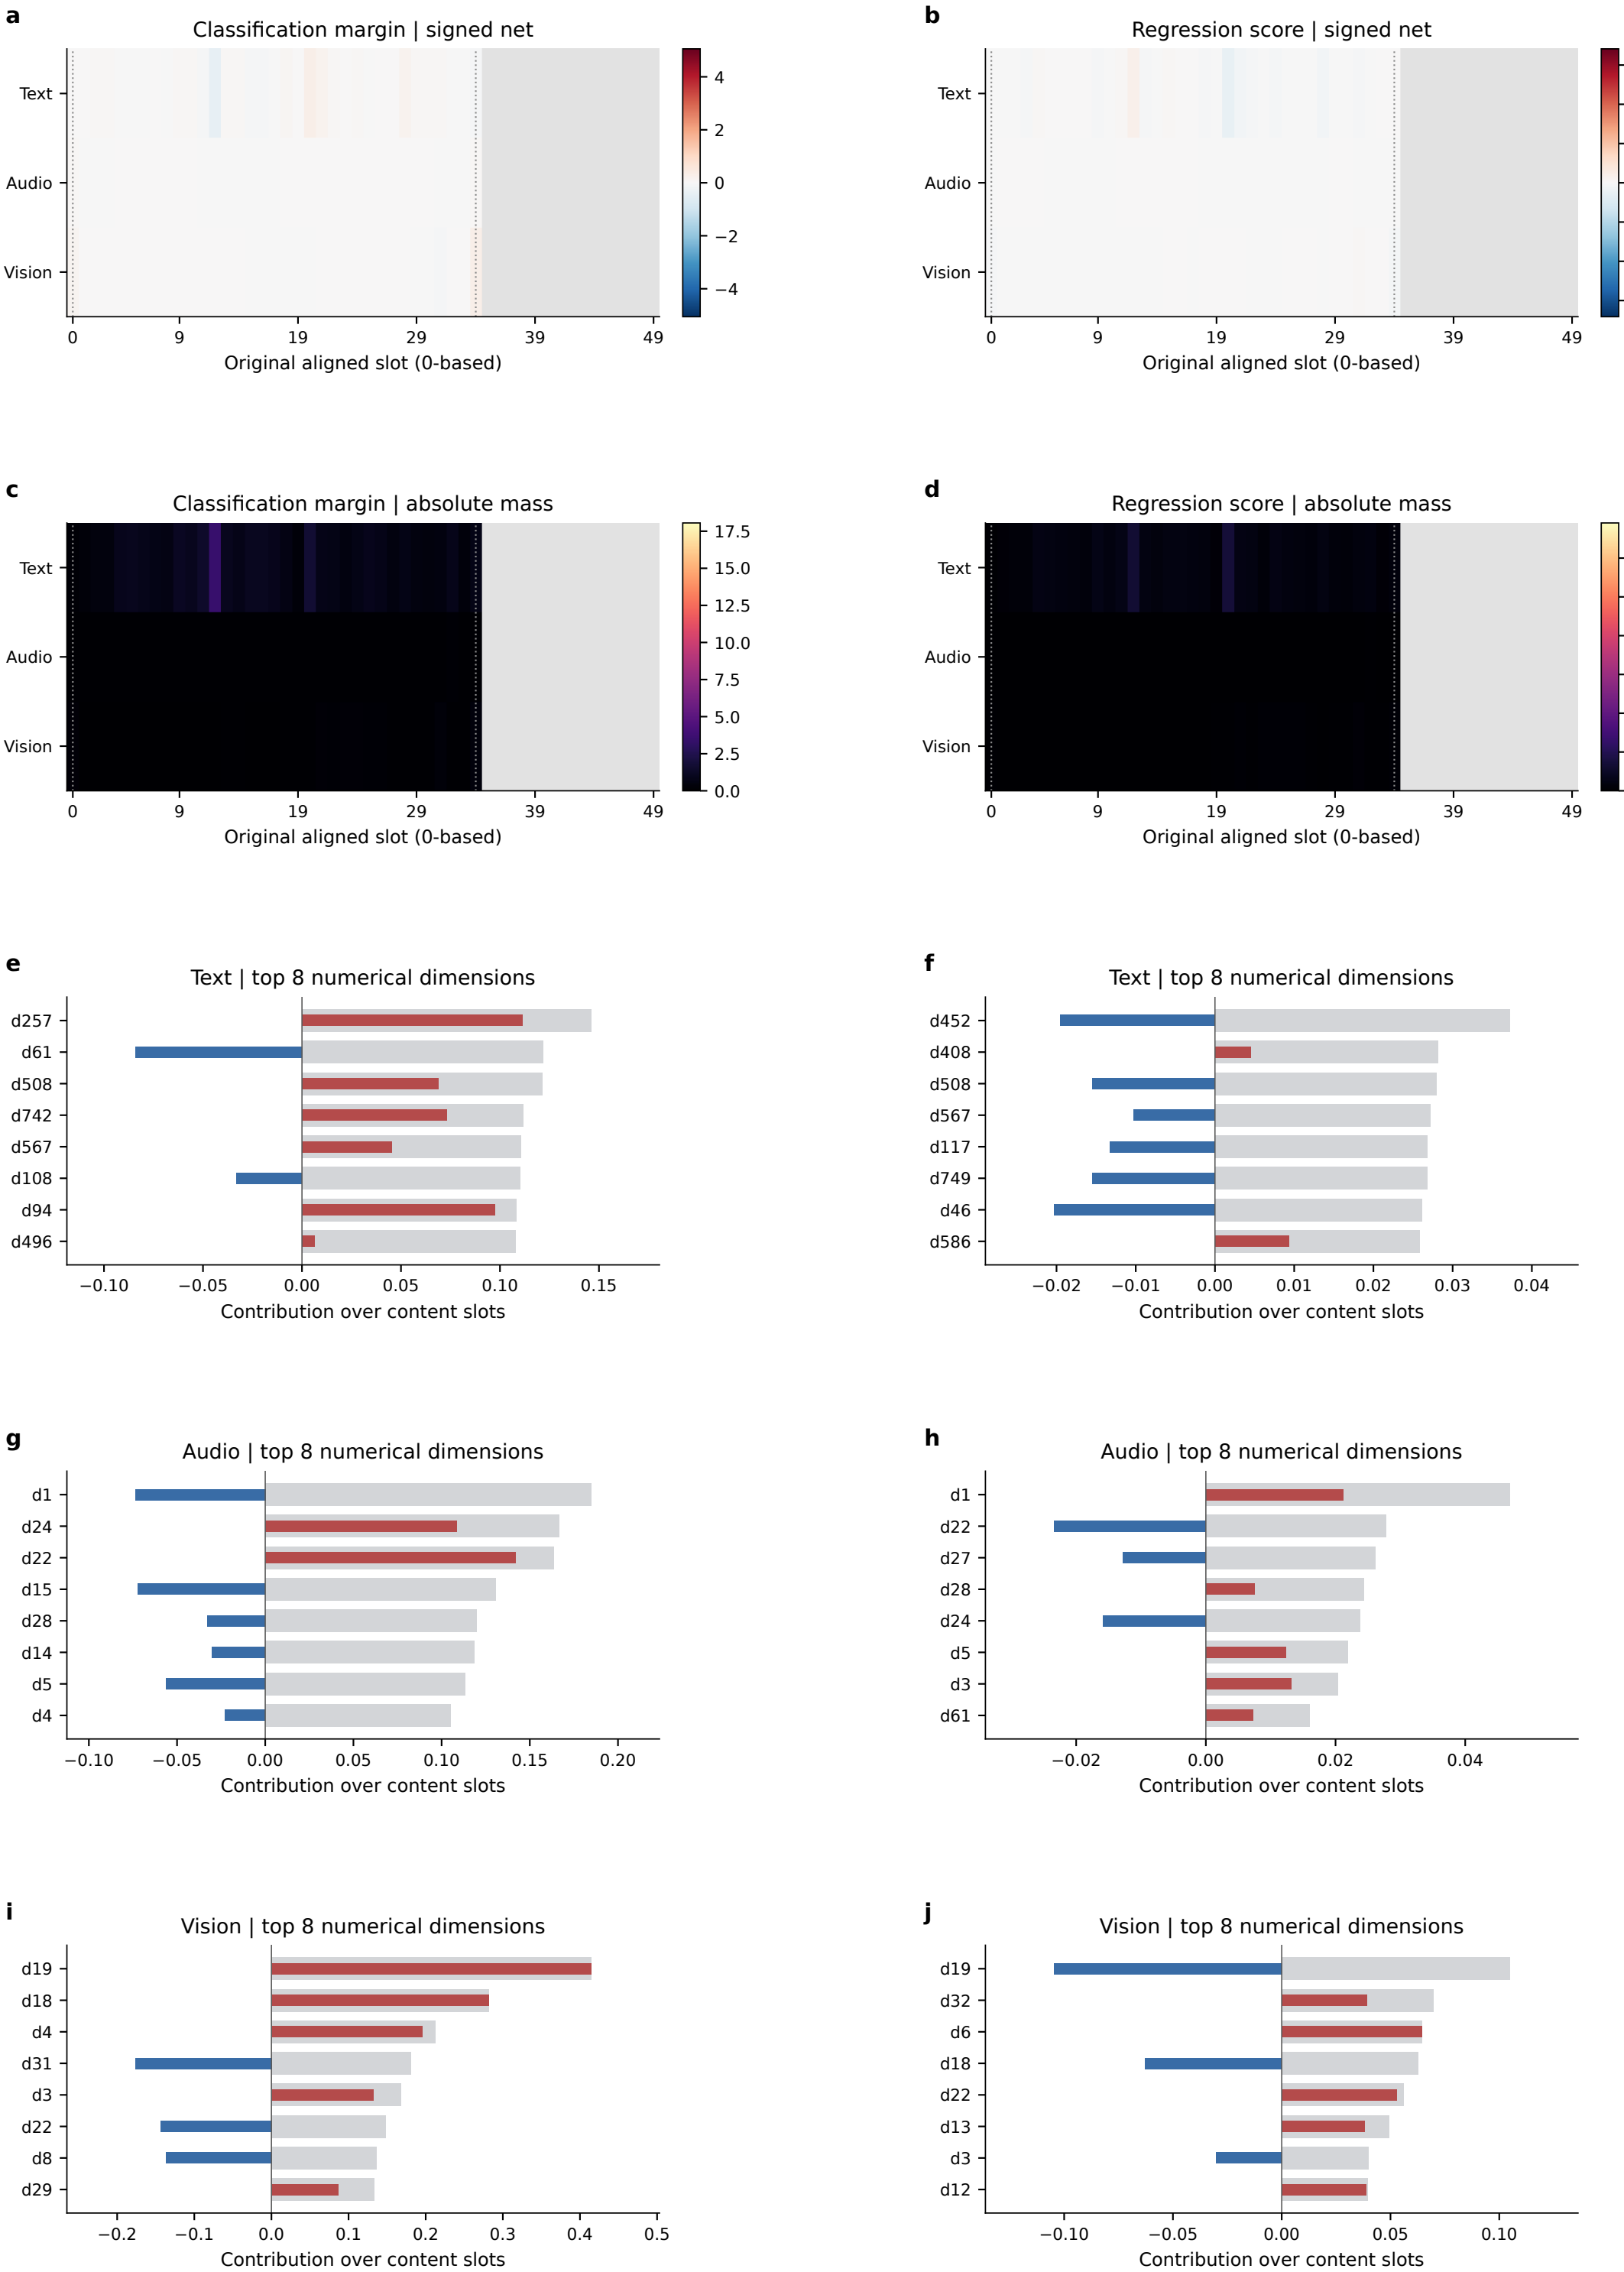
Attachment 4 | Sample 02 | Time-step and feature evidence

Predicted class 0 | Class probabilities [0: 0.571, 1: 0.190, 2: 0.239] | Regression -0.2530



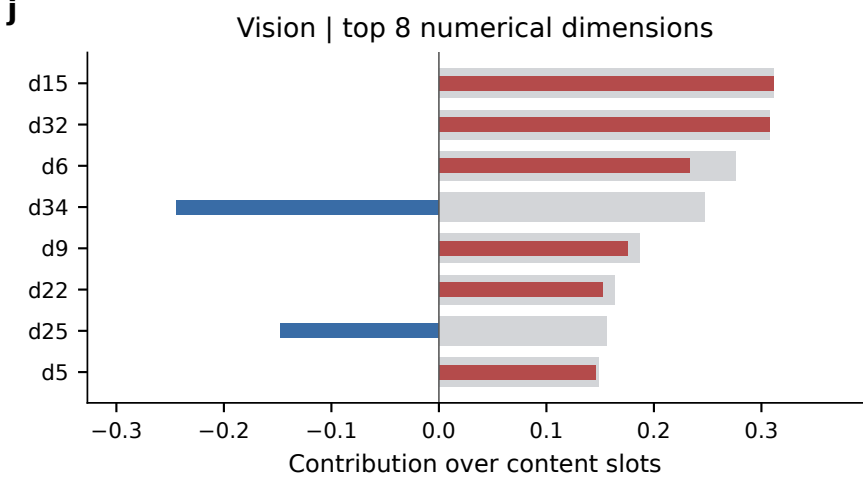
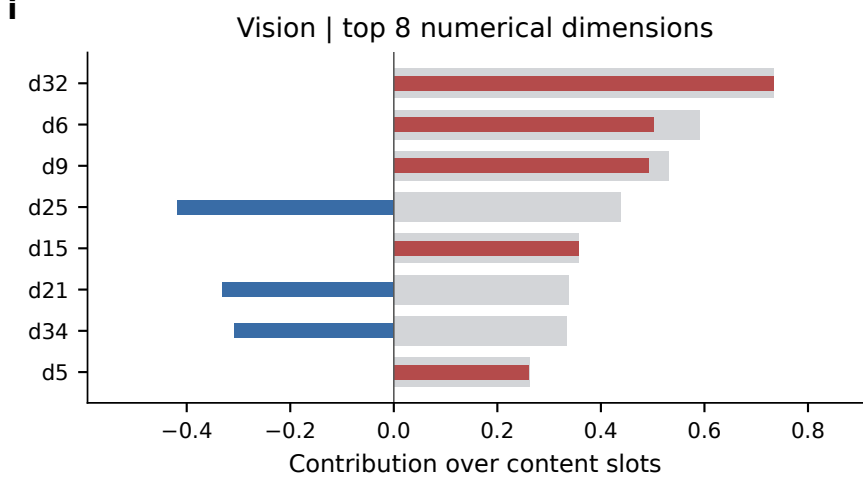
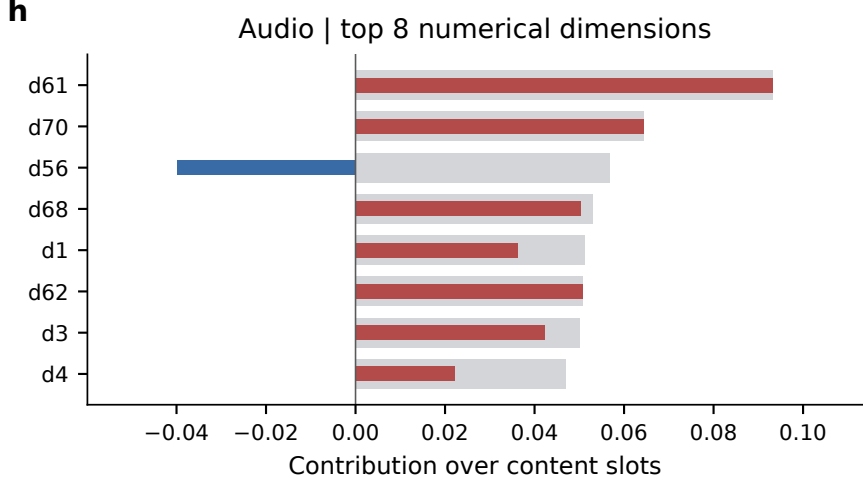
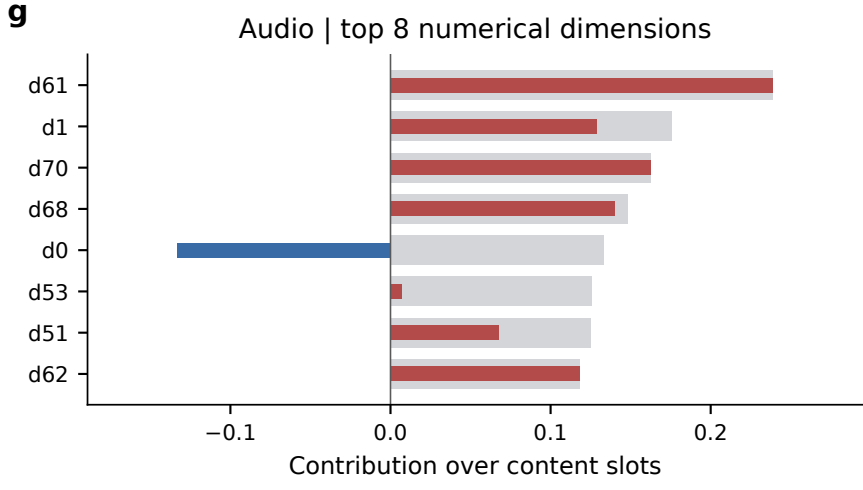
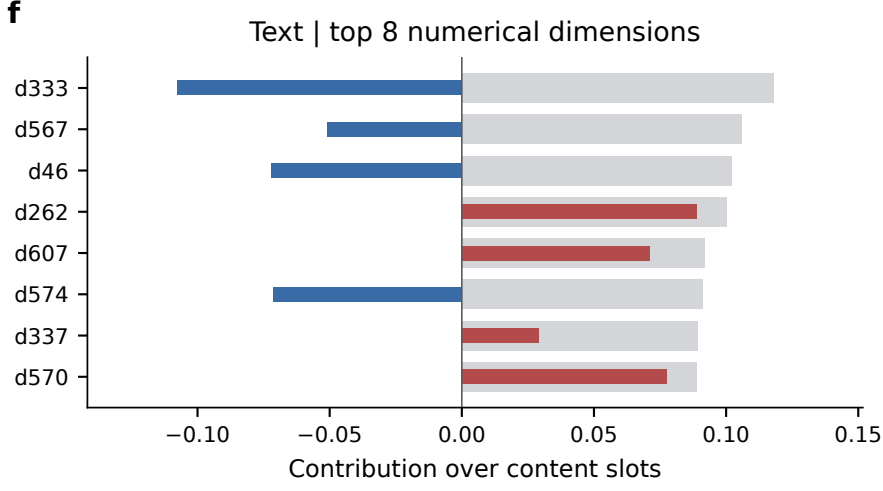
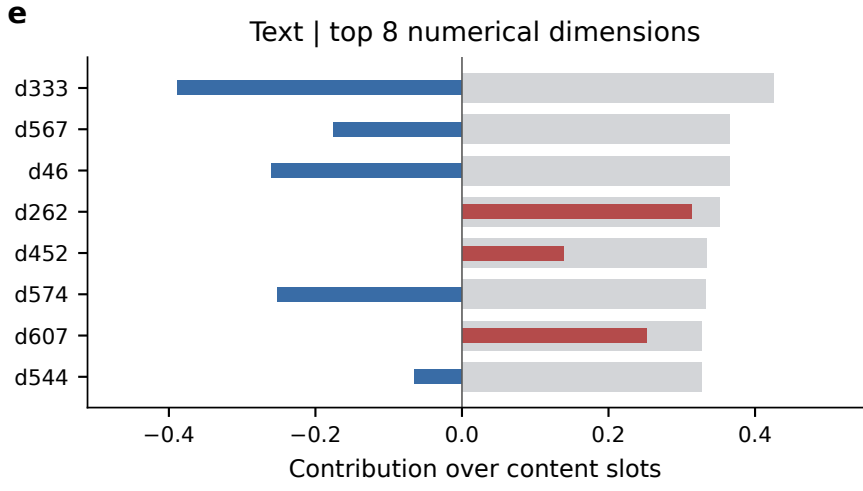
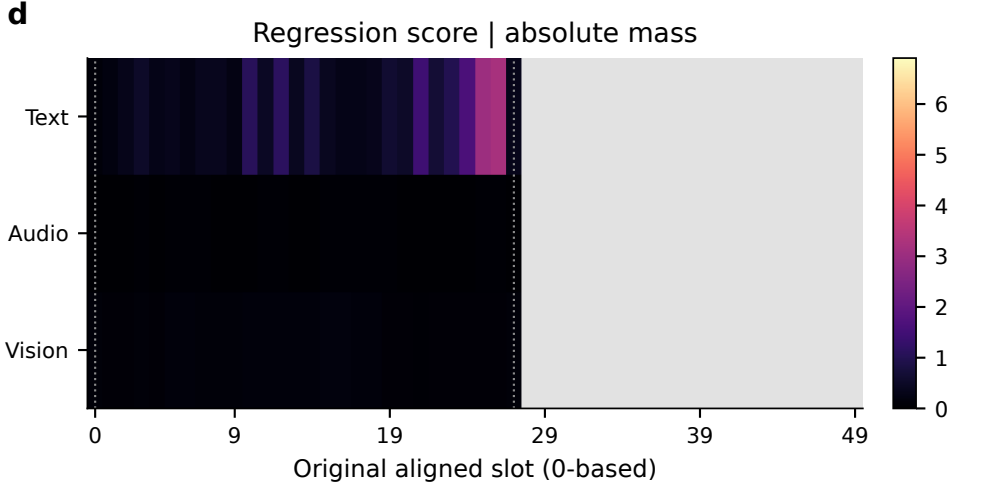
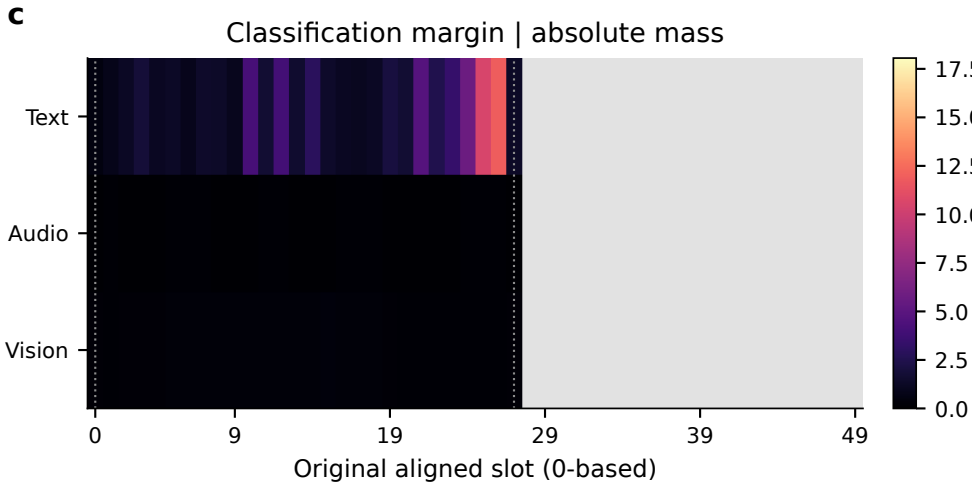
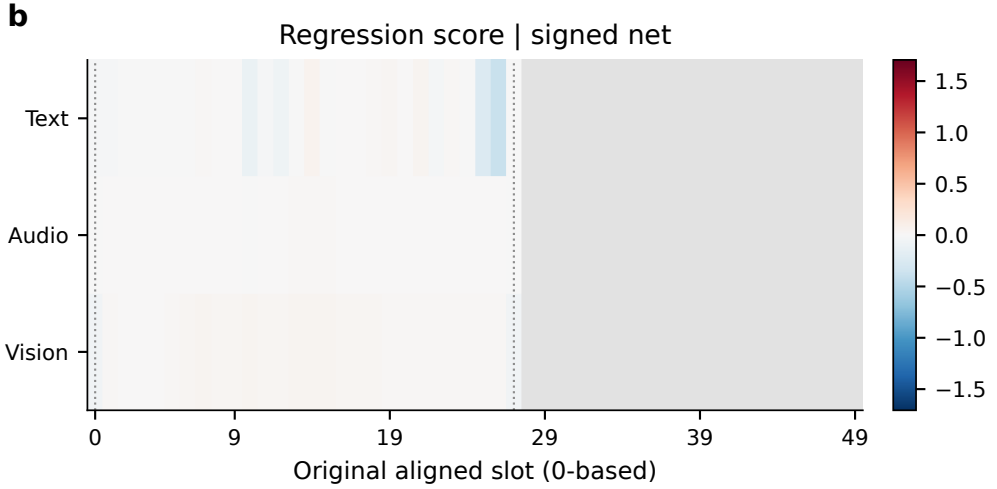
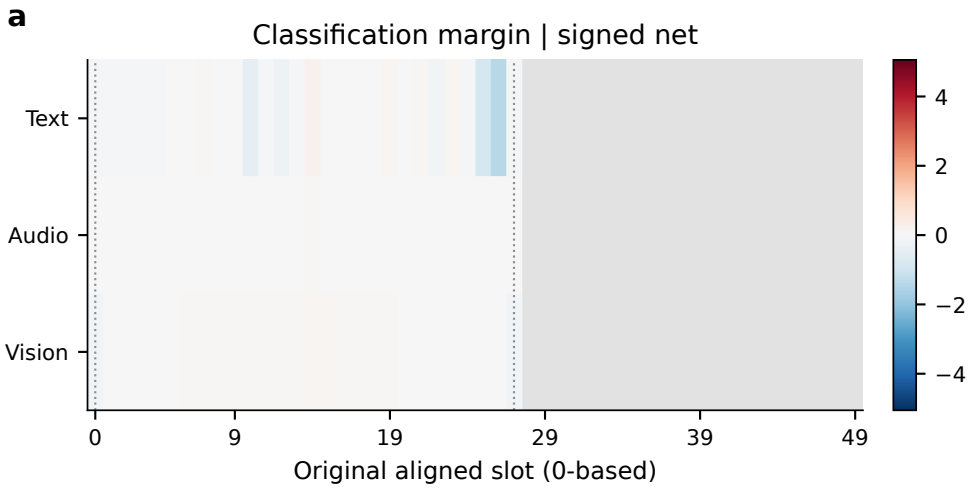
Attachment 4 | Sample 03 | Time-step and feature evidence

Predicted class 1 | Class probabilities [0: 0.251, 1: 0.432, 2: 0.316] | Regression 0.0704



Attachment 4 | Sample 04 | Time-step and feature evidence

Predicted class 2 | Class probabilities [0: 0.309, 1: 0.088, 2: 0.603] | Regression 0.4894

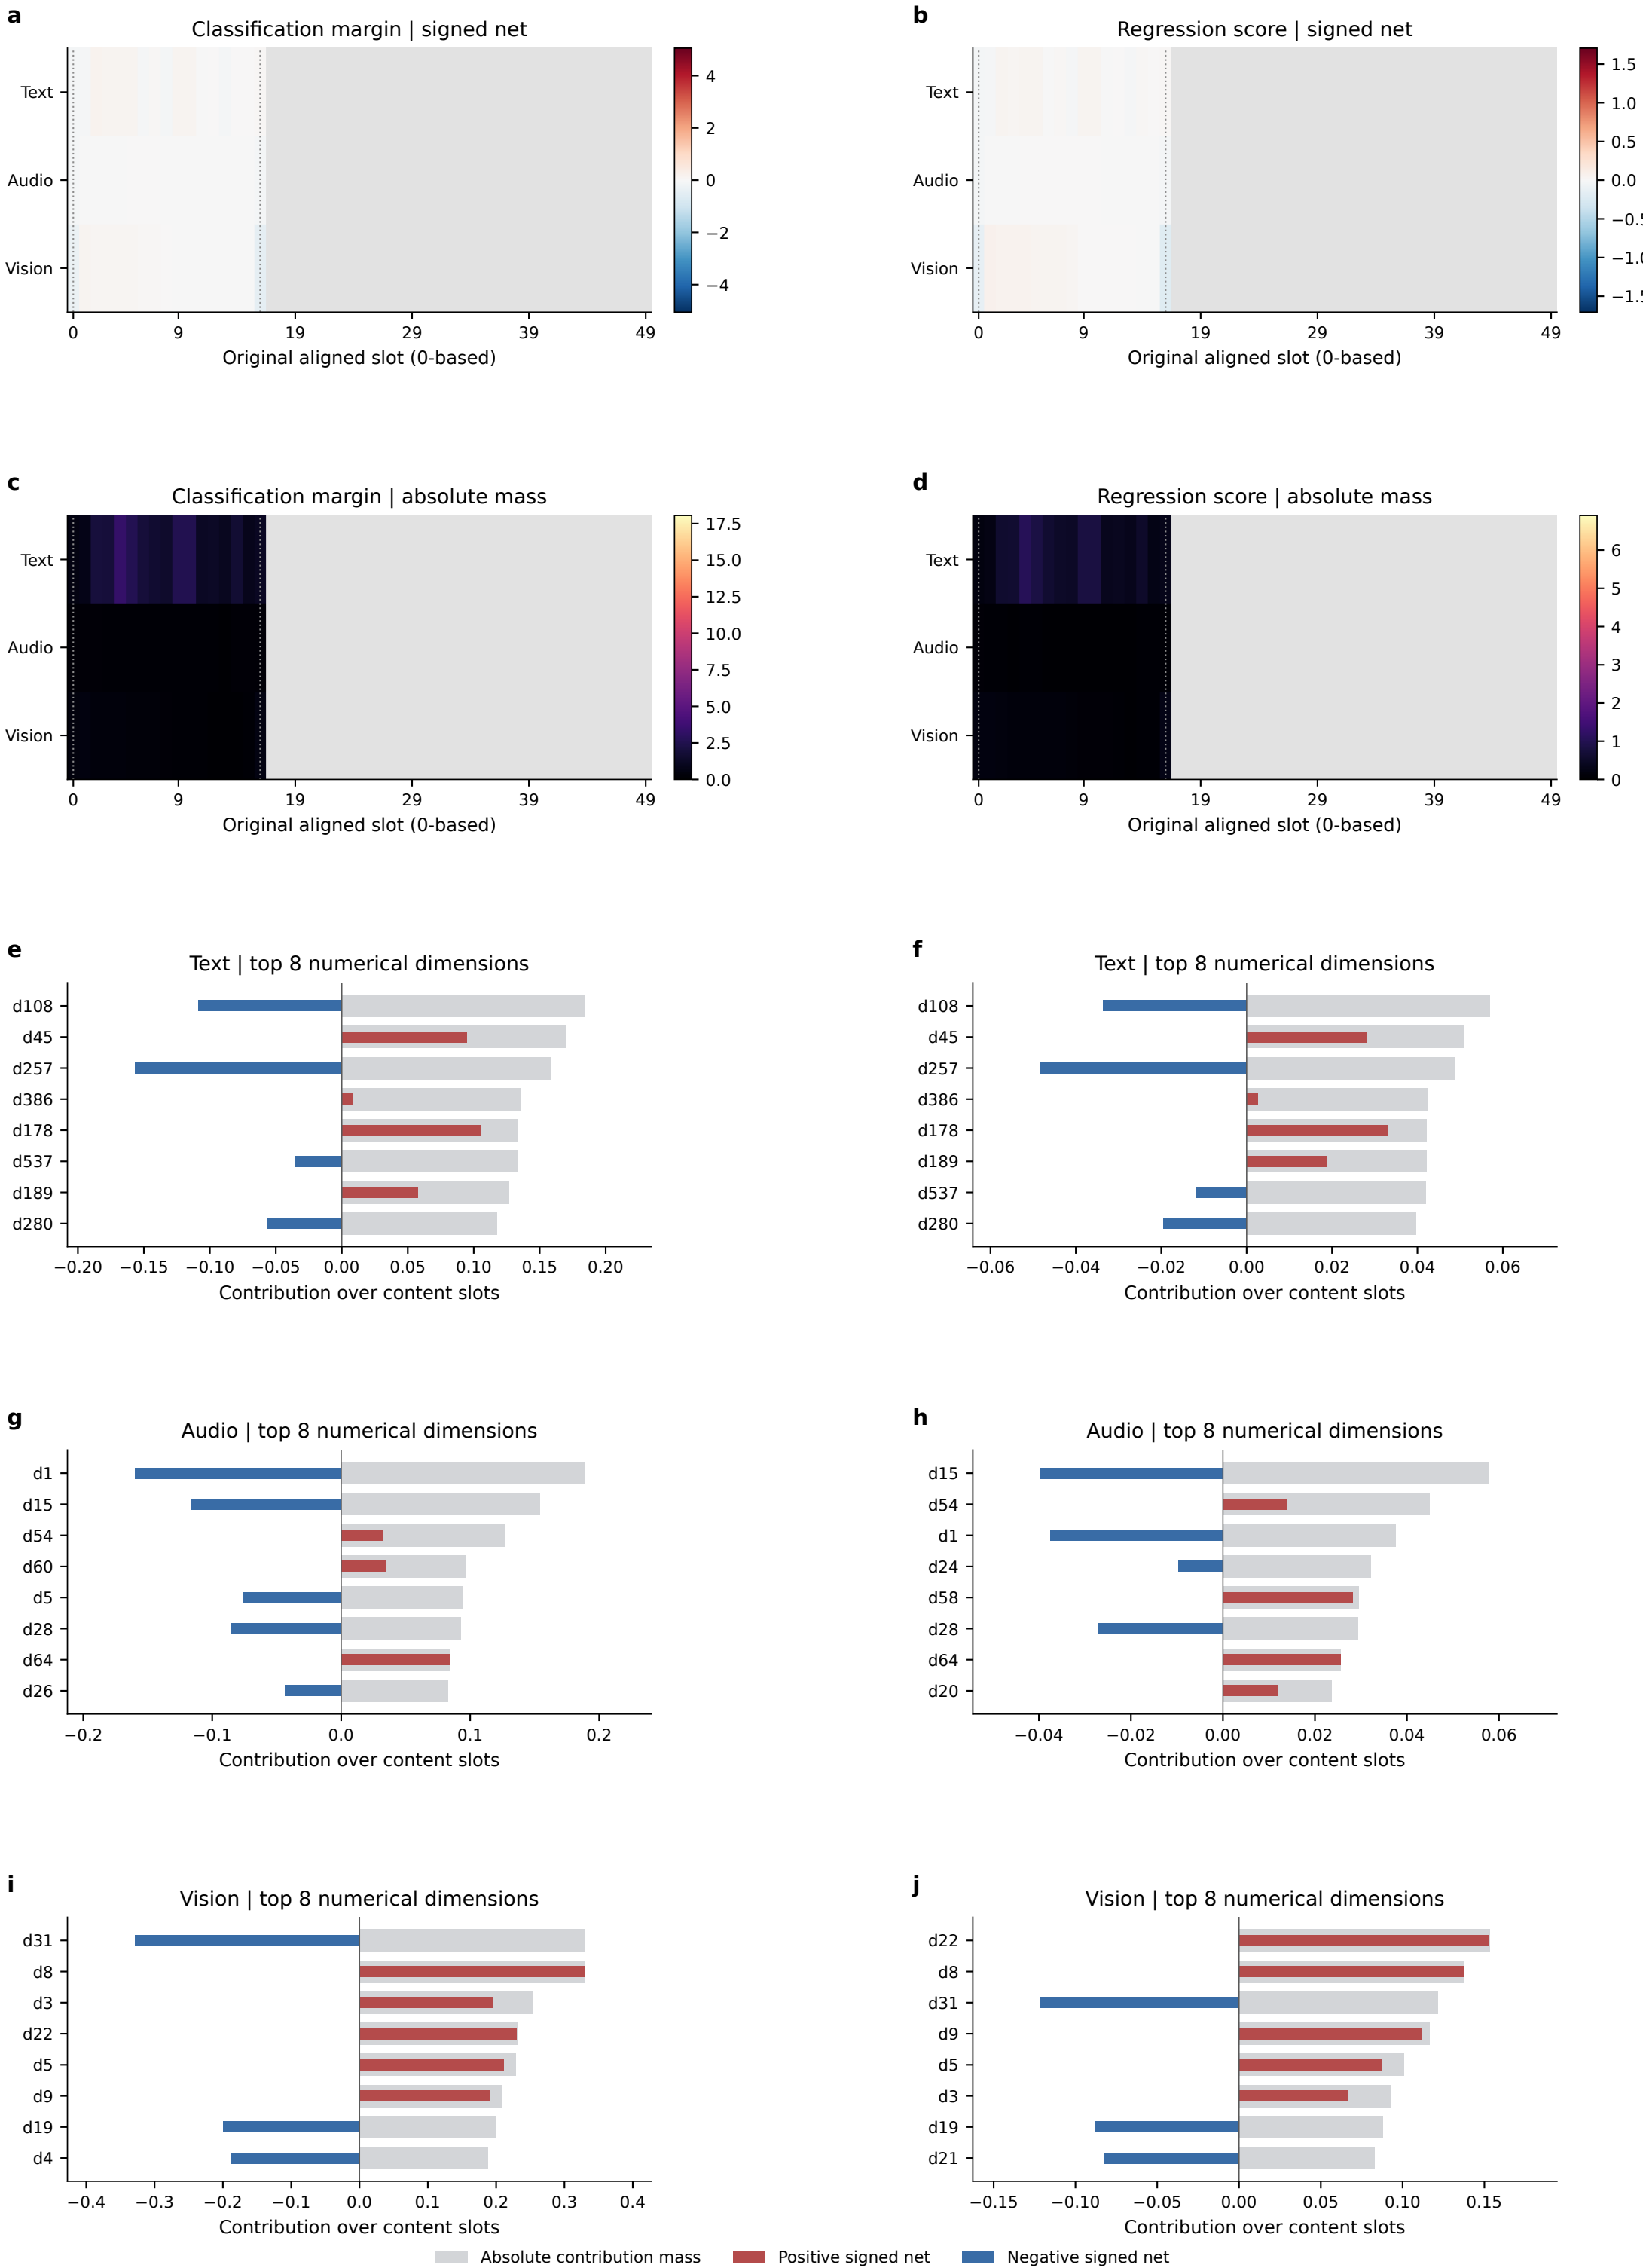


Gray cells: padding; dotted lines: CLS/SEP. Absolute mass

Train-mean reference. Heatmap scales are shared across all 20 samples for each target. Gray cells: padding; dotted lines: CLS/SEP. Absolute mass is not a modality share. Dimension rankings exclude specials and padding. Dimensions are numerical indices, not semantic labels or video timestamps.

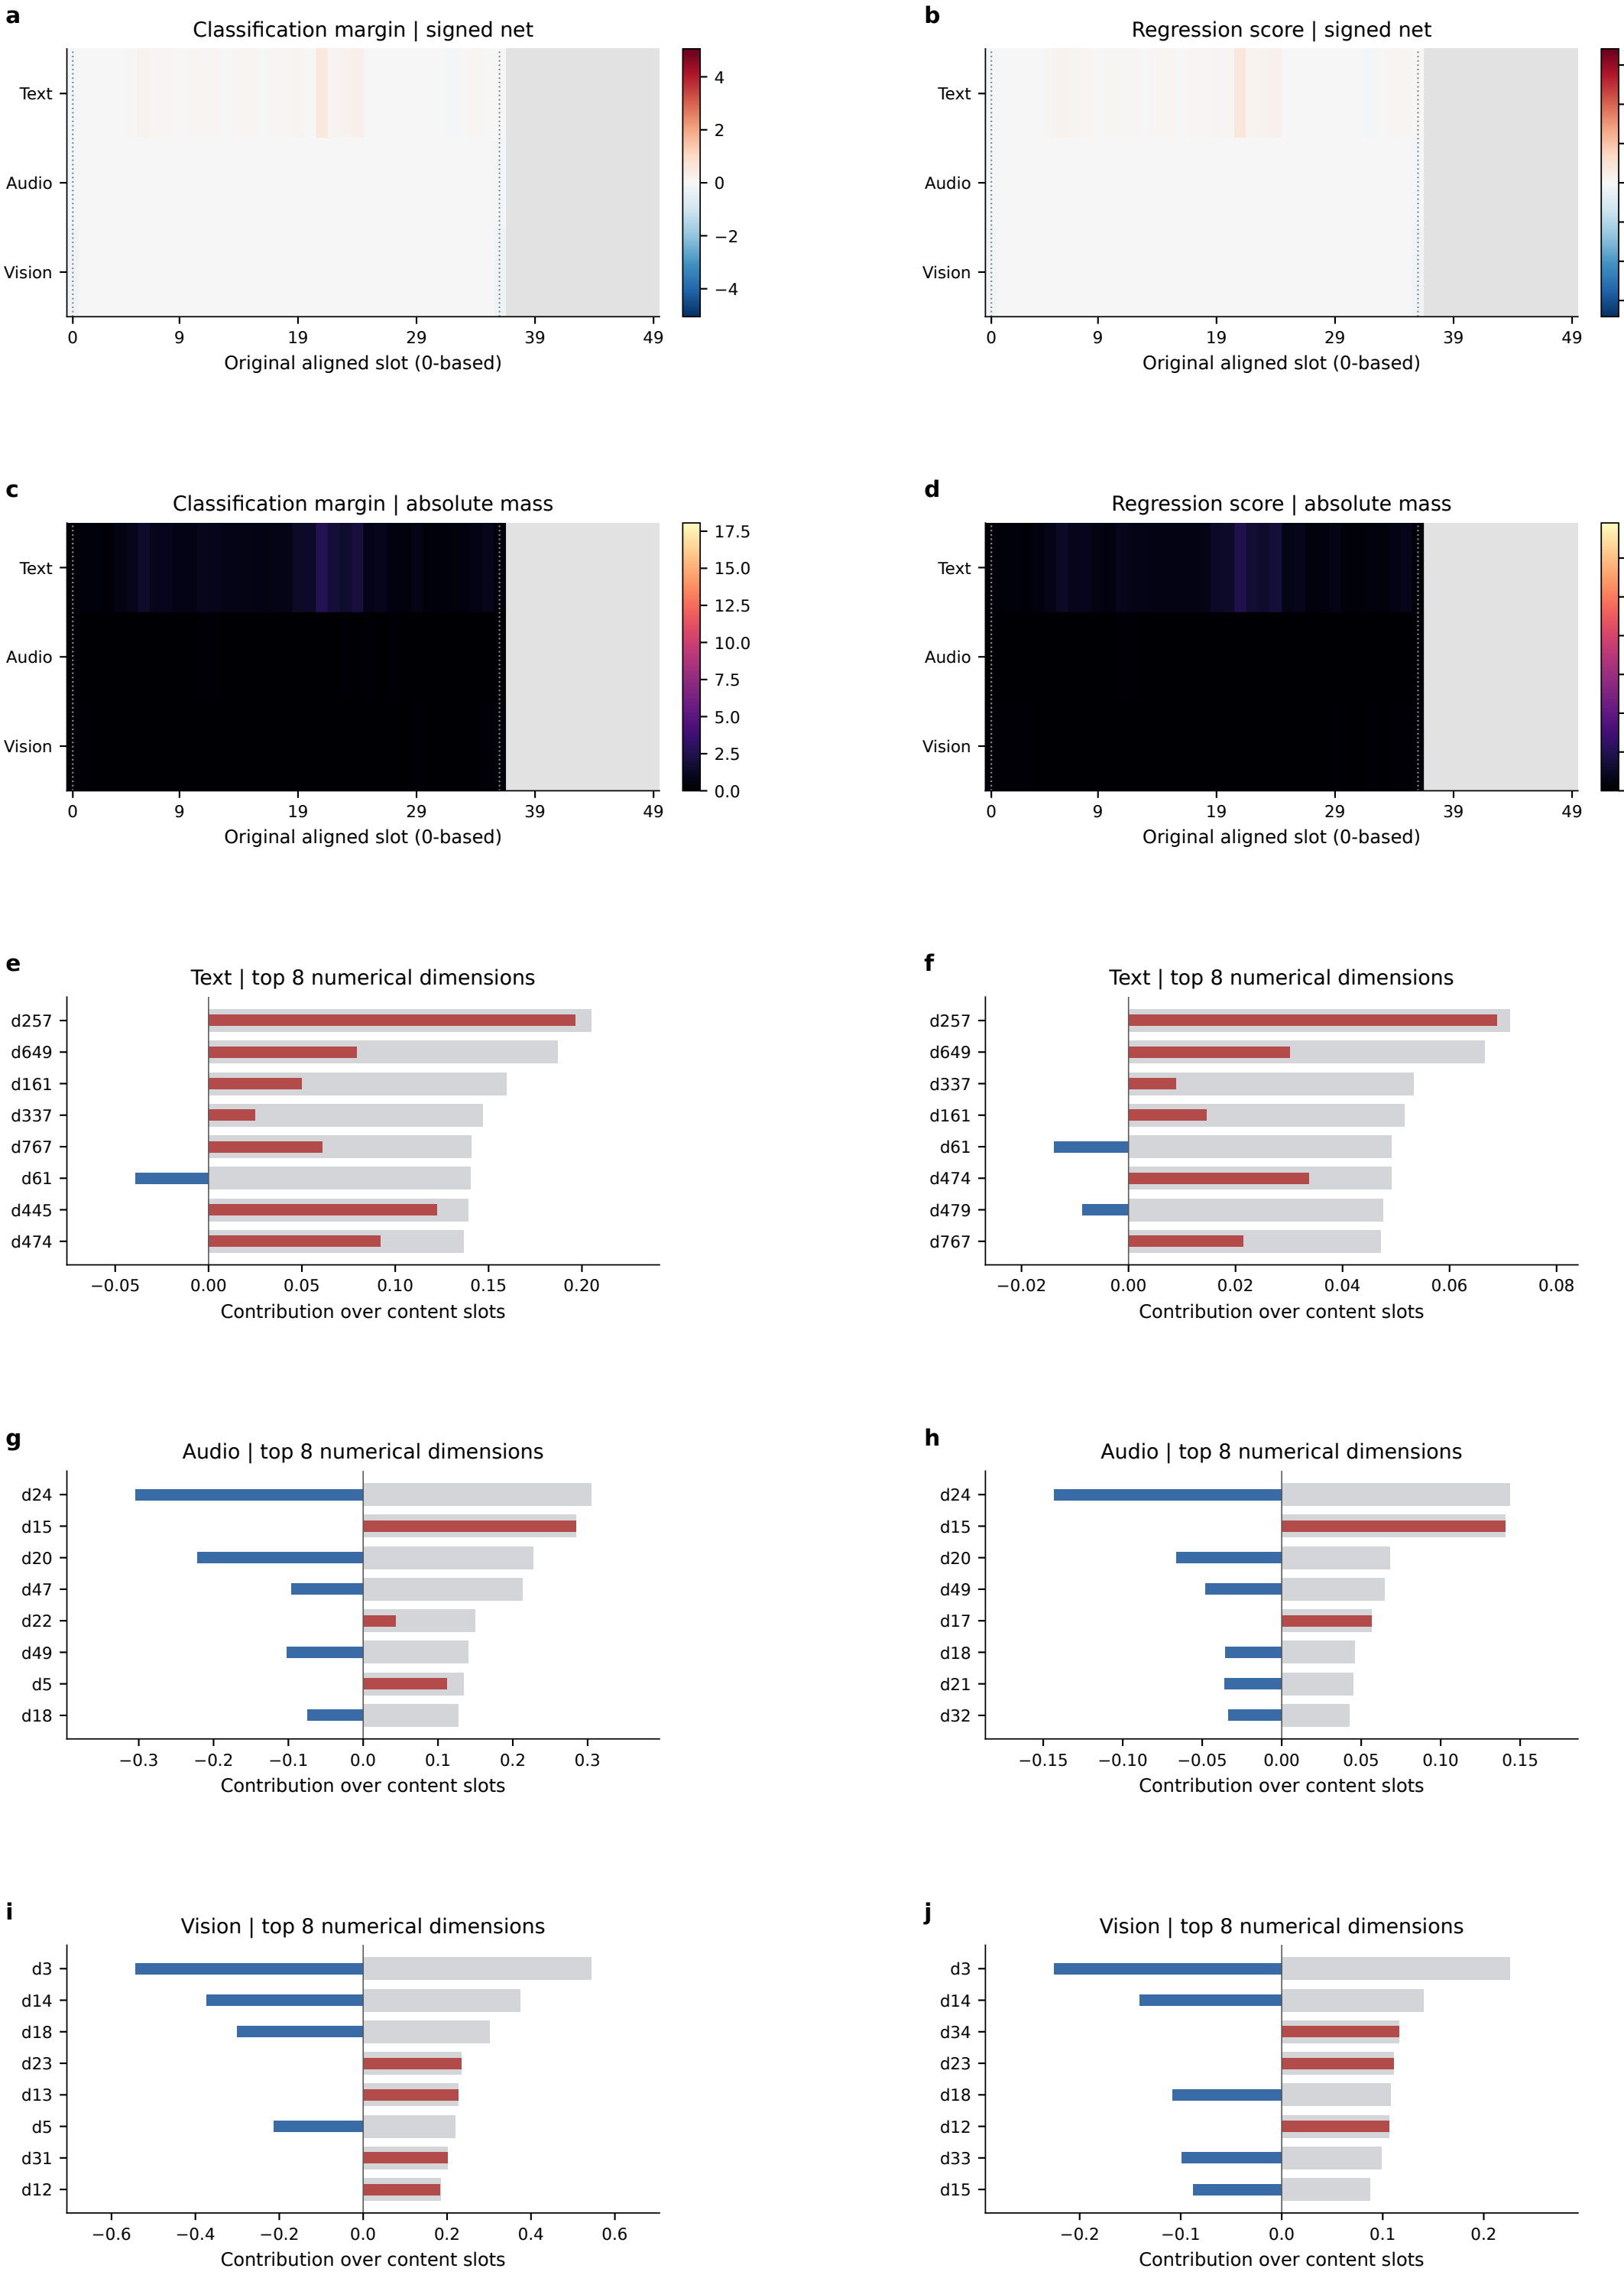
Attachment 4 | Sample 05 | Time-step and feature evidence

Predicted class 2 | Class probabilities [0: 0.031, 1: 0.208, 2: 0.761] | Regression 0.6403



Attachment 4 | Sample 06 | Time-step and feature evidence

Predicted class 2 | Class probabilities [0: 0.027, 1: 0.114, 2: 0.859] | Regression 0.8748

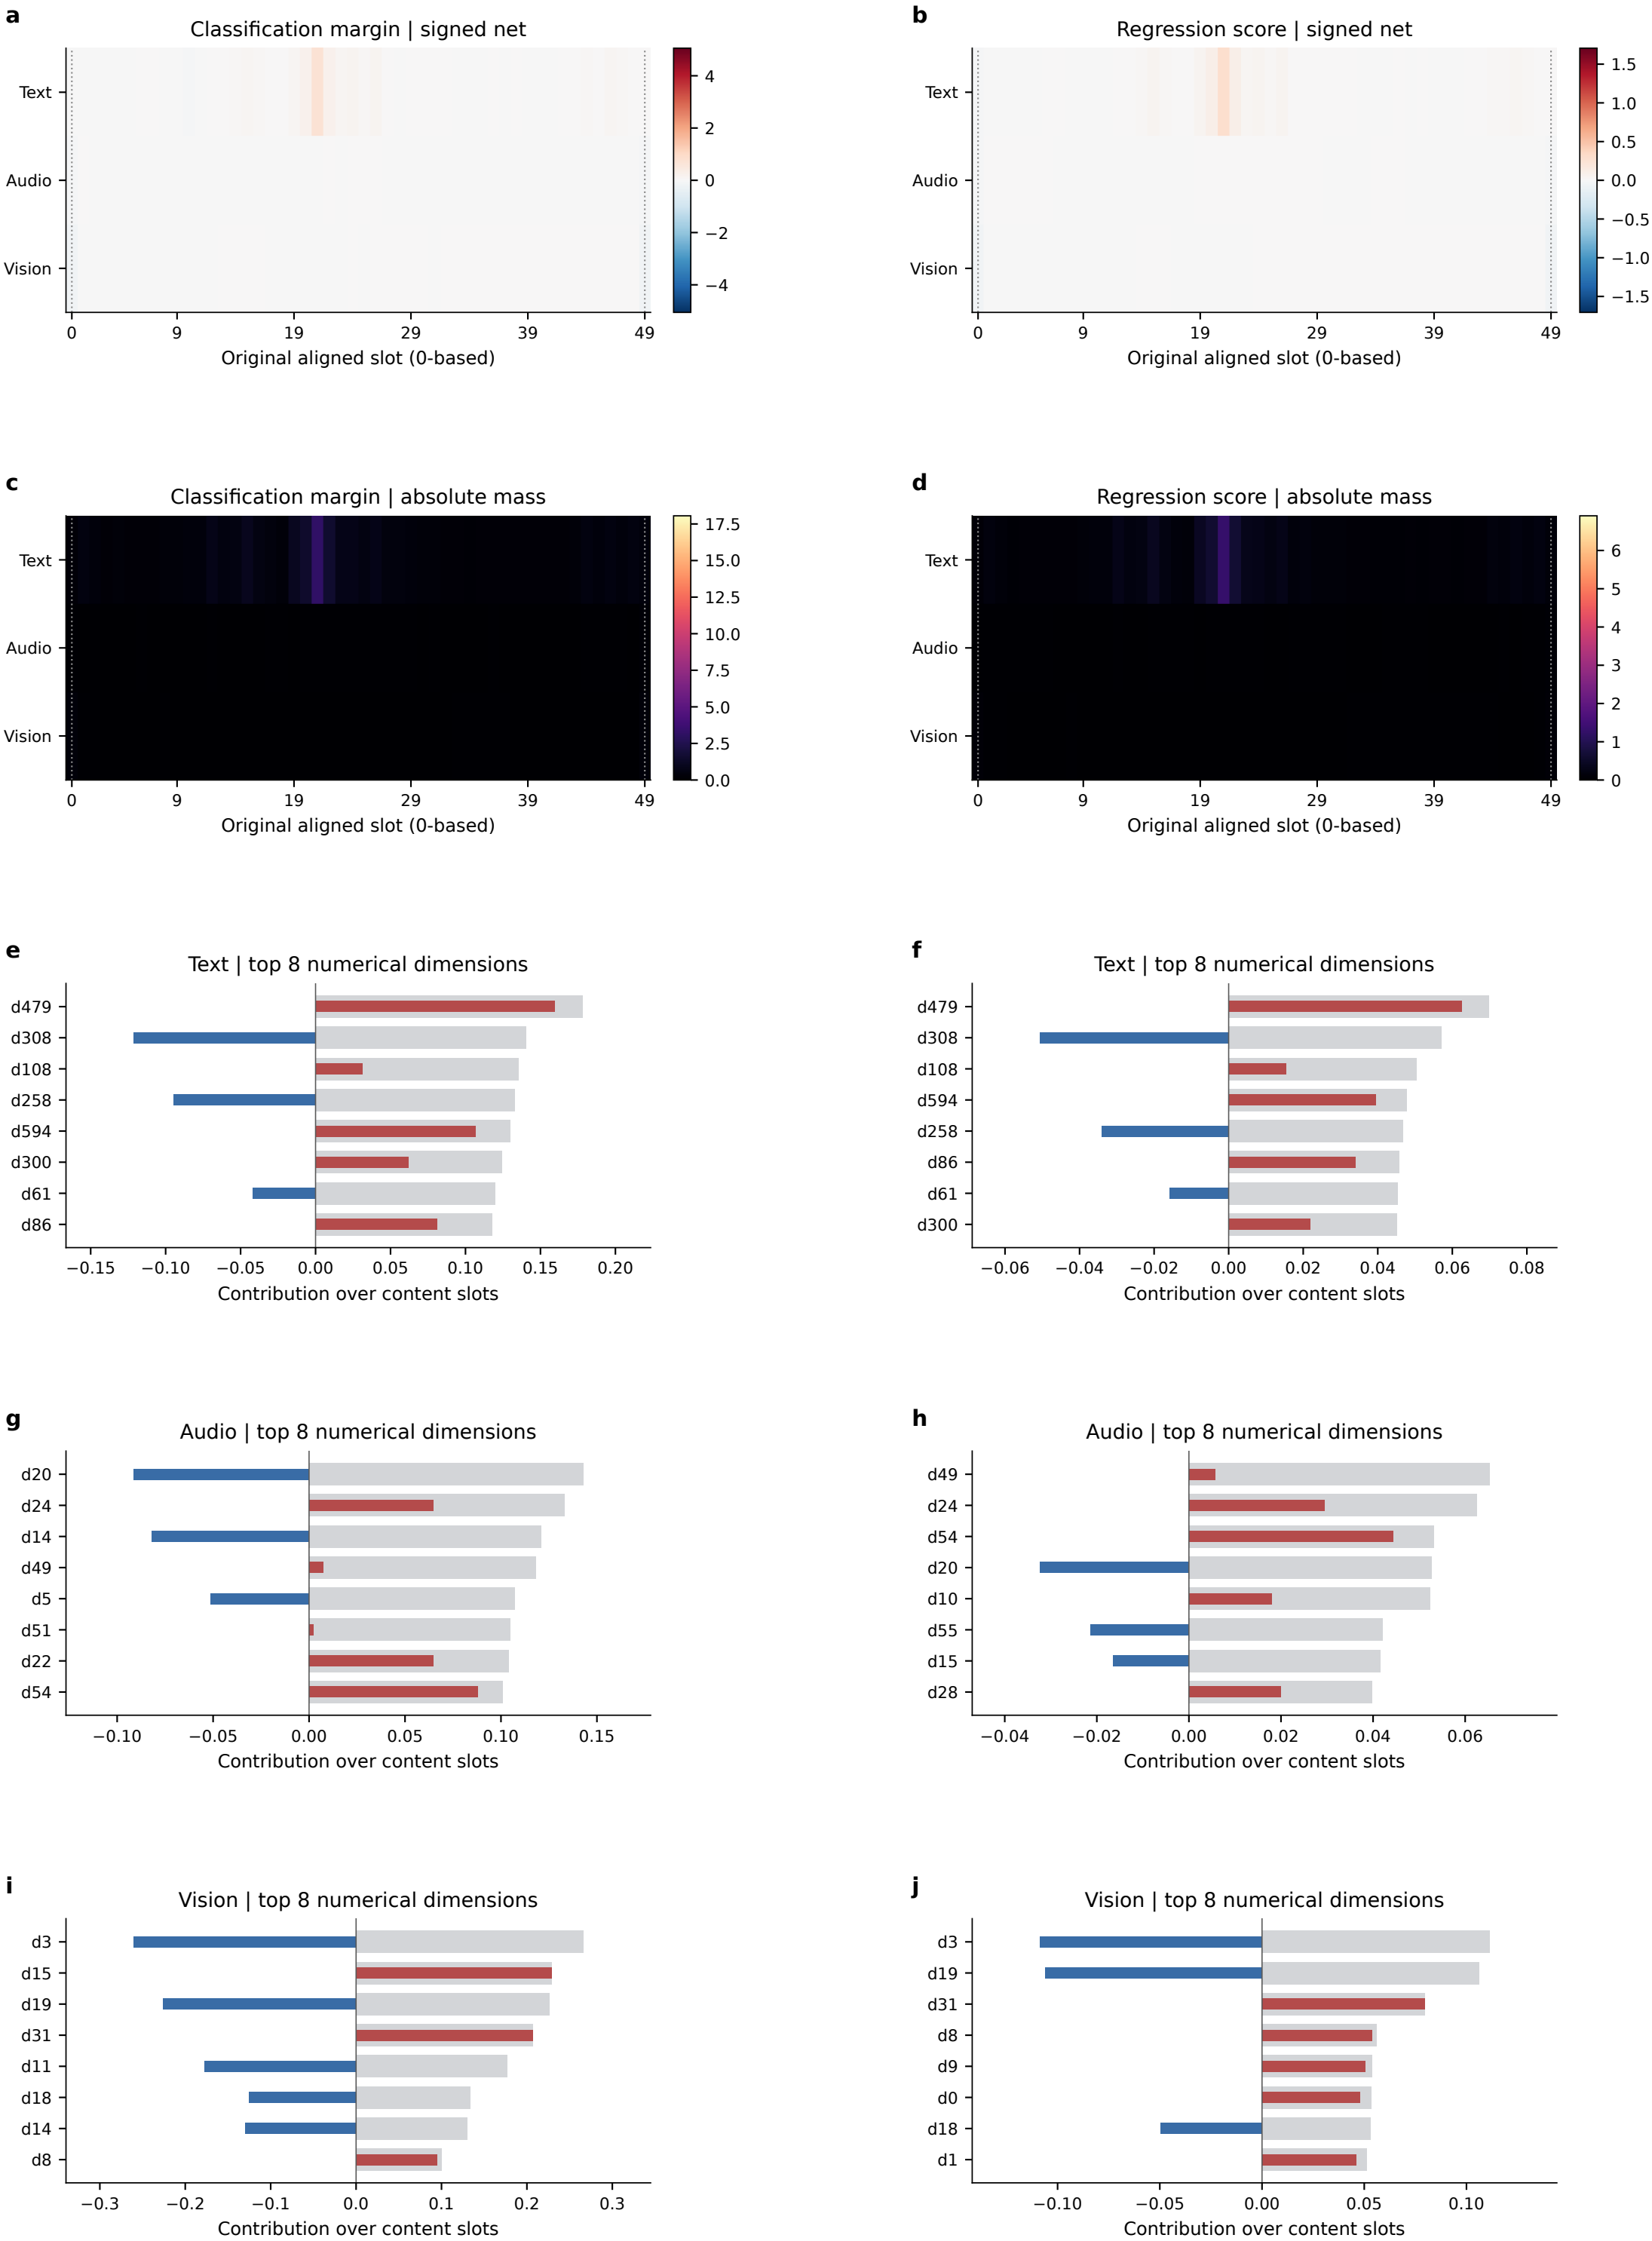


Train-mean reference. Heatmap scales are shared across all 20 samples for each target. Gray cells: padding; dotted lines: CLS/SEP. Absolute mass is not a modality share. Dimension rankings exclude specials and padding. Dimensions are numerical indices, not semantic labels or video timestamps.

Attachment 4 | Sample 07 | Time-step and feature evidence

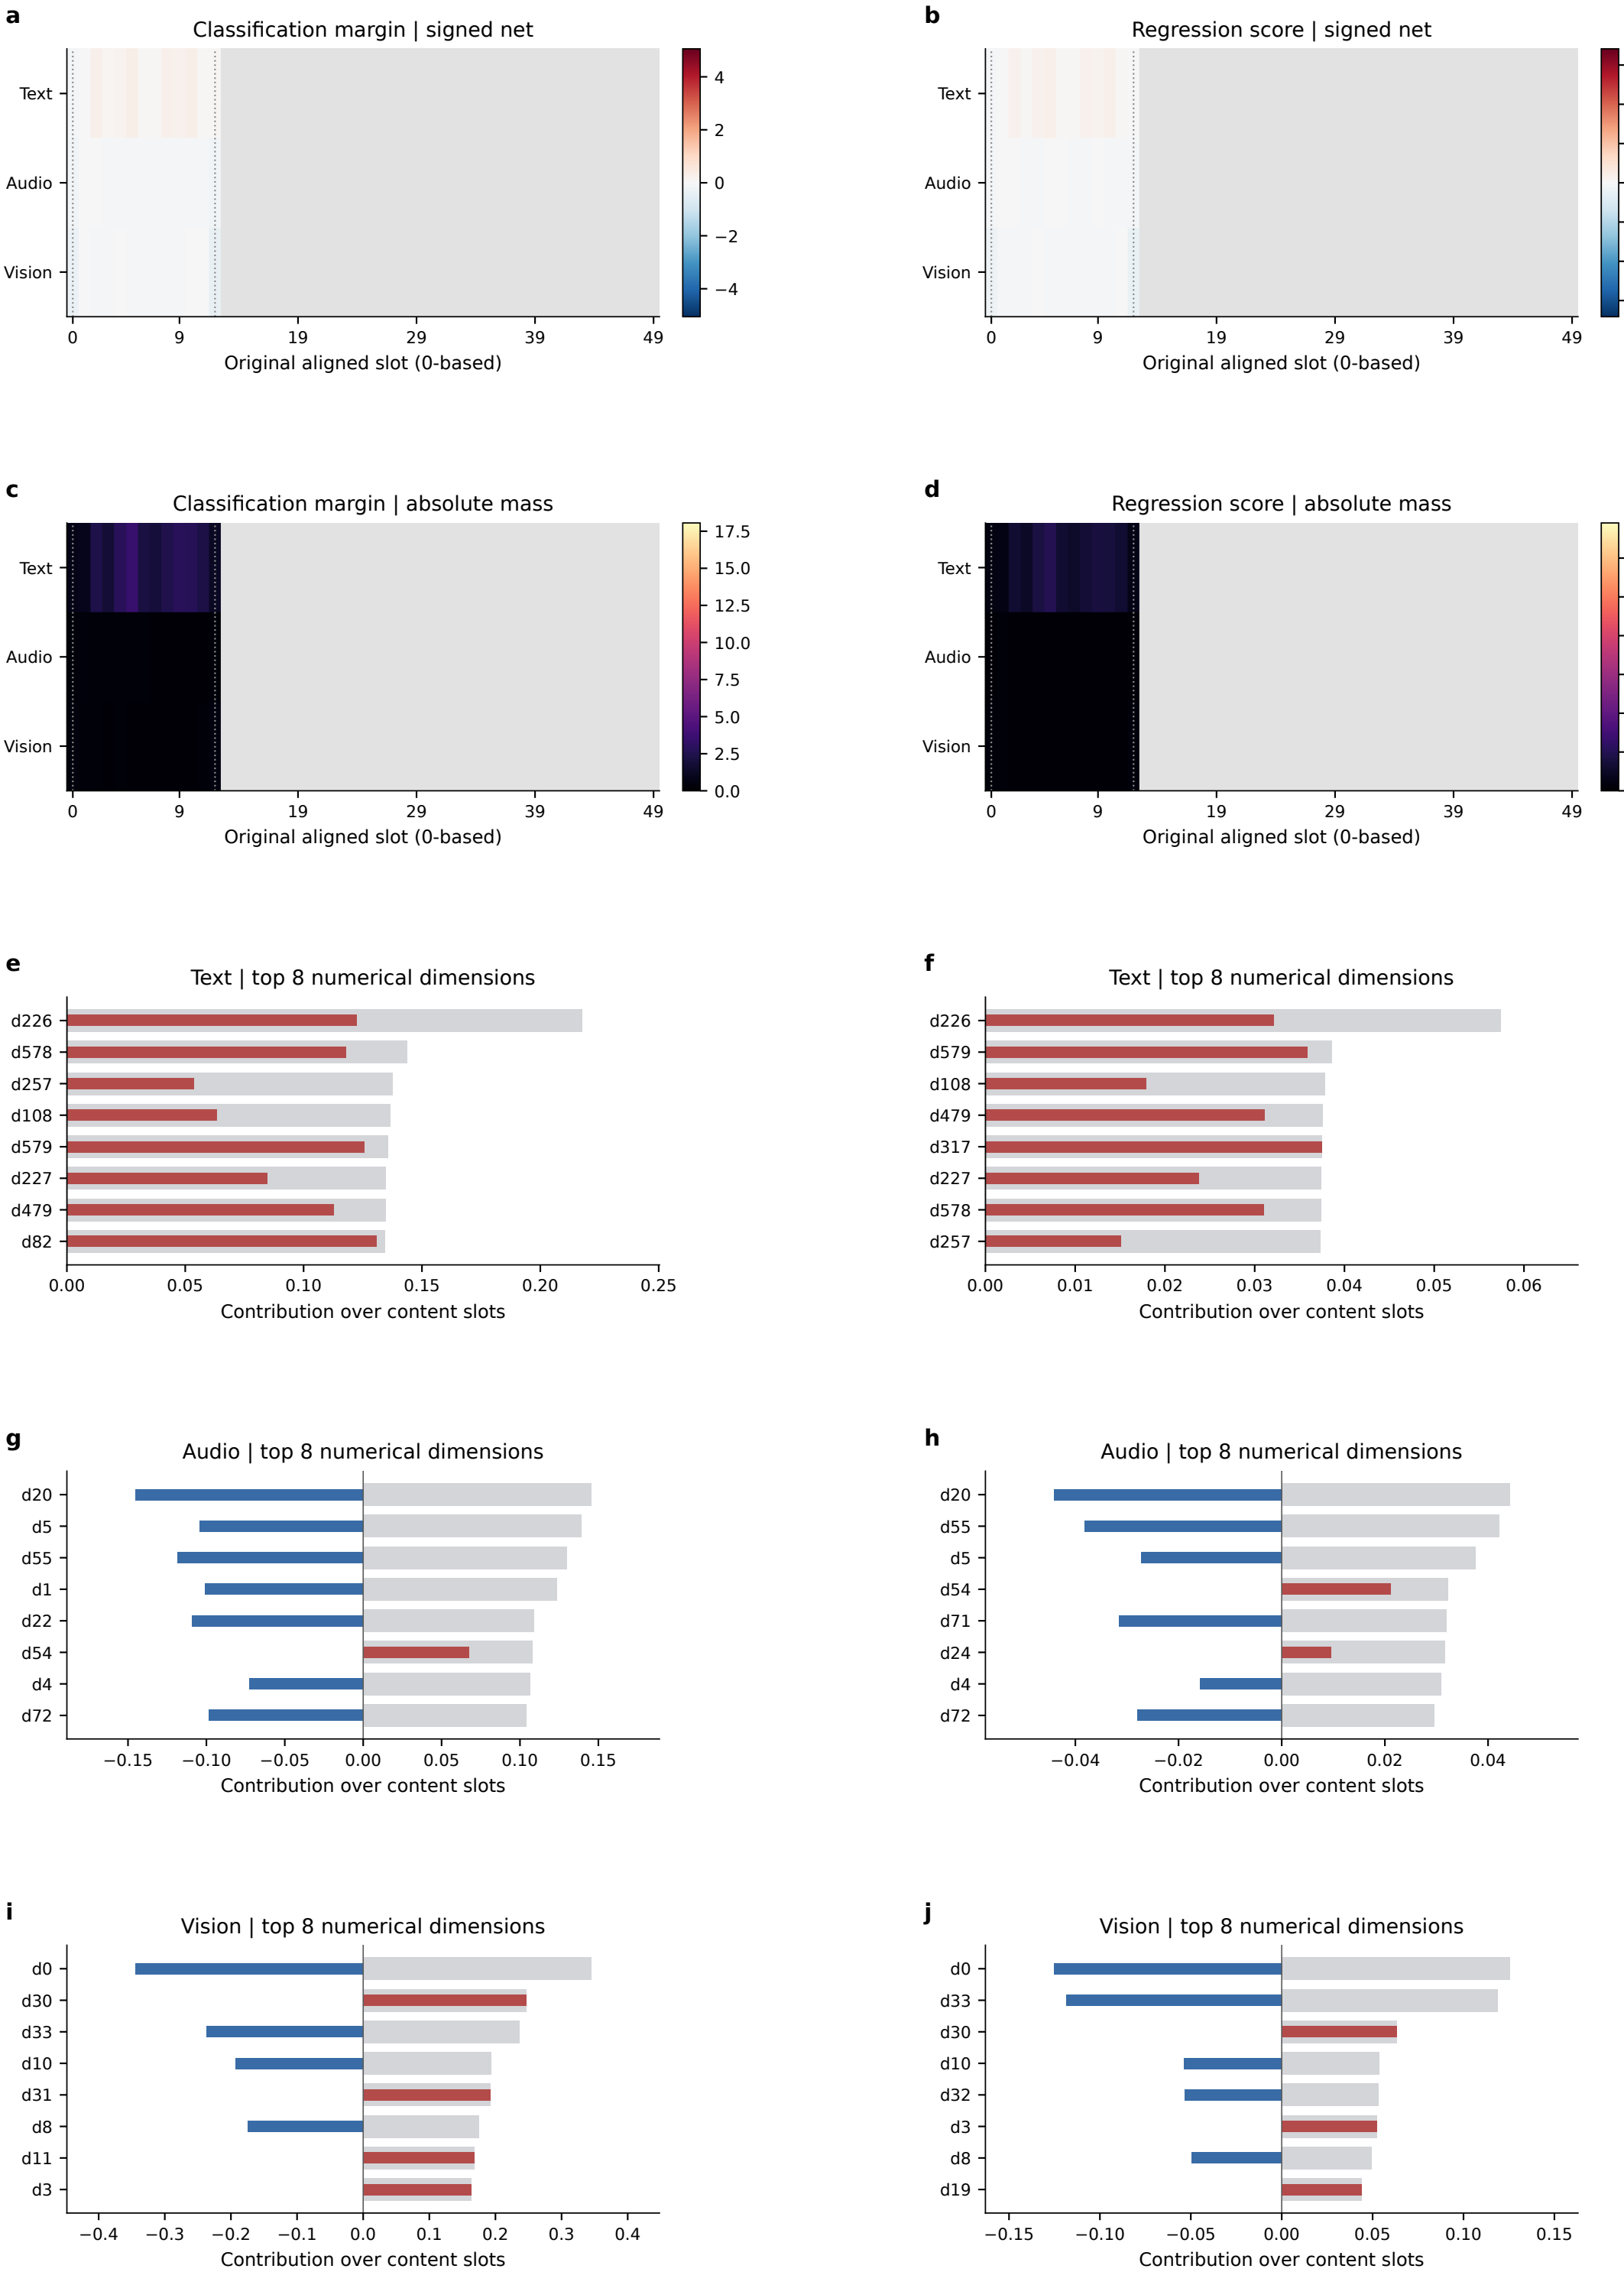
Predicted class 2 | Class probabilities [0: 0.014, 1: 0.053, 2: 0.933] | Regression 1.2134

Text truncated: 48 of 63 content tokens retained.



Attachment 4 | Sample 08 | Time-step and feature evidence

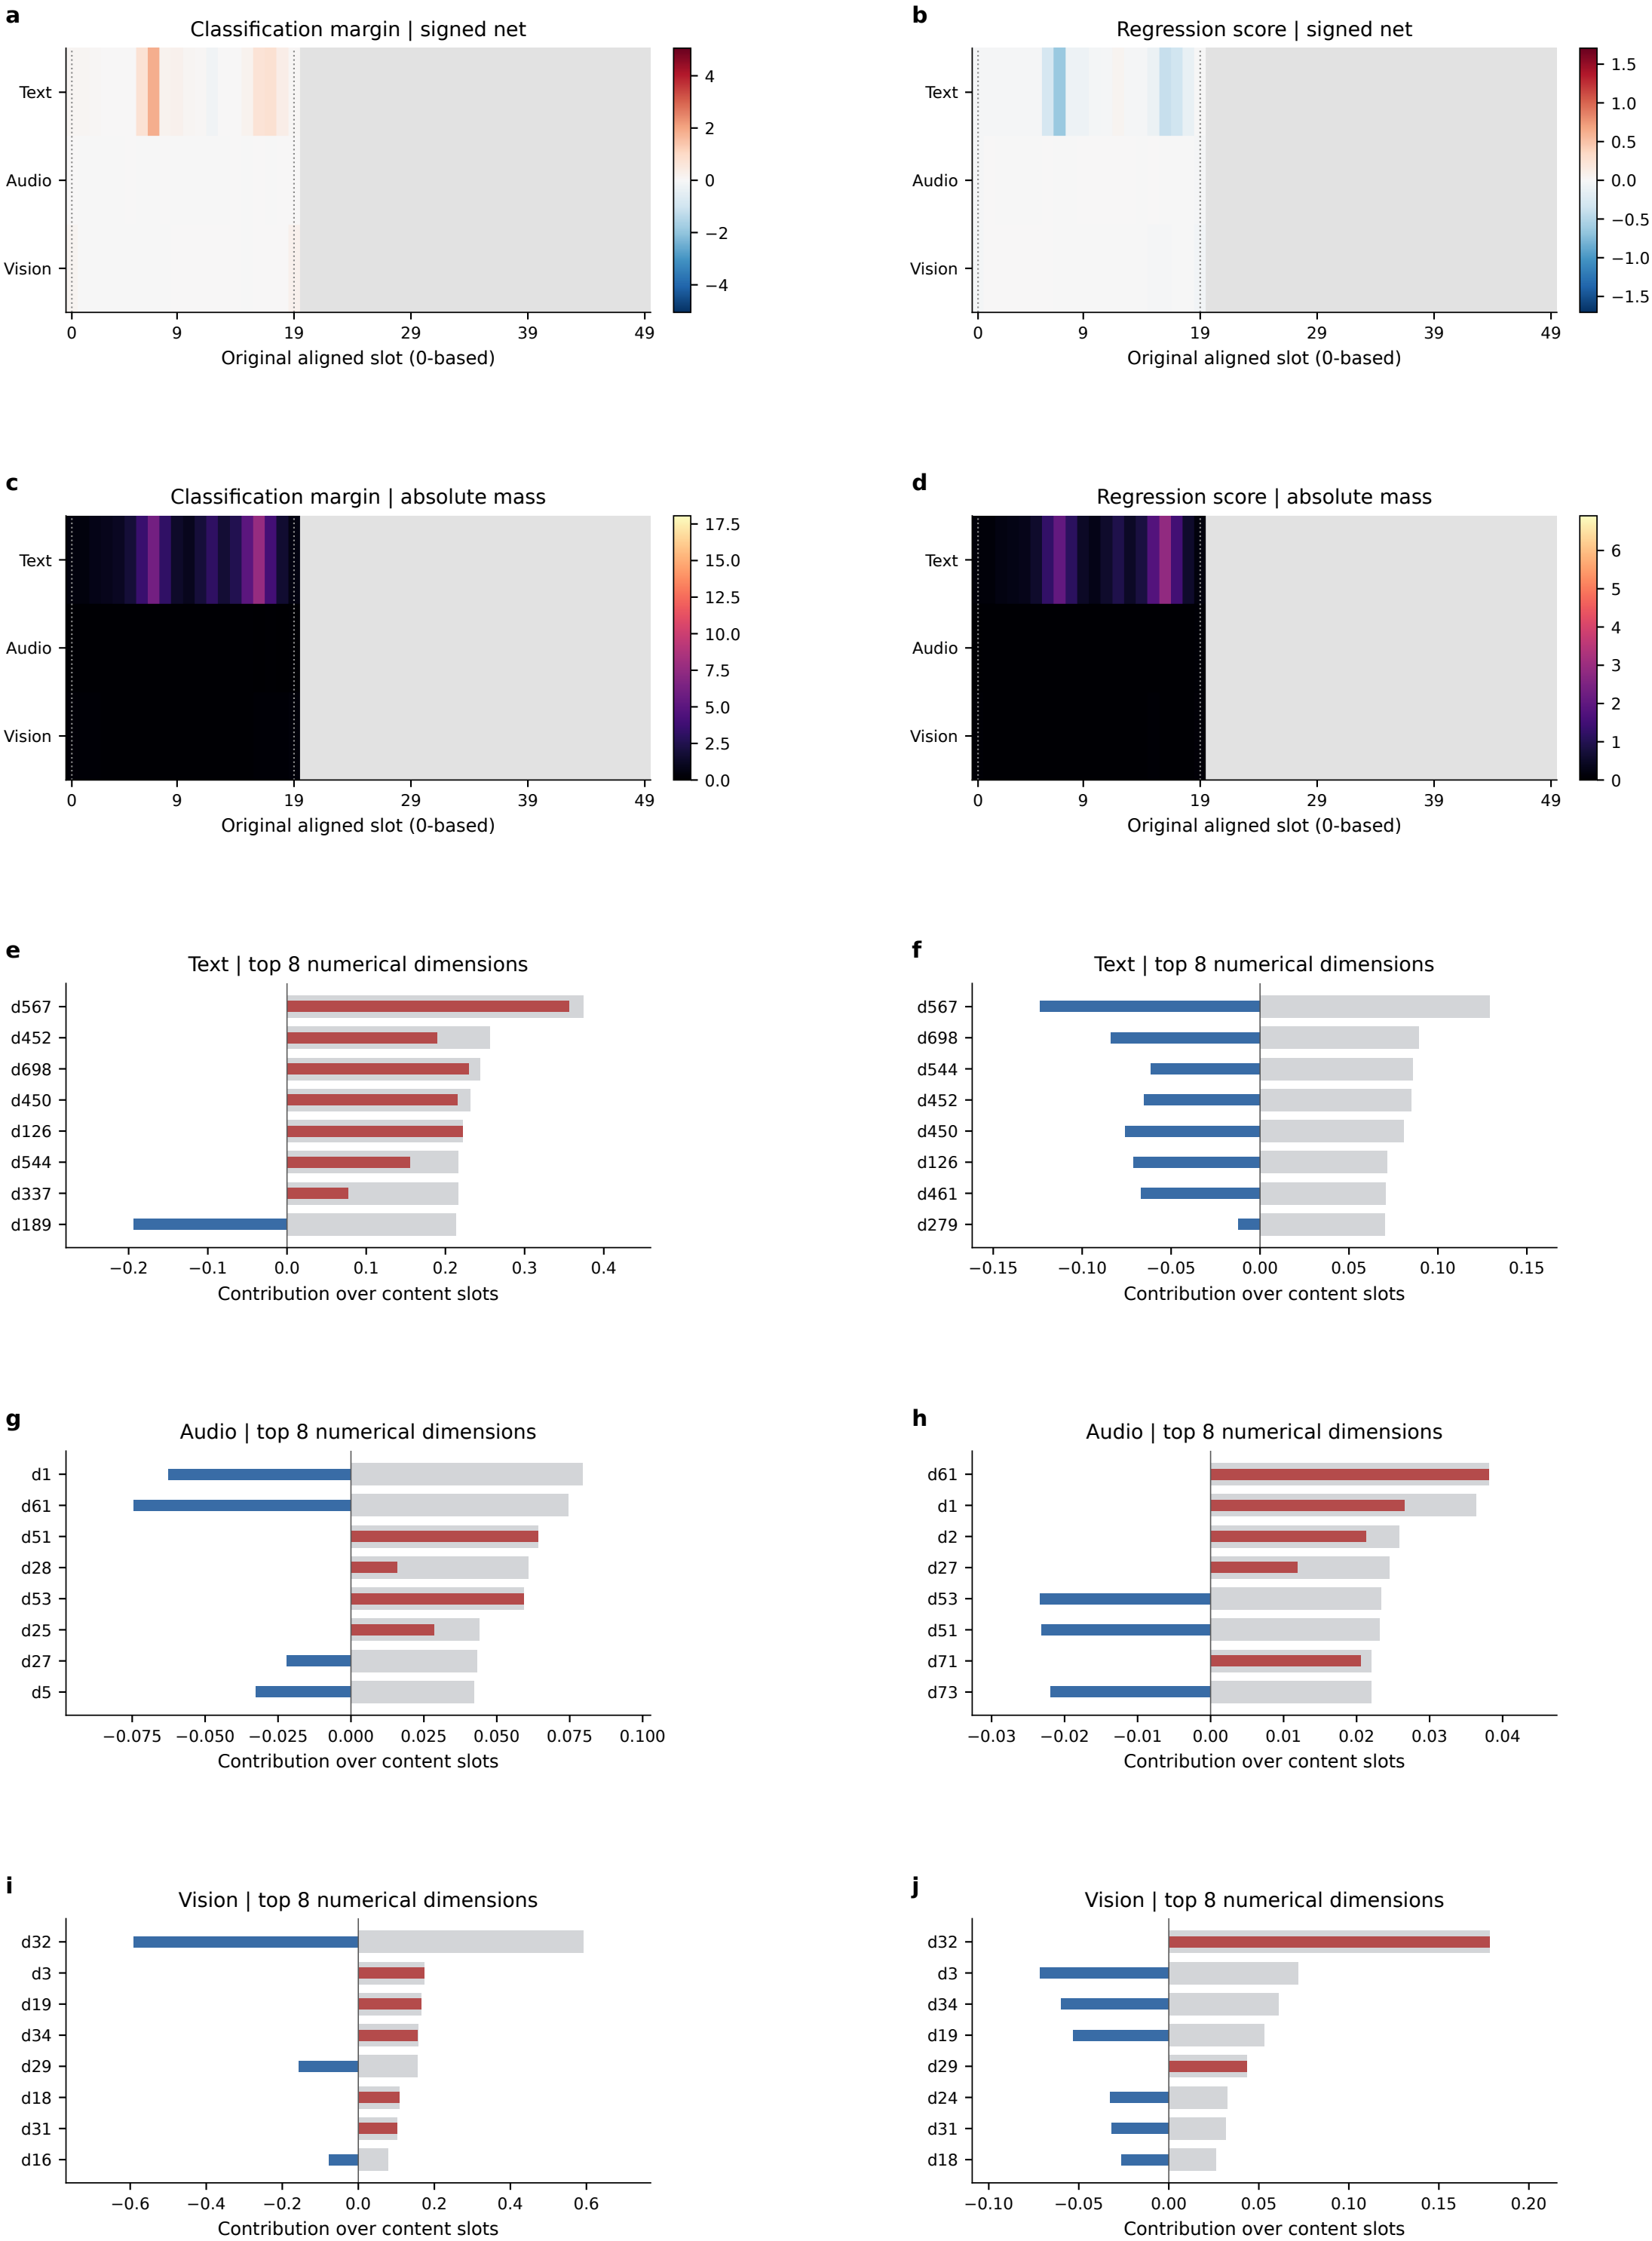
Predicted class 2 | Class probabilities [0: 0.052, 1: 0.247, 2: 0.701] | Regression 0.3236



Train-mean reference. Heatmap scales are shared across all 20 samples for each target. Gray cells: padding; dotted lines: CLS/SEP. Absolute mass is not a modality share. Dimension rankings exclude specials and padding. Dimensions are numerical indices, not semantic labels or video timestamps.

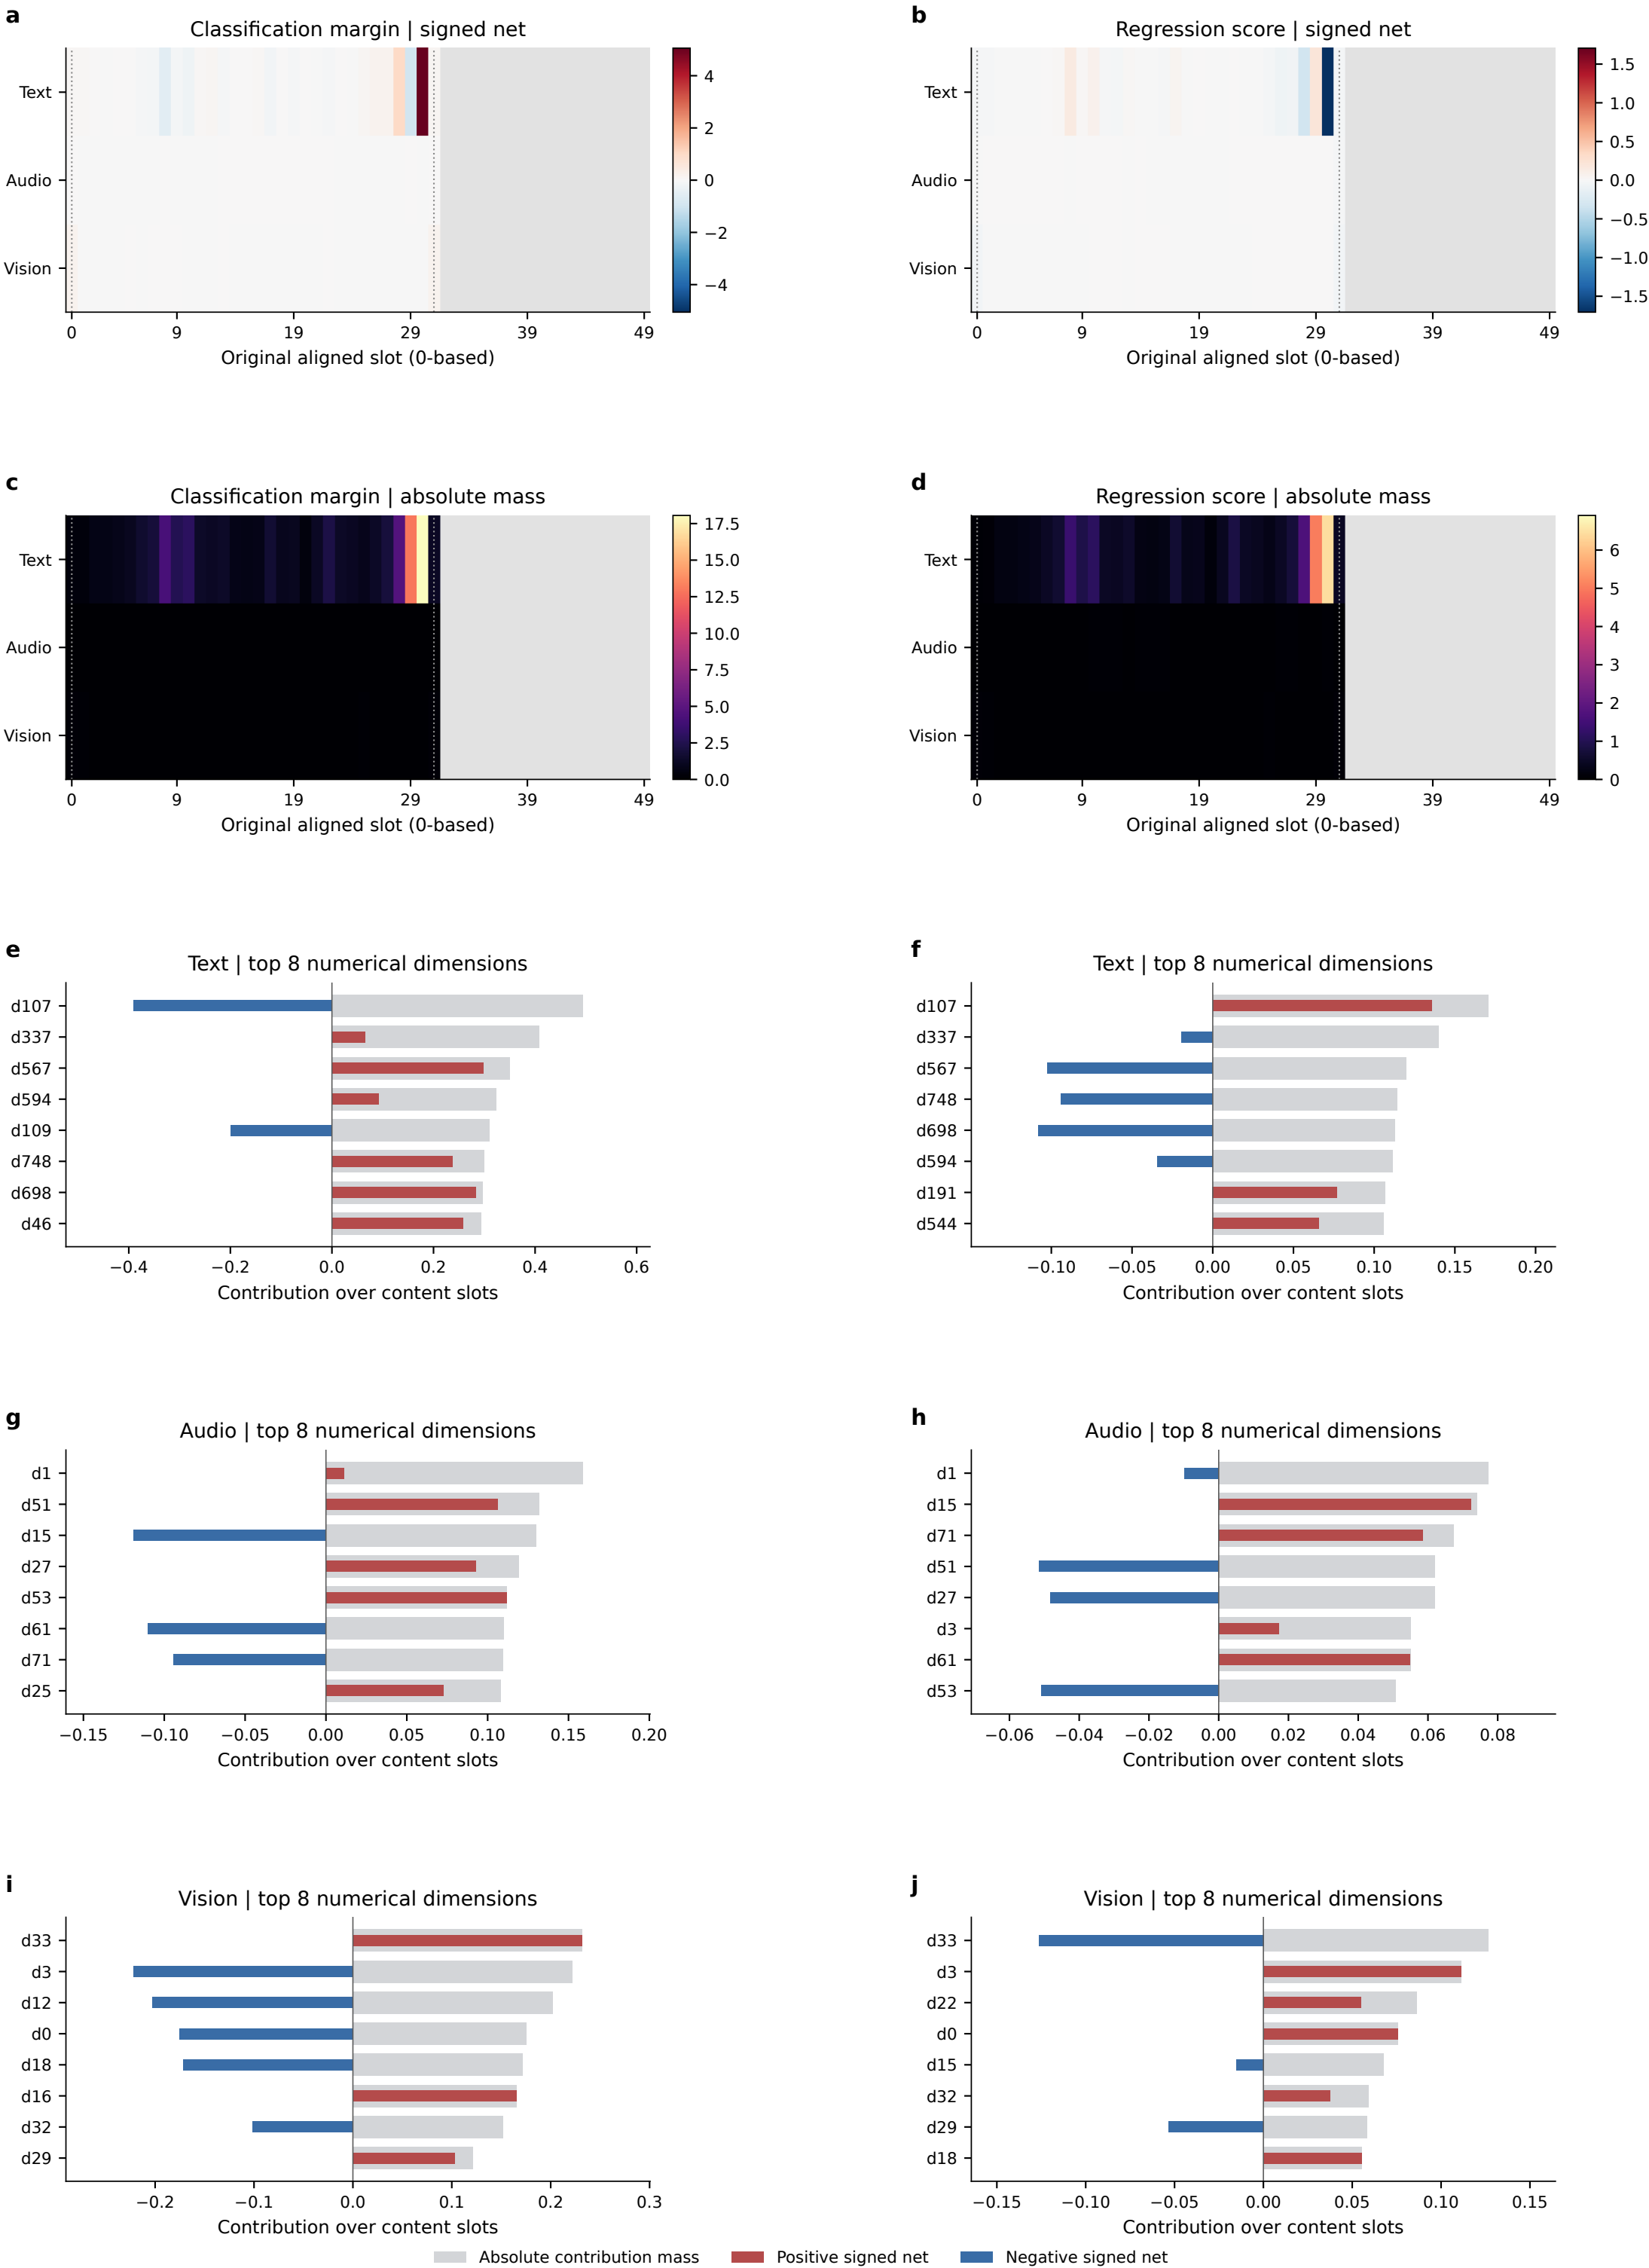
Attachment 4 | Sample 09 | Time-step and feature evidence

Predicted class 0 | Class probabilities [0: 0.922, 1: 0.037, 2: 0.041] | Regression -2.0012



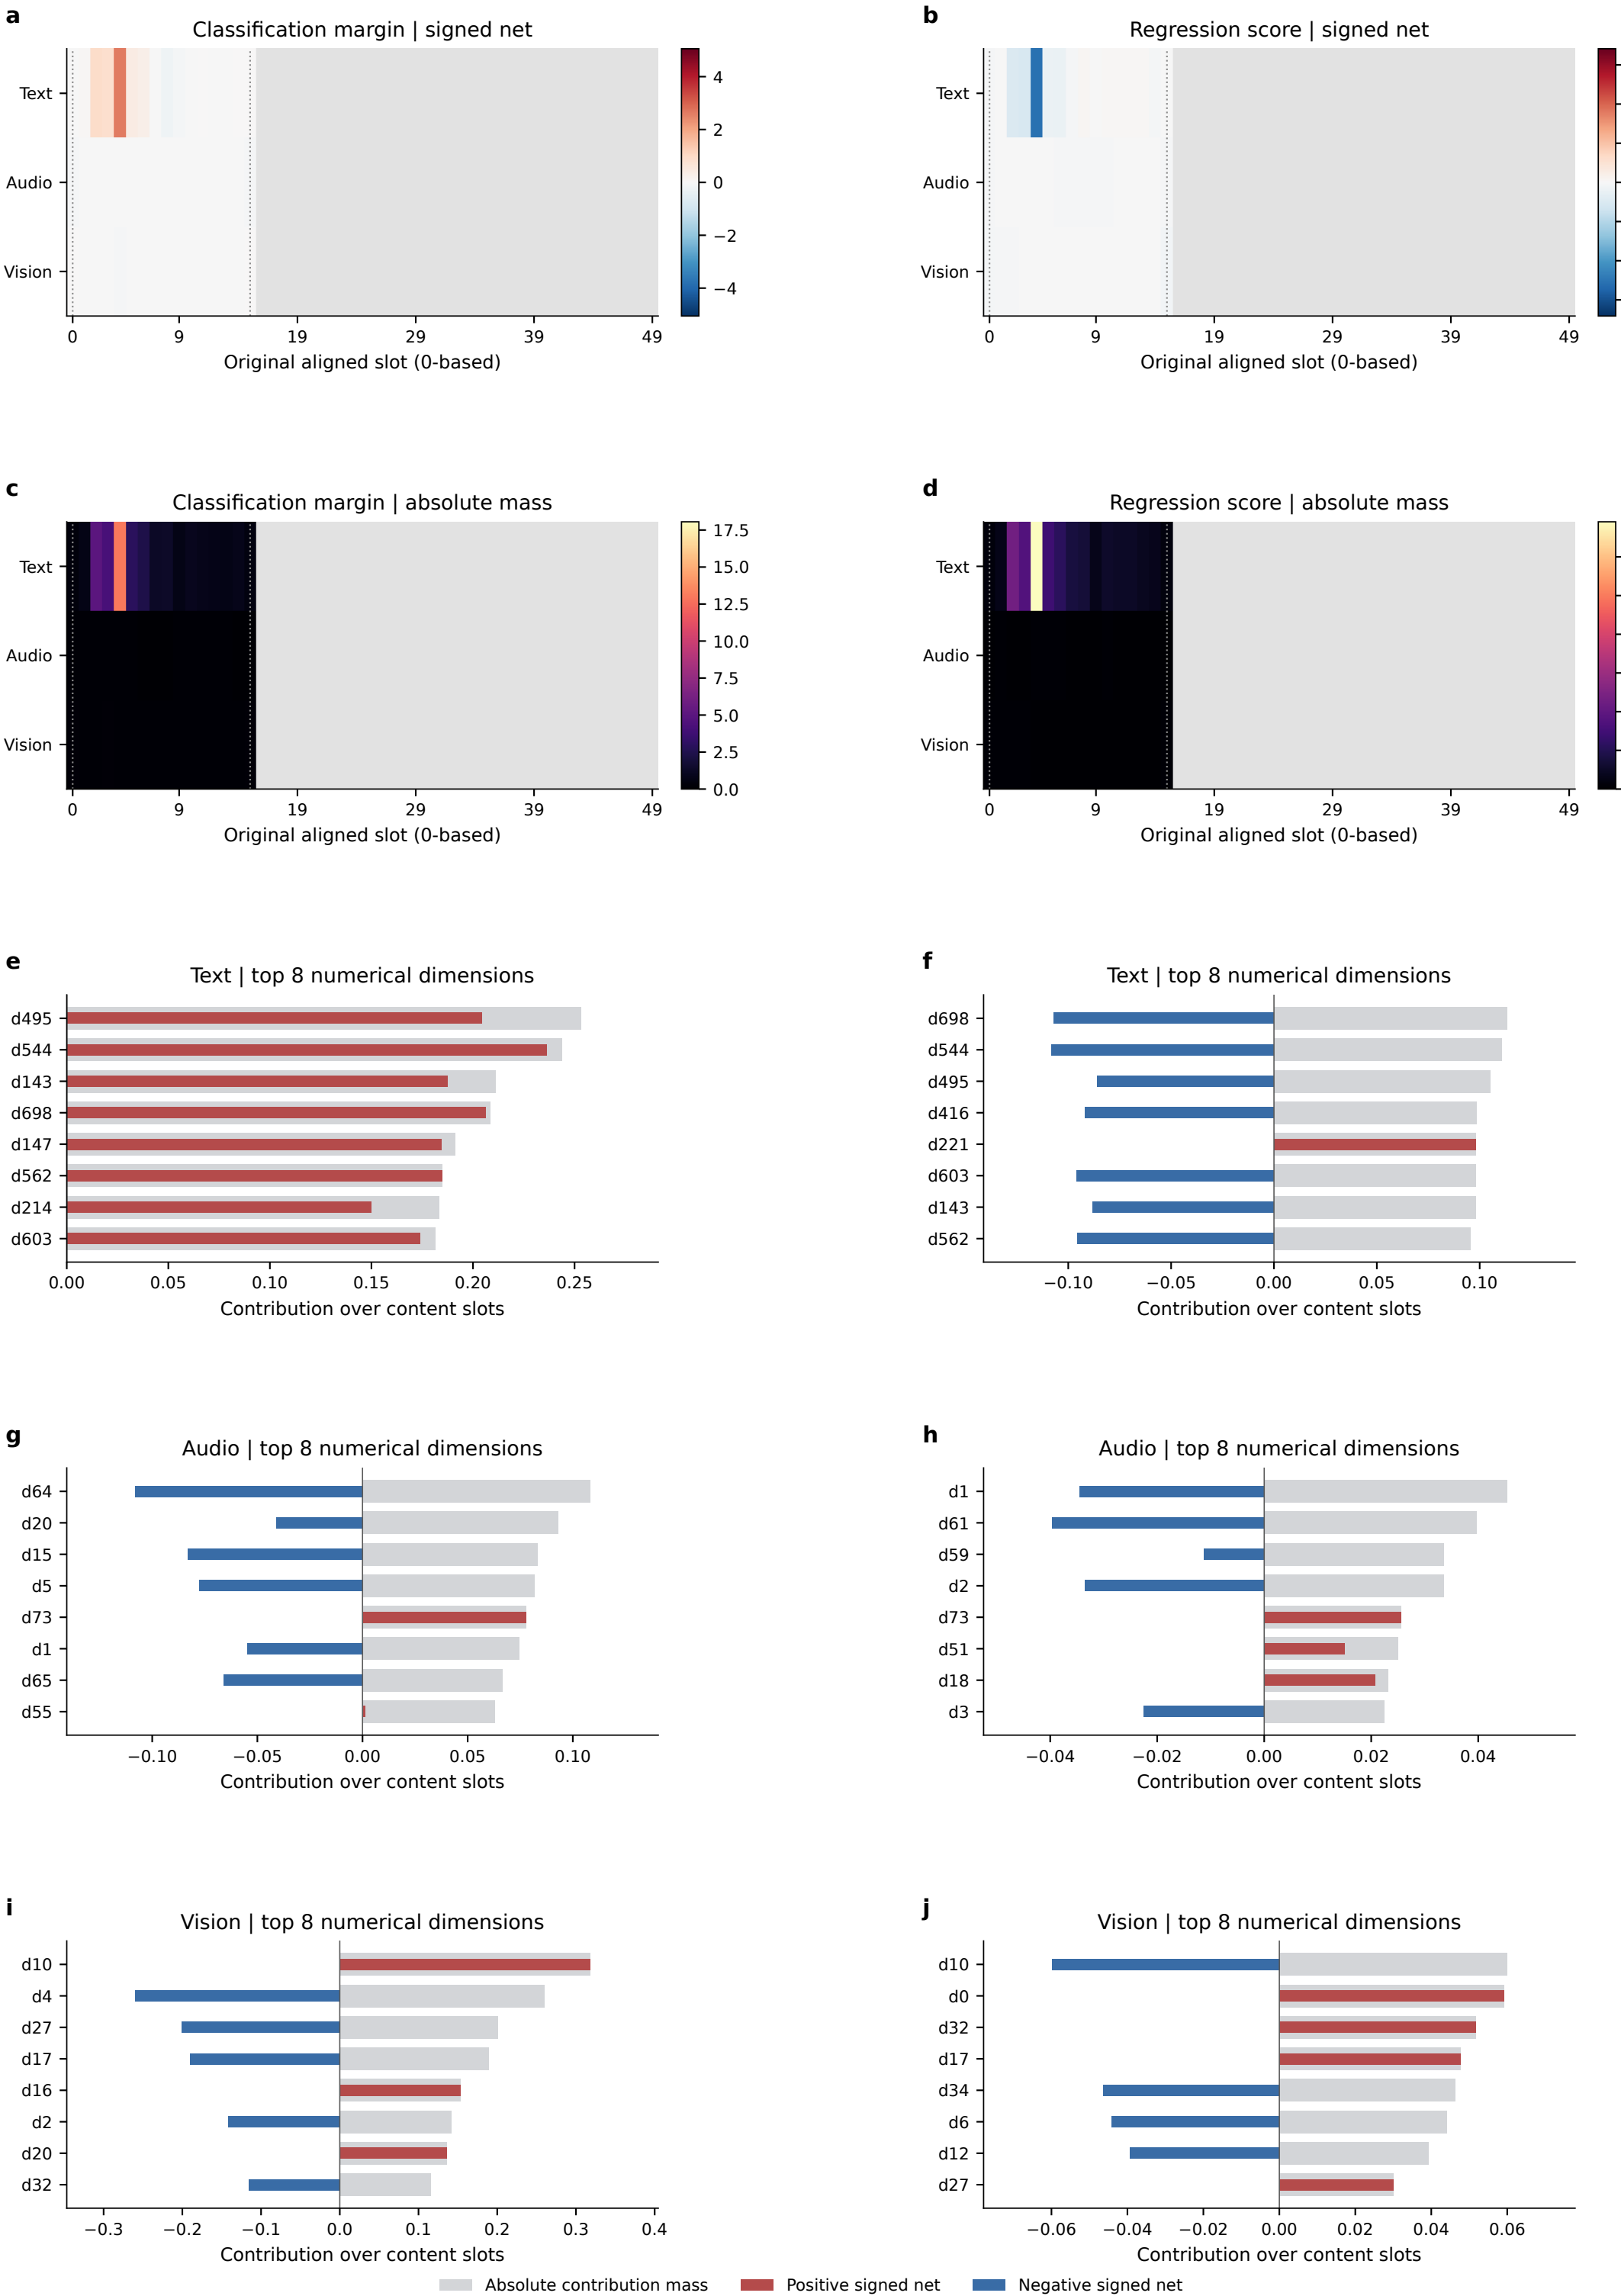
Attachment 4 | Sample 10 | Time-step and feature evidence

Predicted class 0 | Class probabilities [0: 0.892, 1: 0.033, 2: 0.076] | Regression -1.4864



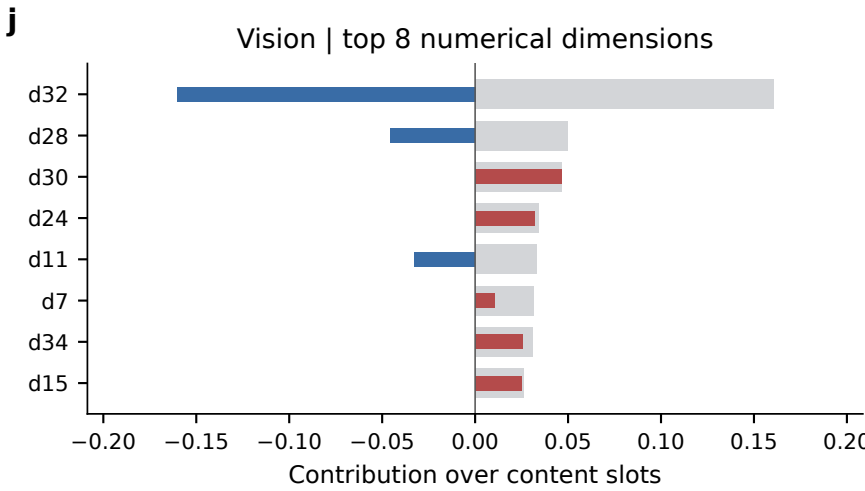
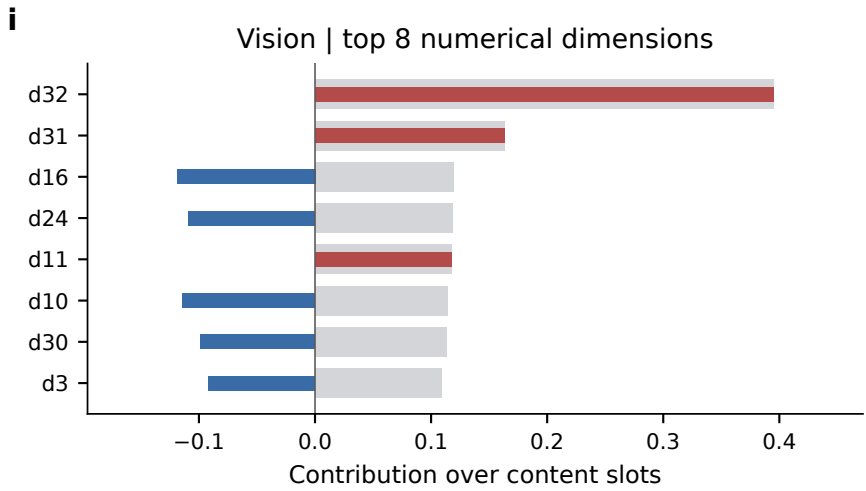
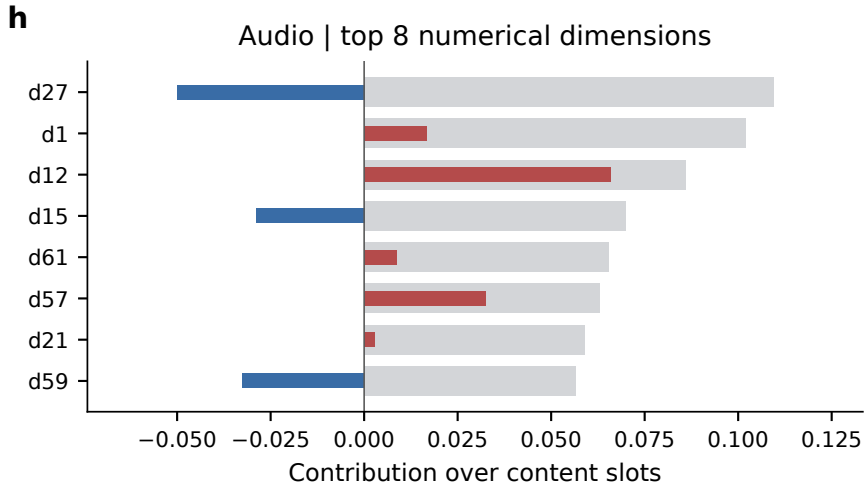
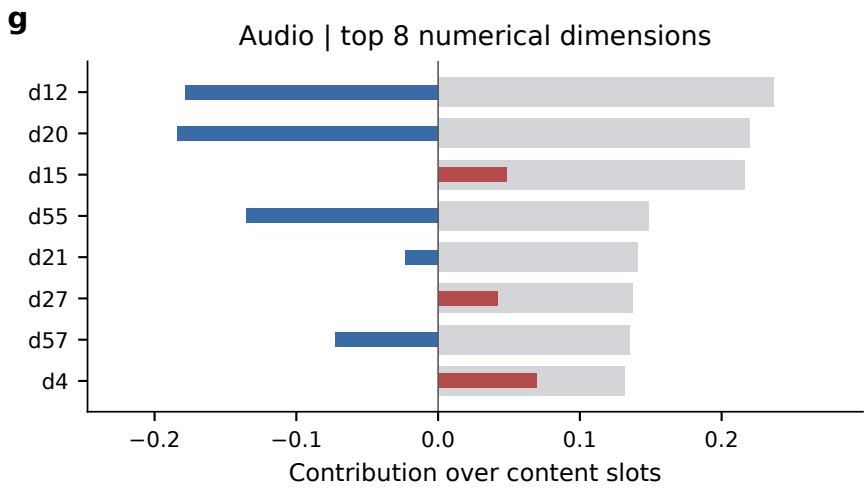
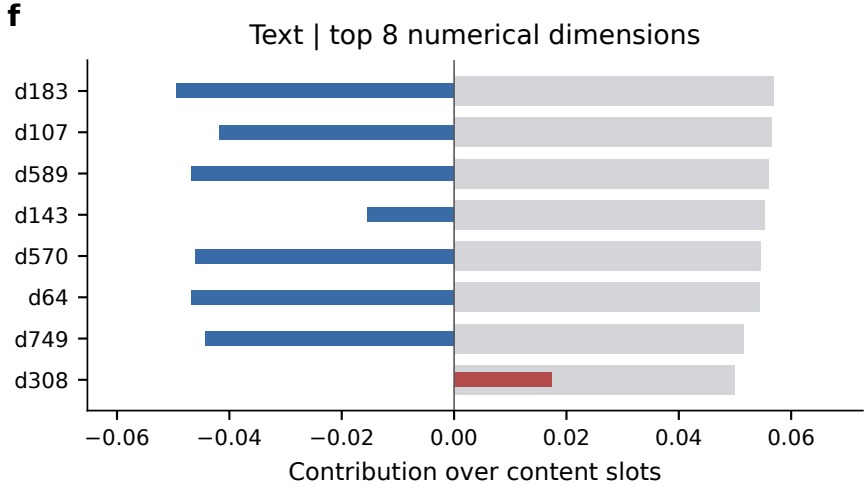
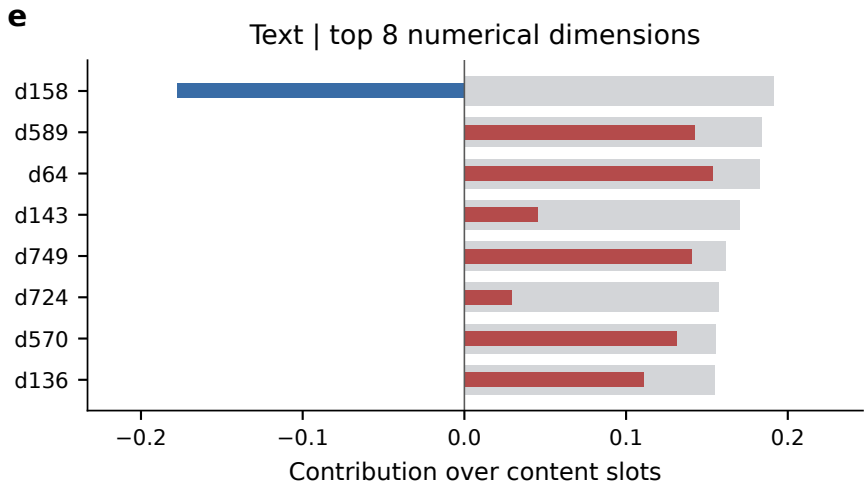
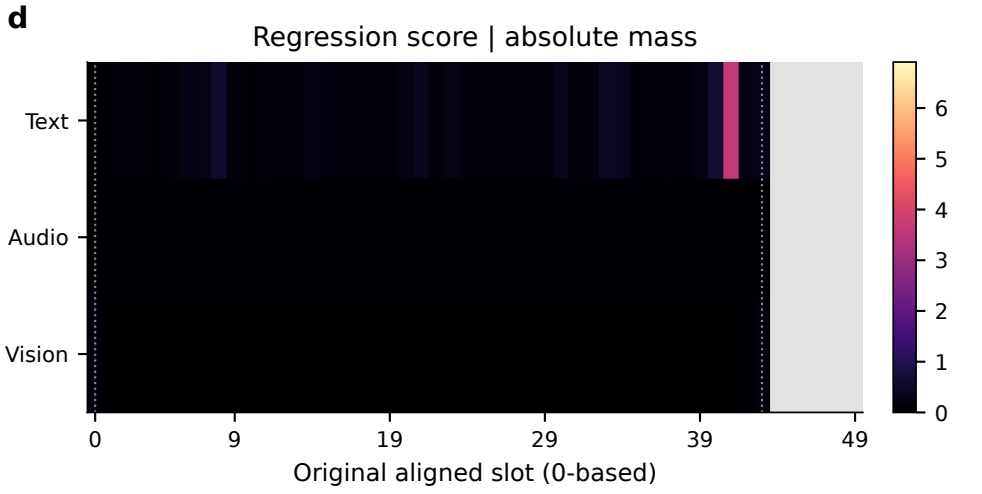
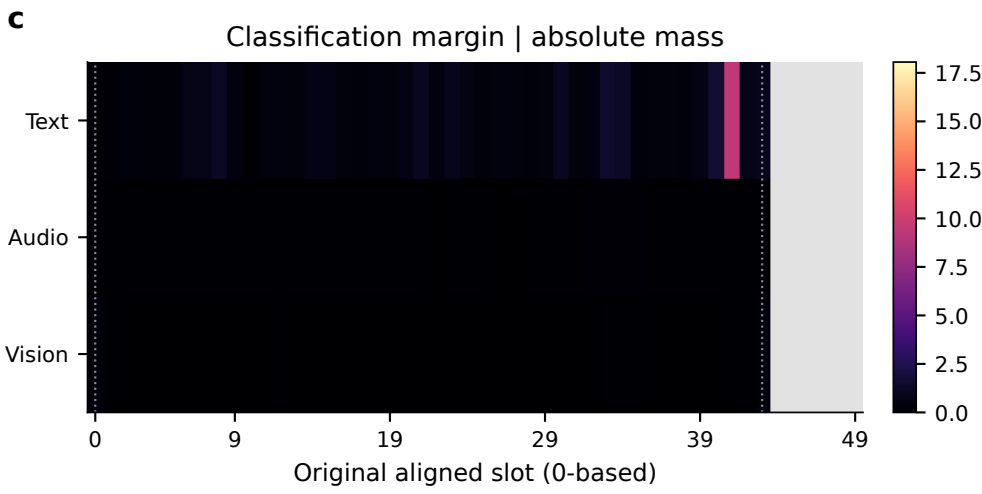
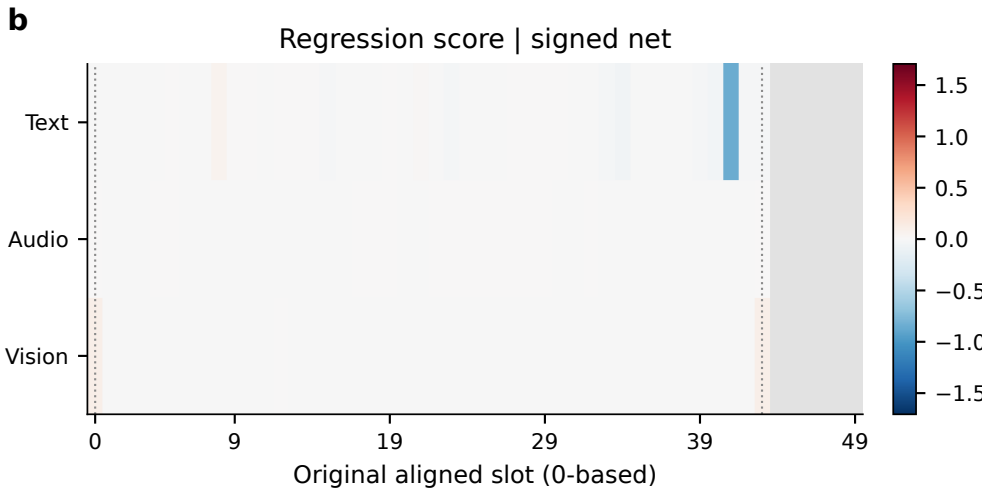
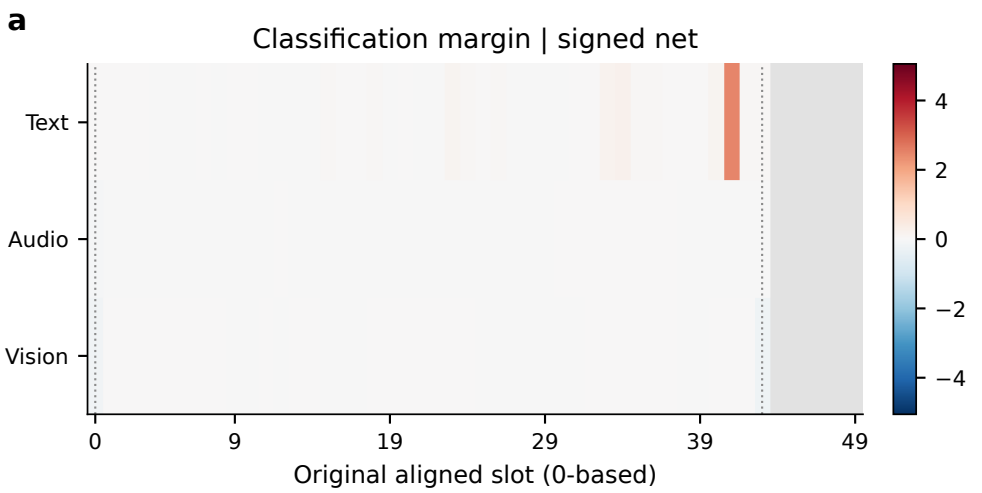
Attachment 4 | Sample 11 | Time-step and feature evidence

Predicted class 0 | Class probabilities [0: 0.912, 1: 0.049, 2: 0.039] | Regression -1.9479



Attachment 4 | Sample 12 | Time-step and feature evidence

Predicted class 0 | Class probabilities [0: 0.790, 1: 0.128, 2: 0.082] | Regression -0.6350



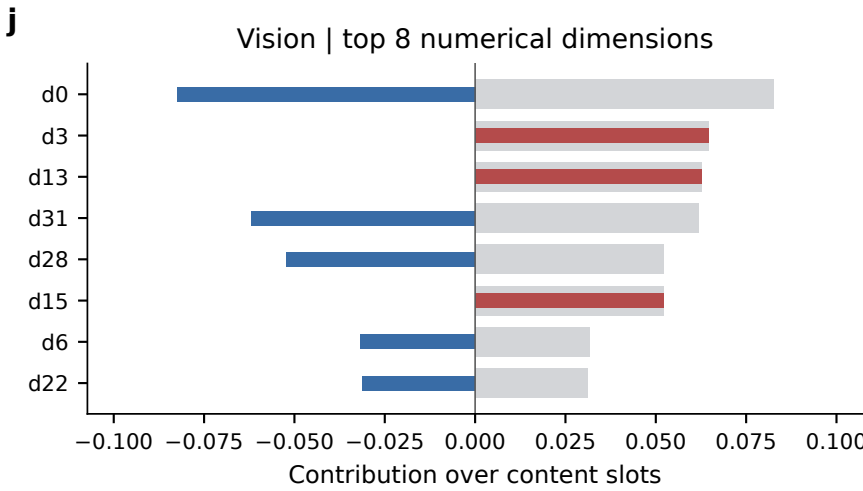
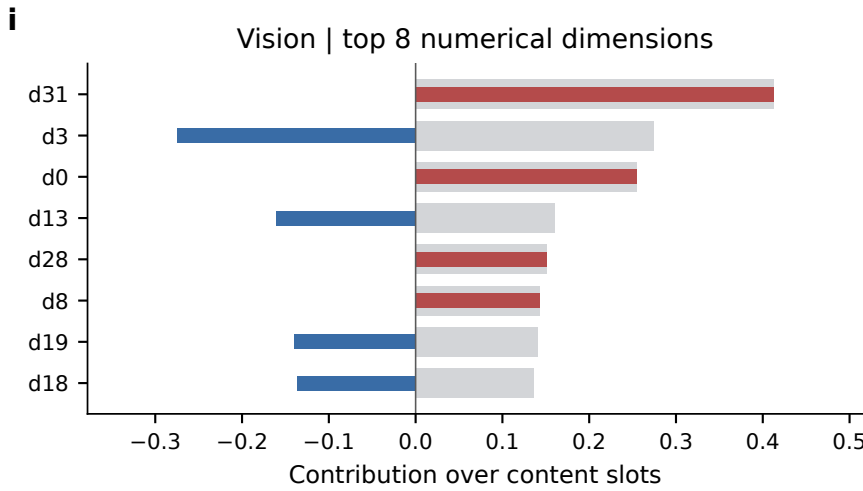
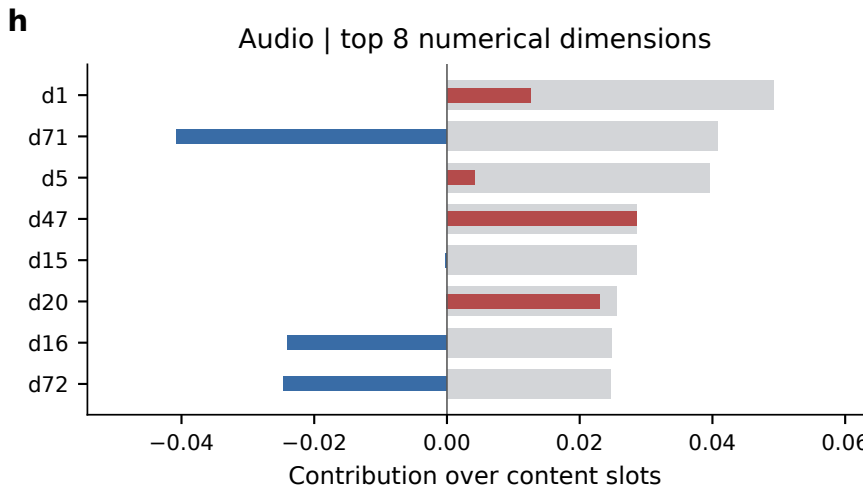
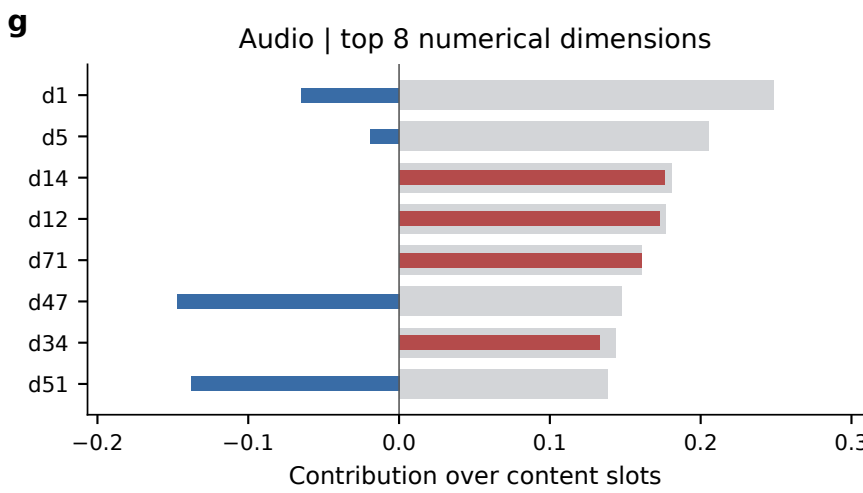
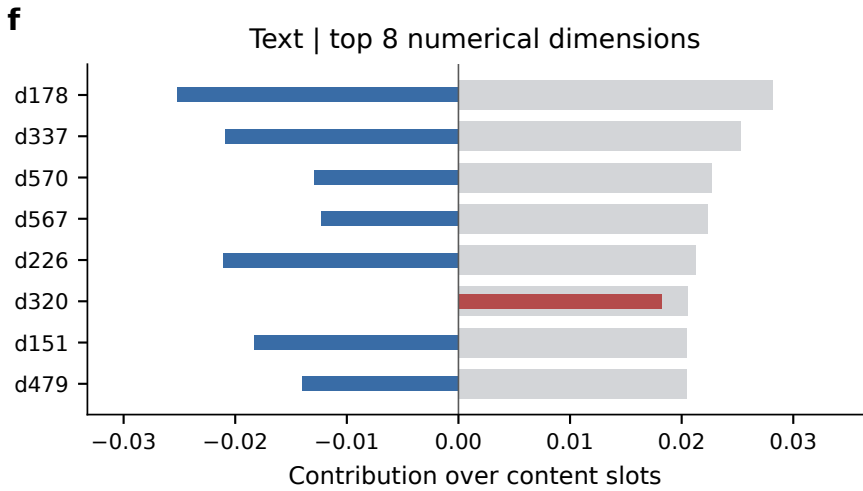
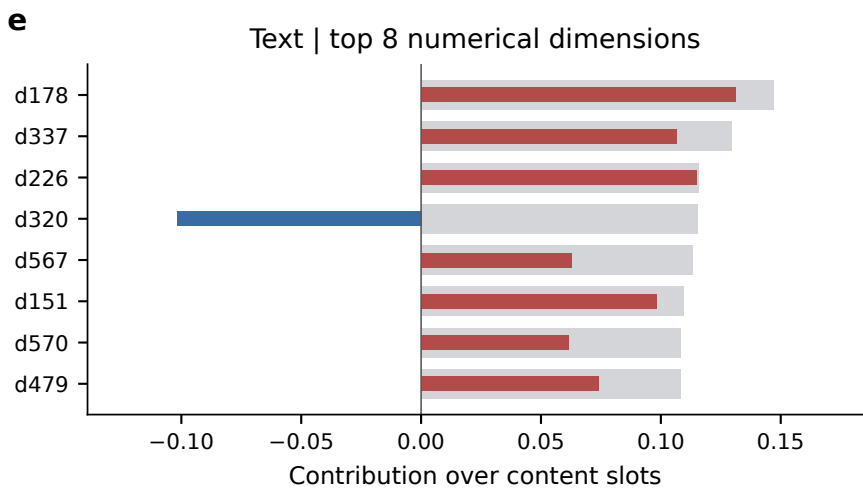
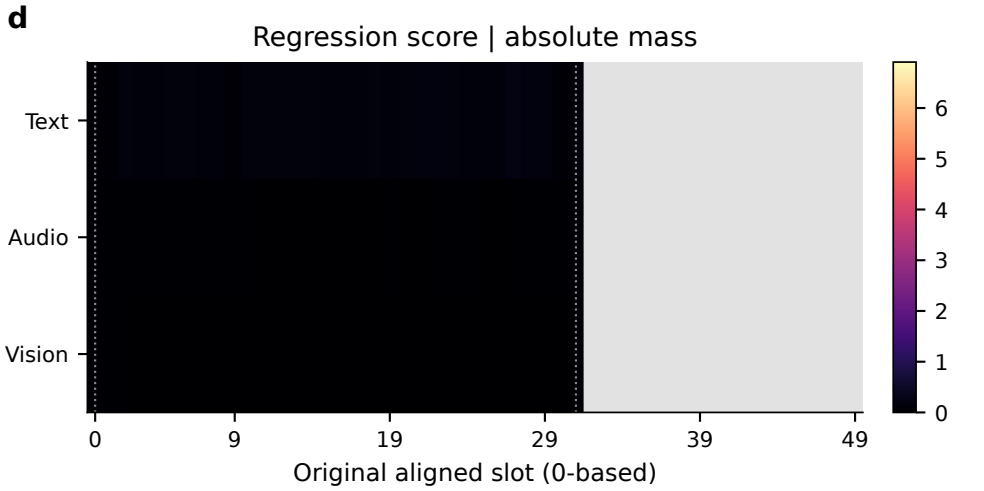
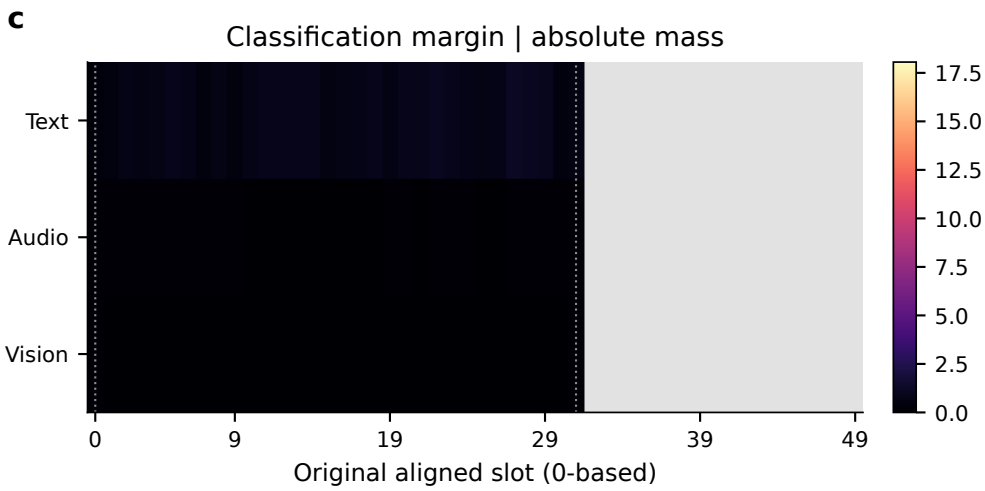
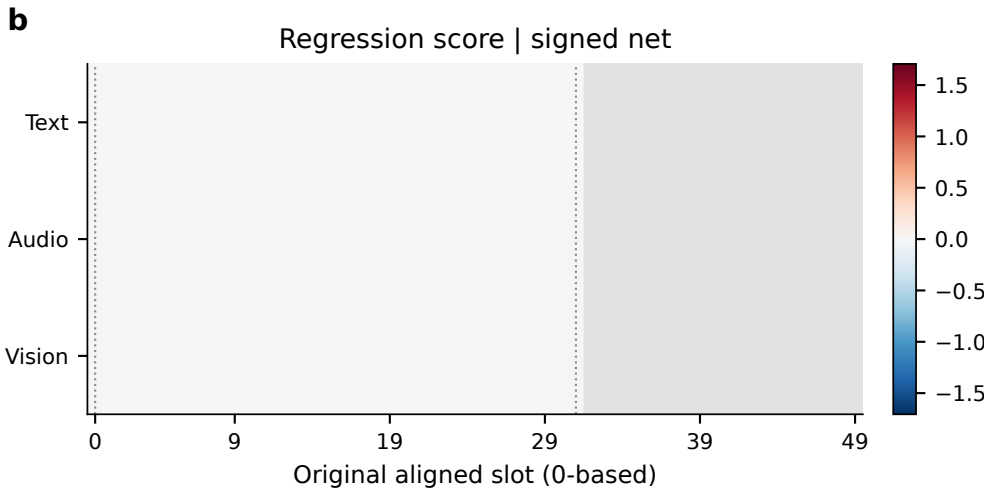
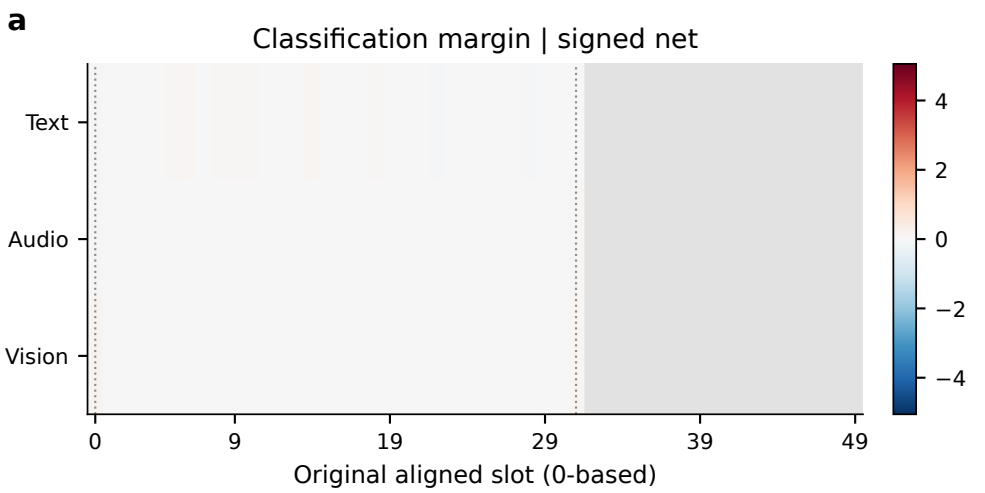
Gray cells: Absolute contribution mass Positive signed net Negative signed net

Train-mean reference. Heatmap scales are shared across all 20 samples for each target. Gray cells: padding; dotted lines: CLS/SEP. Absolute mass is not a modality share. Dimension rankings exclude specials and padding. Dimensions are numerical indices, not semantic labels or video timestamps.

Attachment 4 | Sample 13 | Time-step and feature evidence

Predicted class 1 | Class probabilities [0: 0.087, 1: 0.613, 2: 0.300] | Regression 0.0550

Raw vision is all zero; visual attribution explains zero-input behavior.

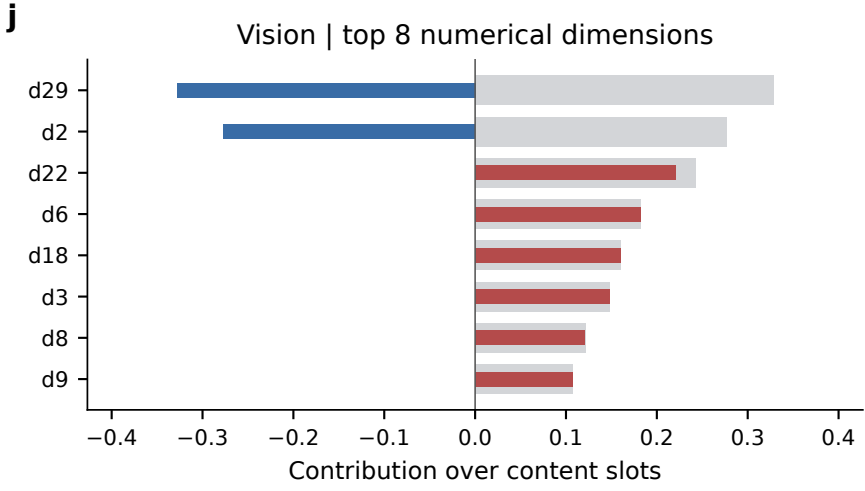
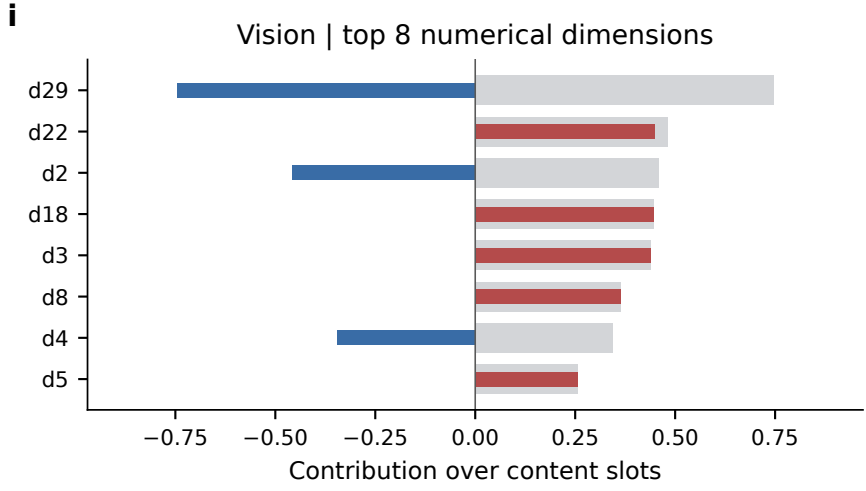
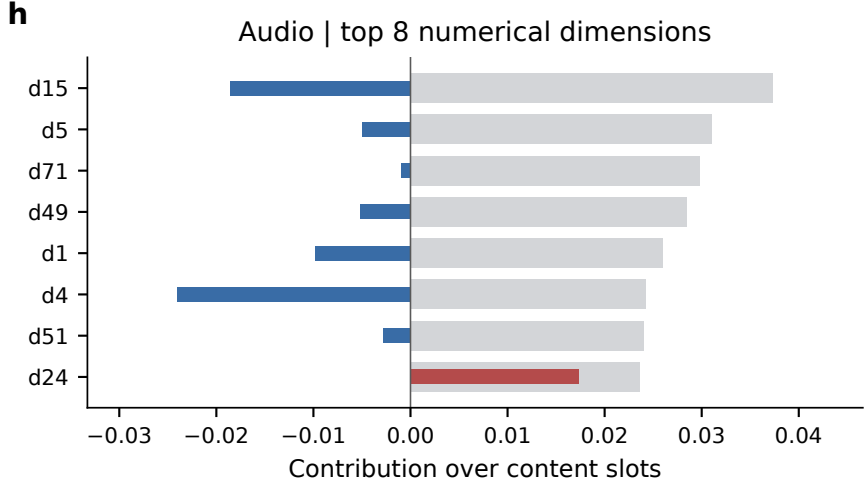
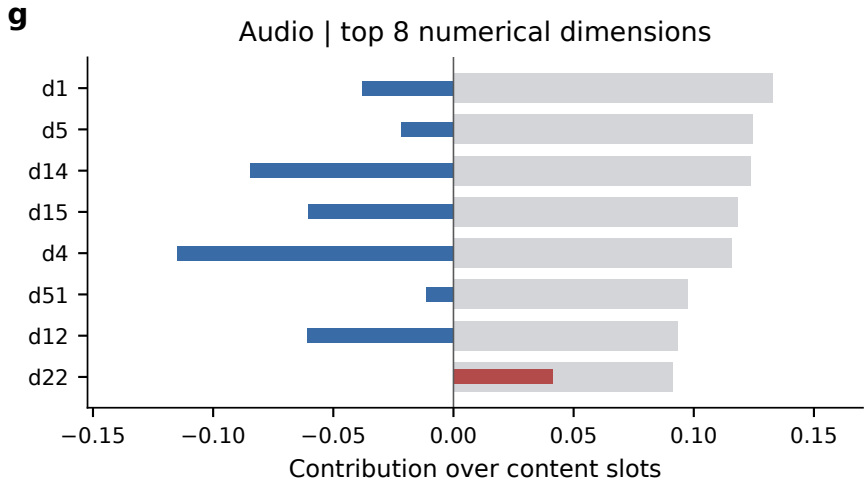
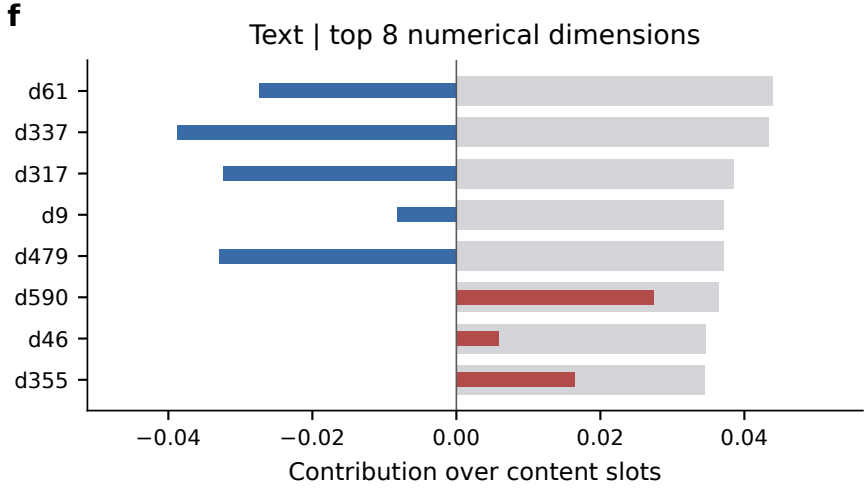
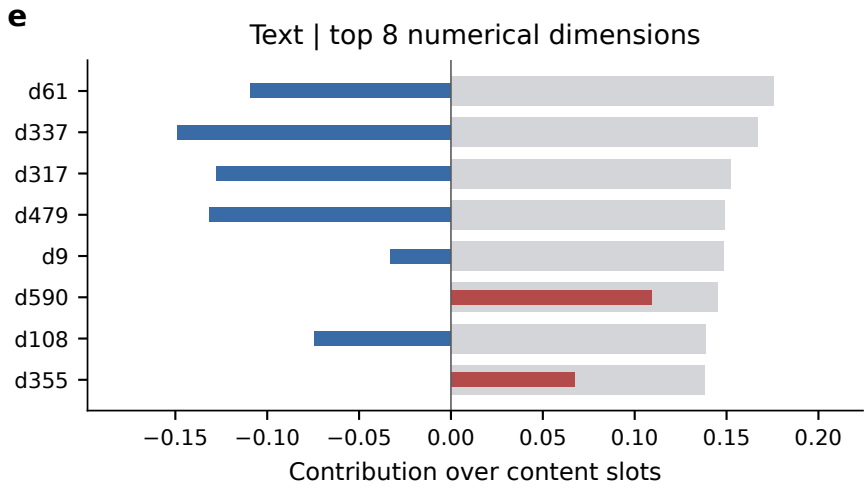
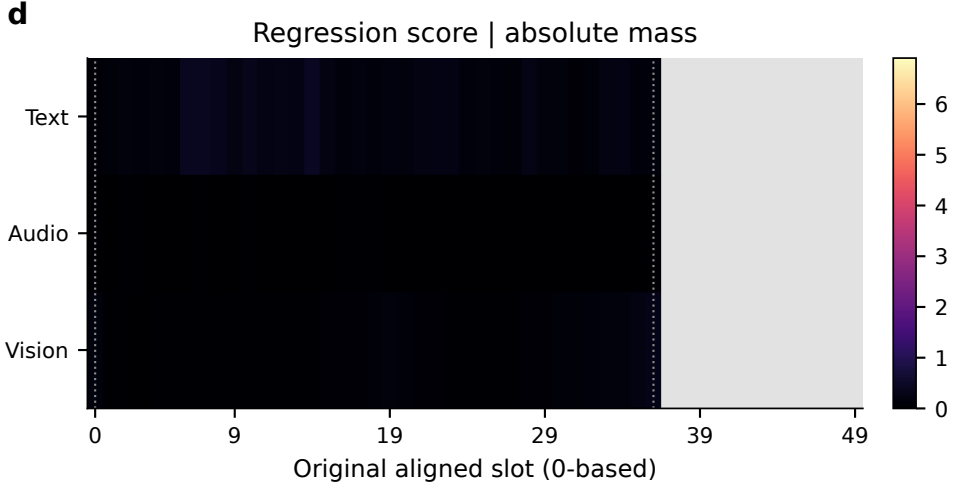
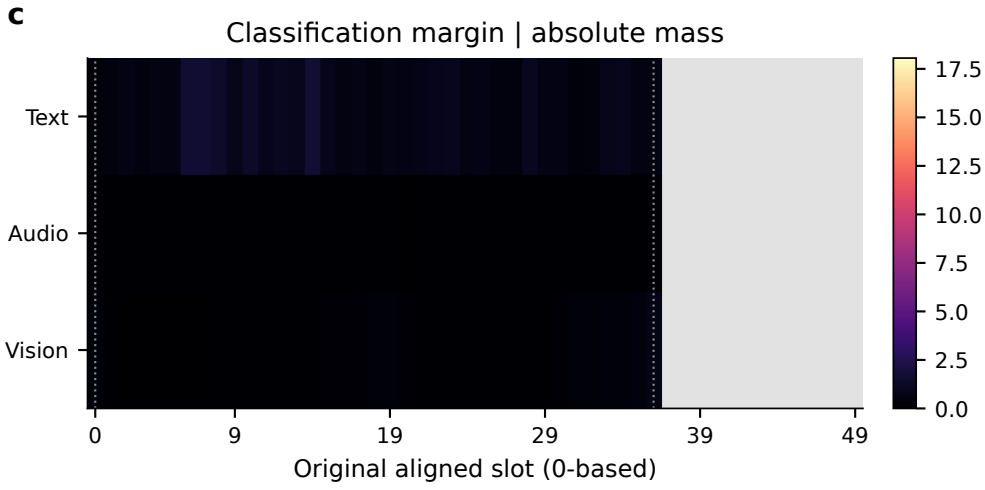
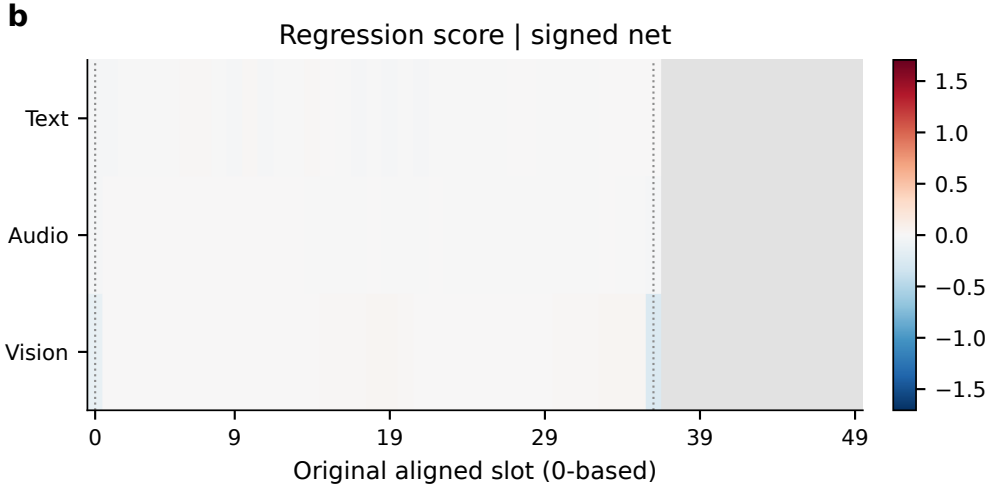
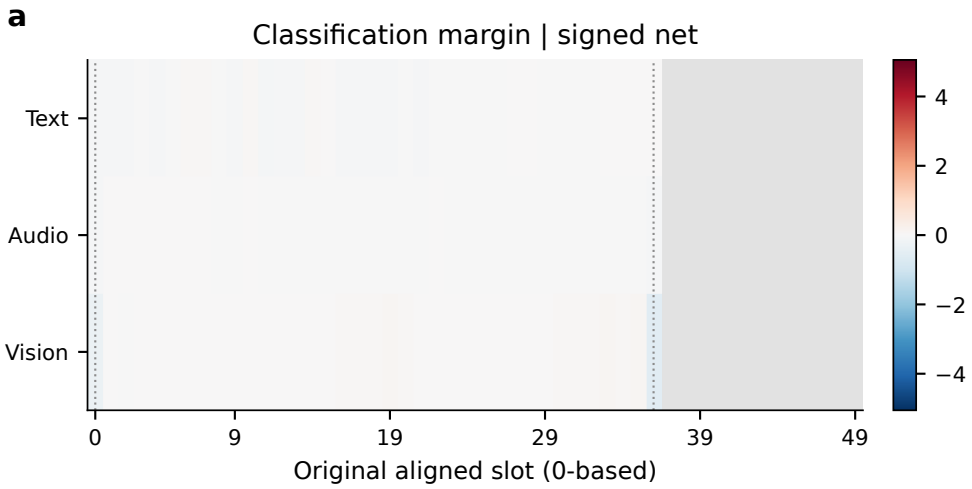


Gray cells: padding; dotted lines: CLS/SEP. Absolute mass

Train-mean reference. Heatmap scales are shared across all 20 samples for each target. Gray cells: padding; dotted lines: CLS/SEP. Absolute mass is not a modality share. Dimension rankings exclude specials and padding. Dimensions are numerical indices, not semantic labels or video timestamps.

Attachment 4 | Sample 14 | Time-step and feature evidence

Predicted class 2 | Class probabilities [0: 0.052, 1: 0.368, 2: 0.579] | Regression 0.3282

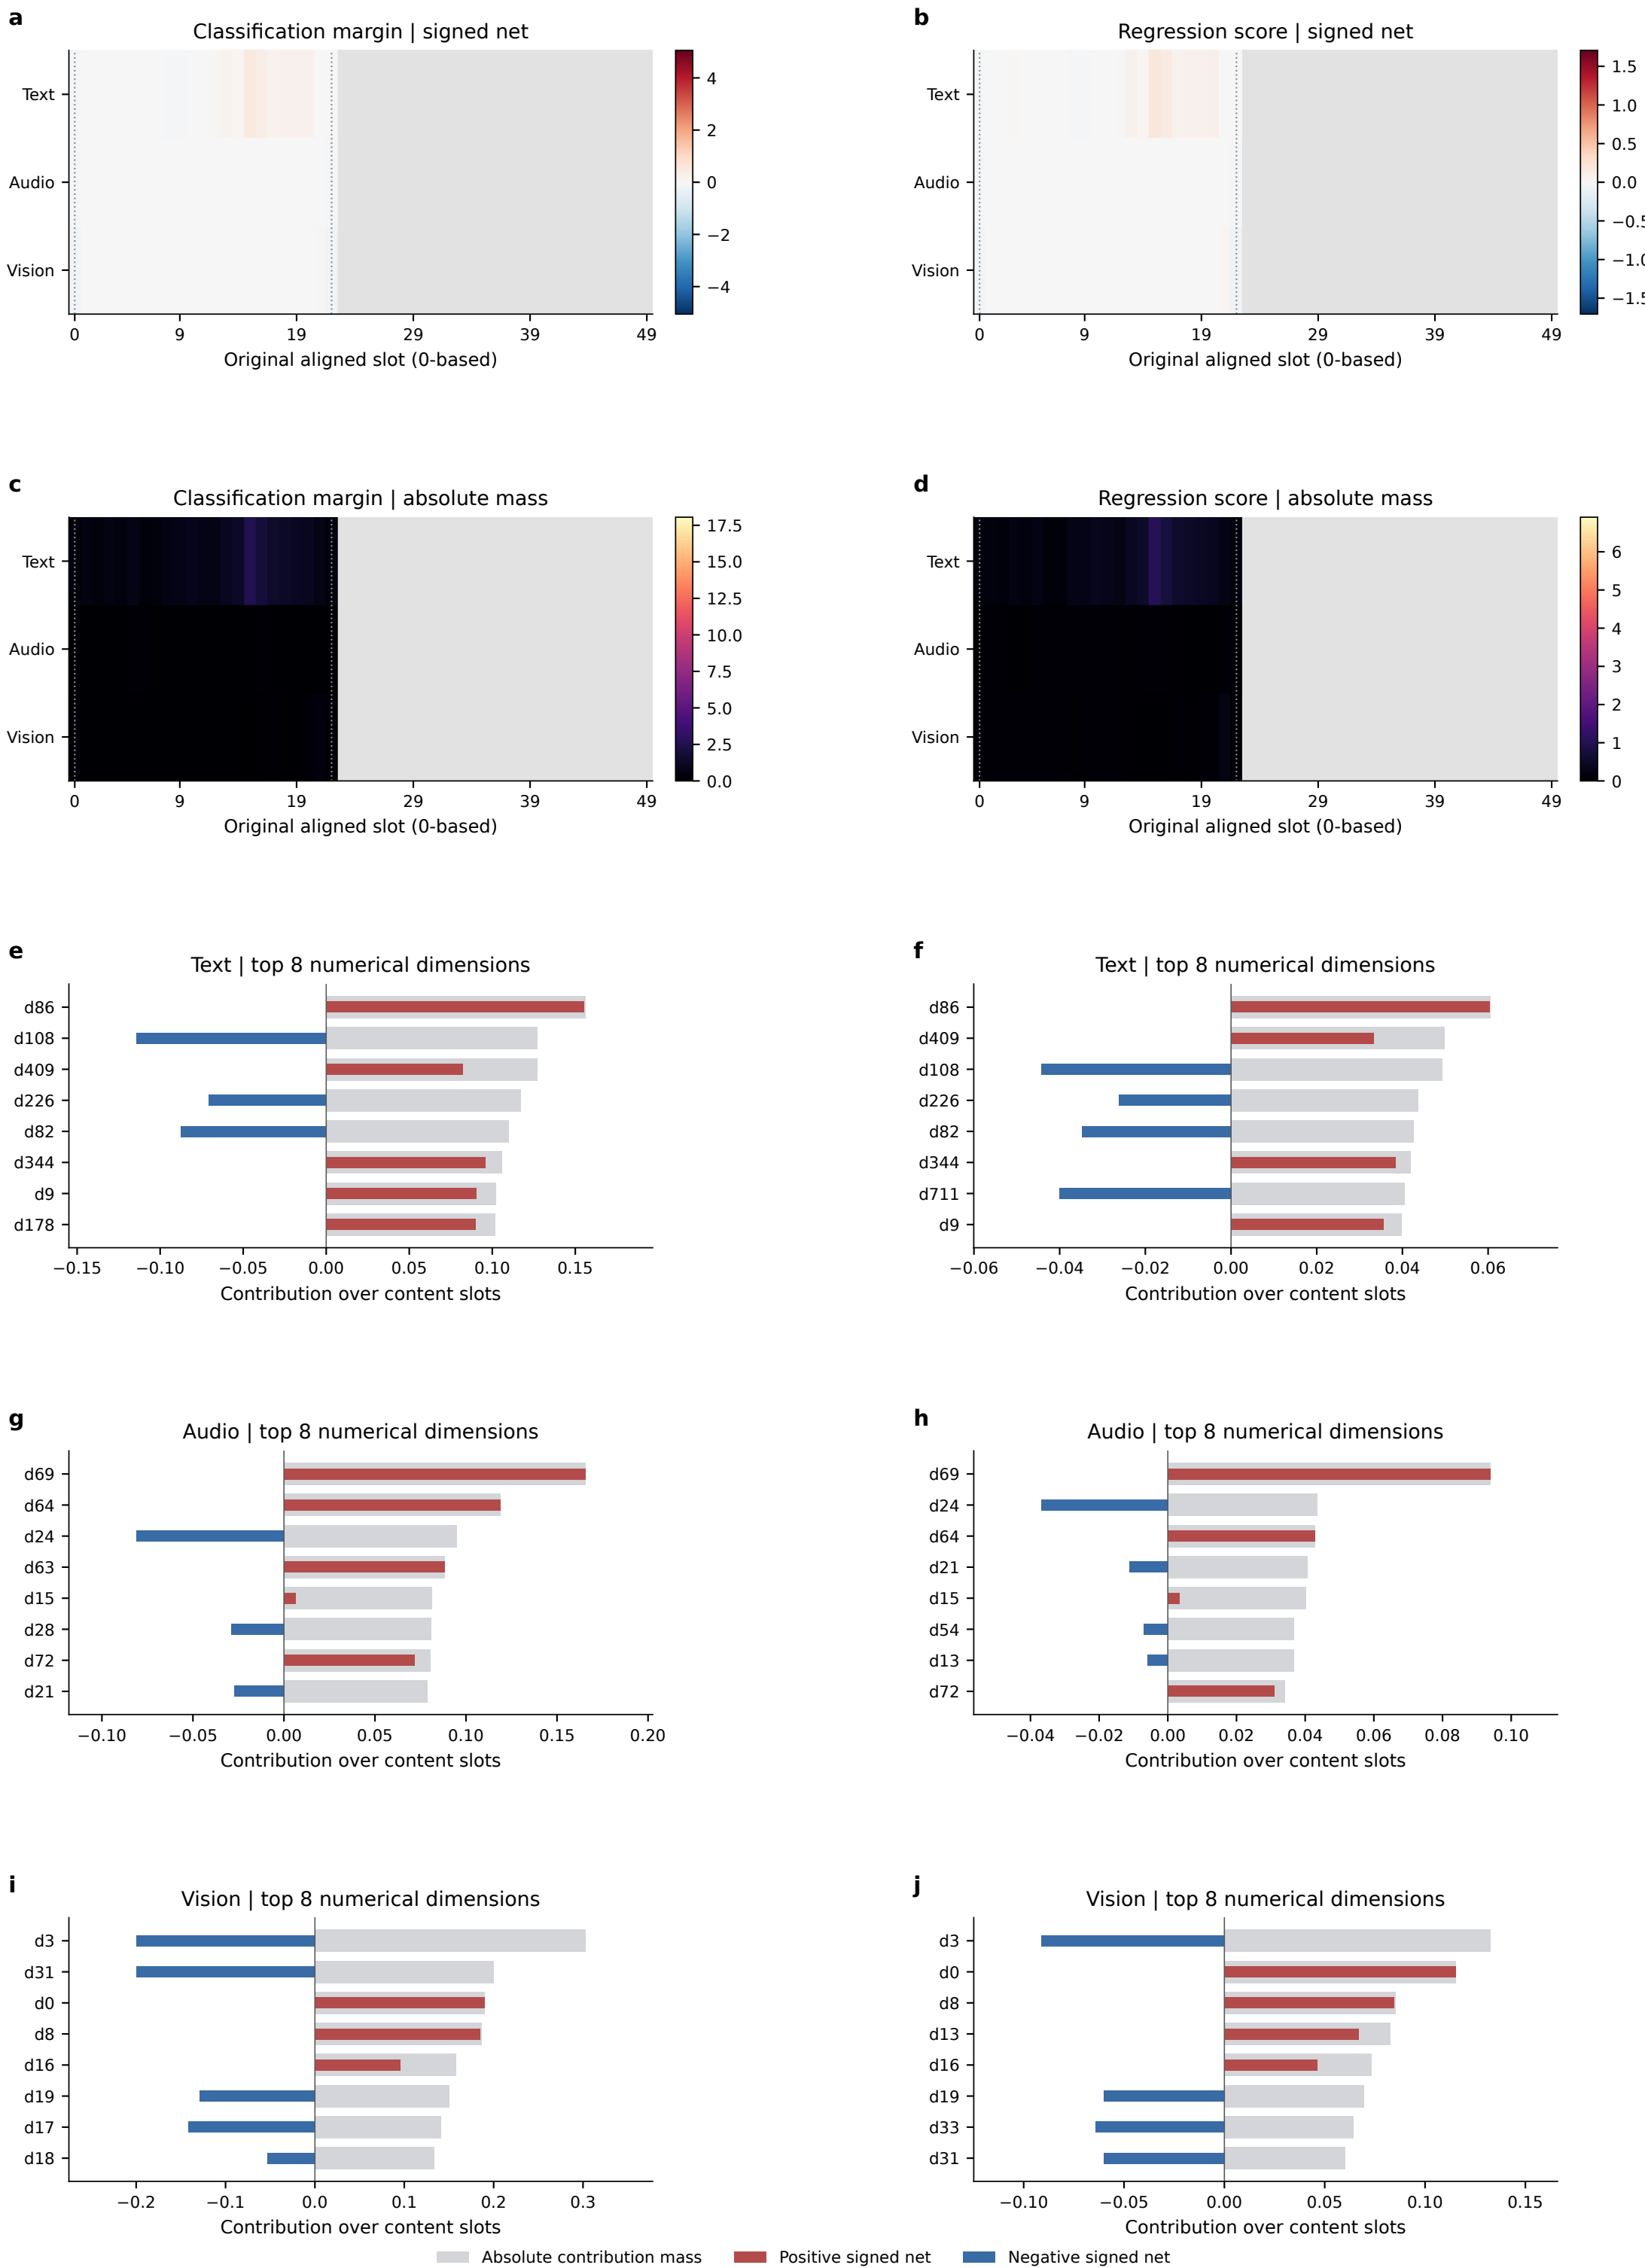


Gray: Absolute contribution mass Red: Positive signed net Blue: Negative signed net

Train-mean reference. Heatmap scales are shared across all 20 samples for each target. Gray cells: padding; dotted lines: CLS/SEP. Absolute mass is not a modality share. Dimension rankings exclude specials and padding. Dimensions are numerical indices, not semantic labels or video timestamps.

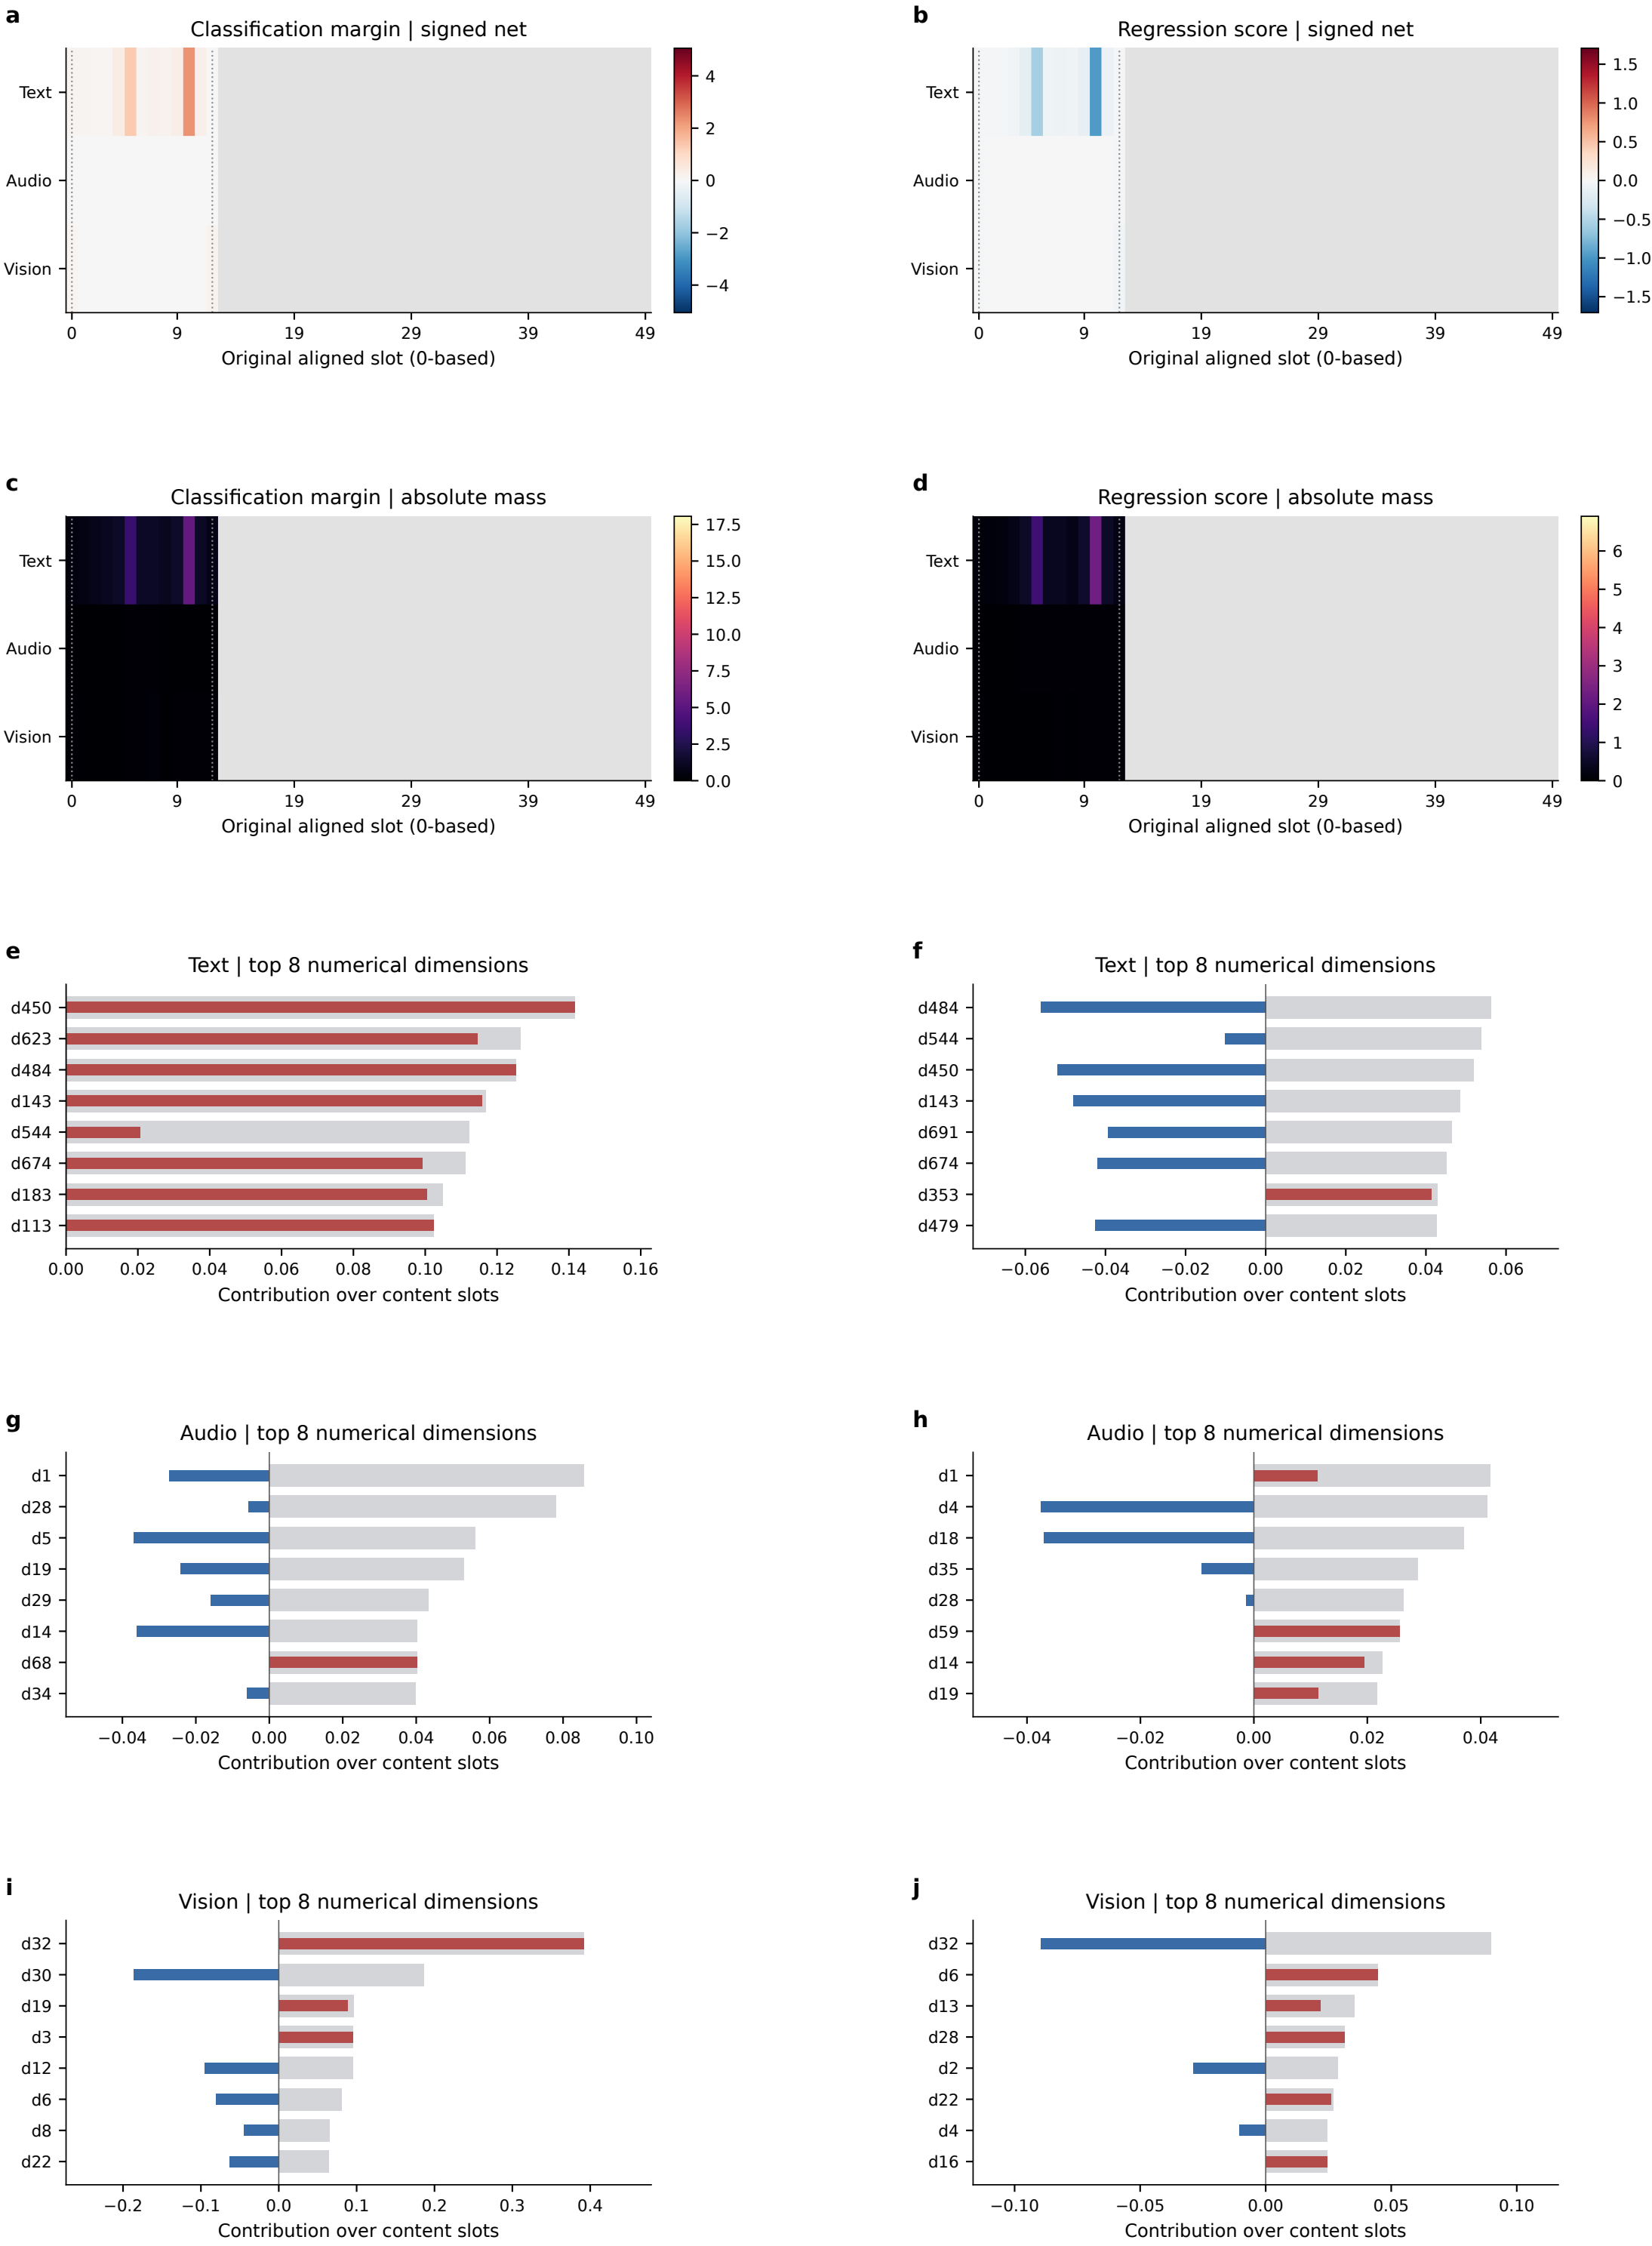
Attachment 4 | Sample 15 | Time-step and feature evidence

Predicted class 2 | Class probabilities [0: 0.022, 1: 0.089, 2: 0.889] | Regression 1.0632



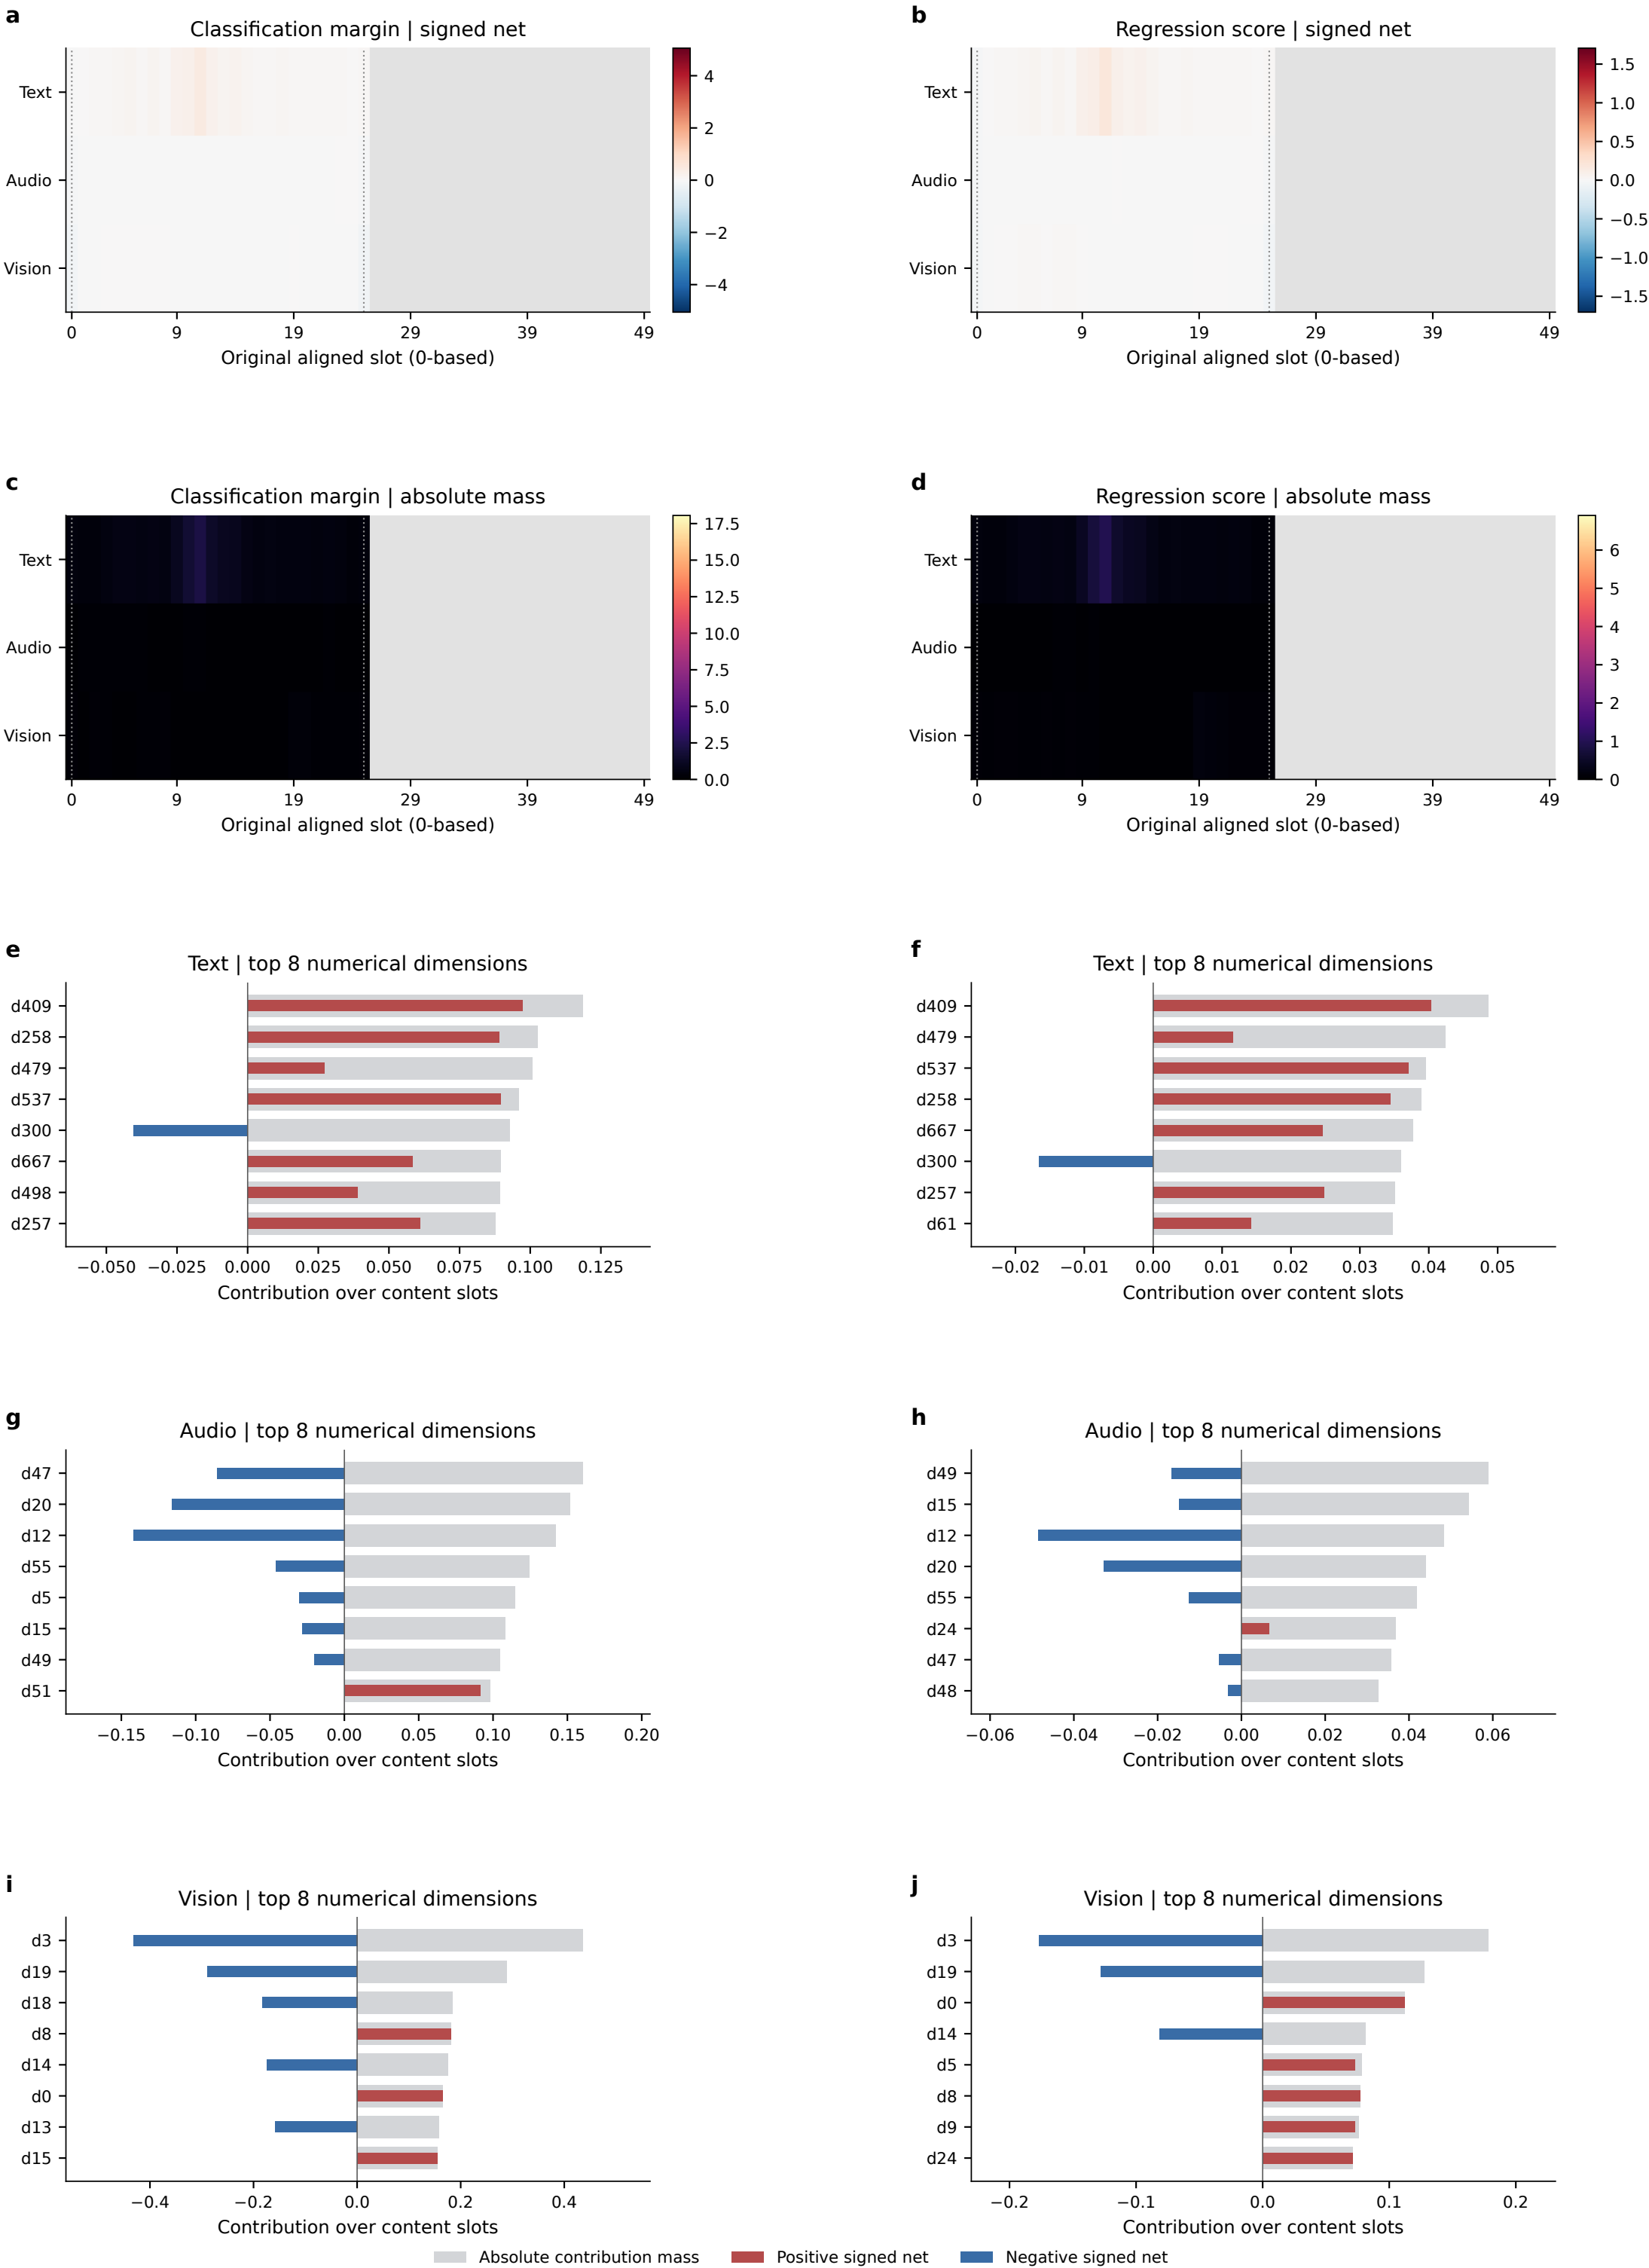
Attachment 4 | Sample 16 | Time-step and feature evidence

Predicted class 0 | Class probabilities [0: 0.949, 1: 0.023, 2: 0.028] | Regression -2.0530



Attachment 4 | Sample 17 | Time-step and feature evidence

Predicted class 2 | Class probabilities [0: 0.015, 1: 0.041, 2: 0.944] | Regression 1.3862



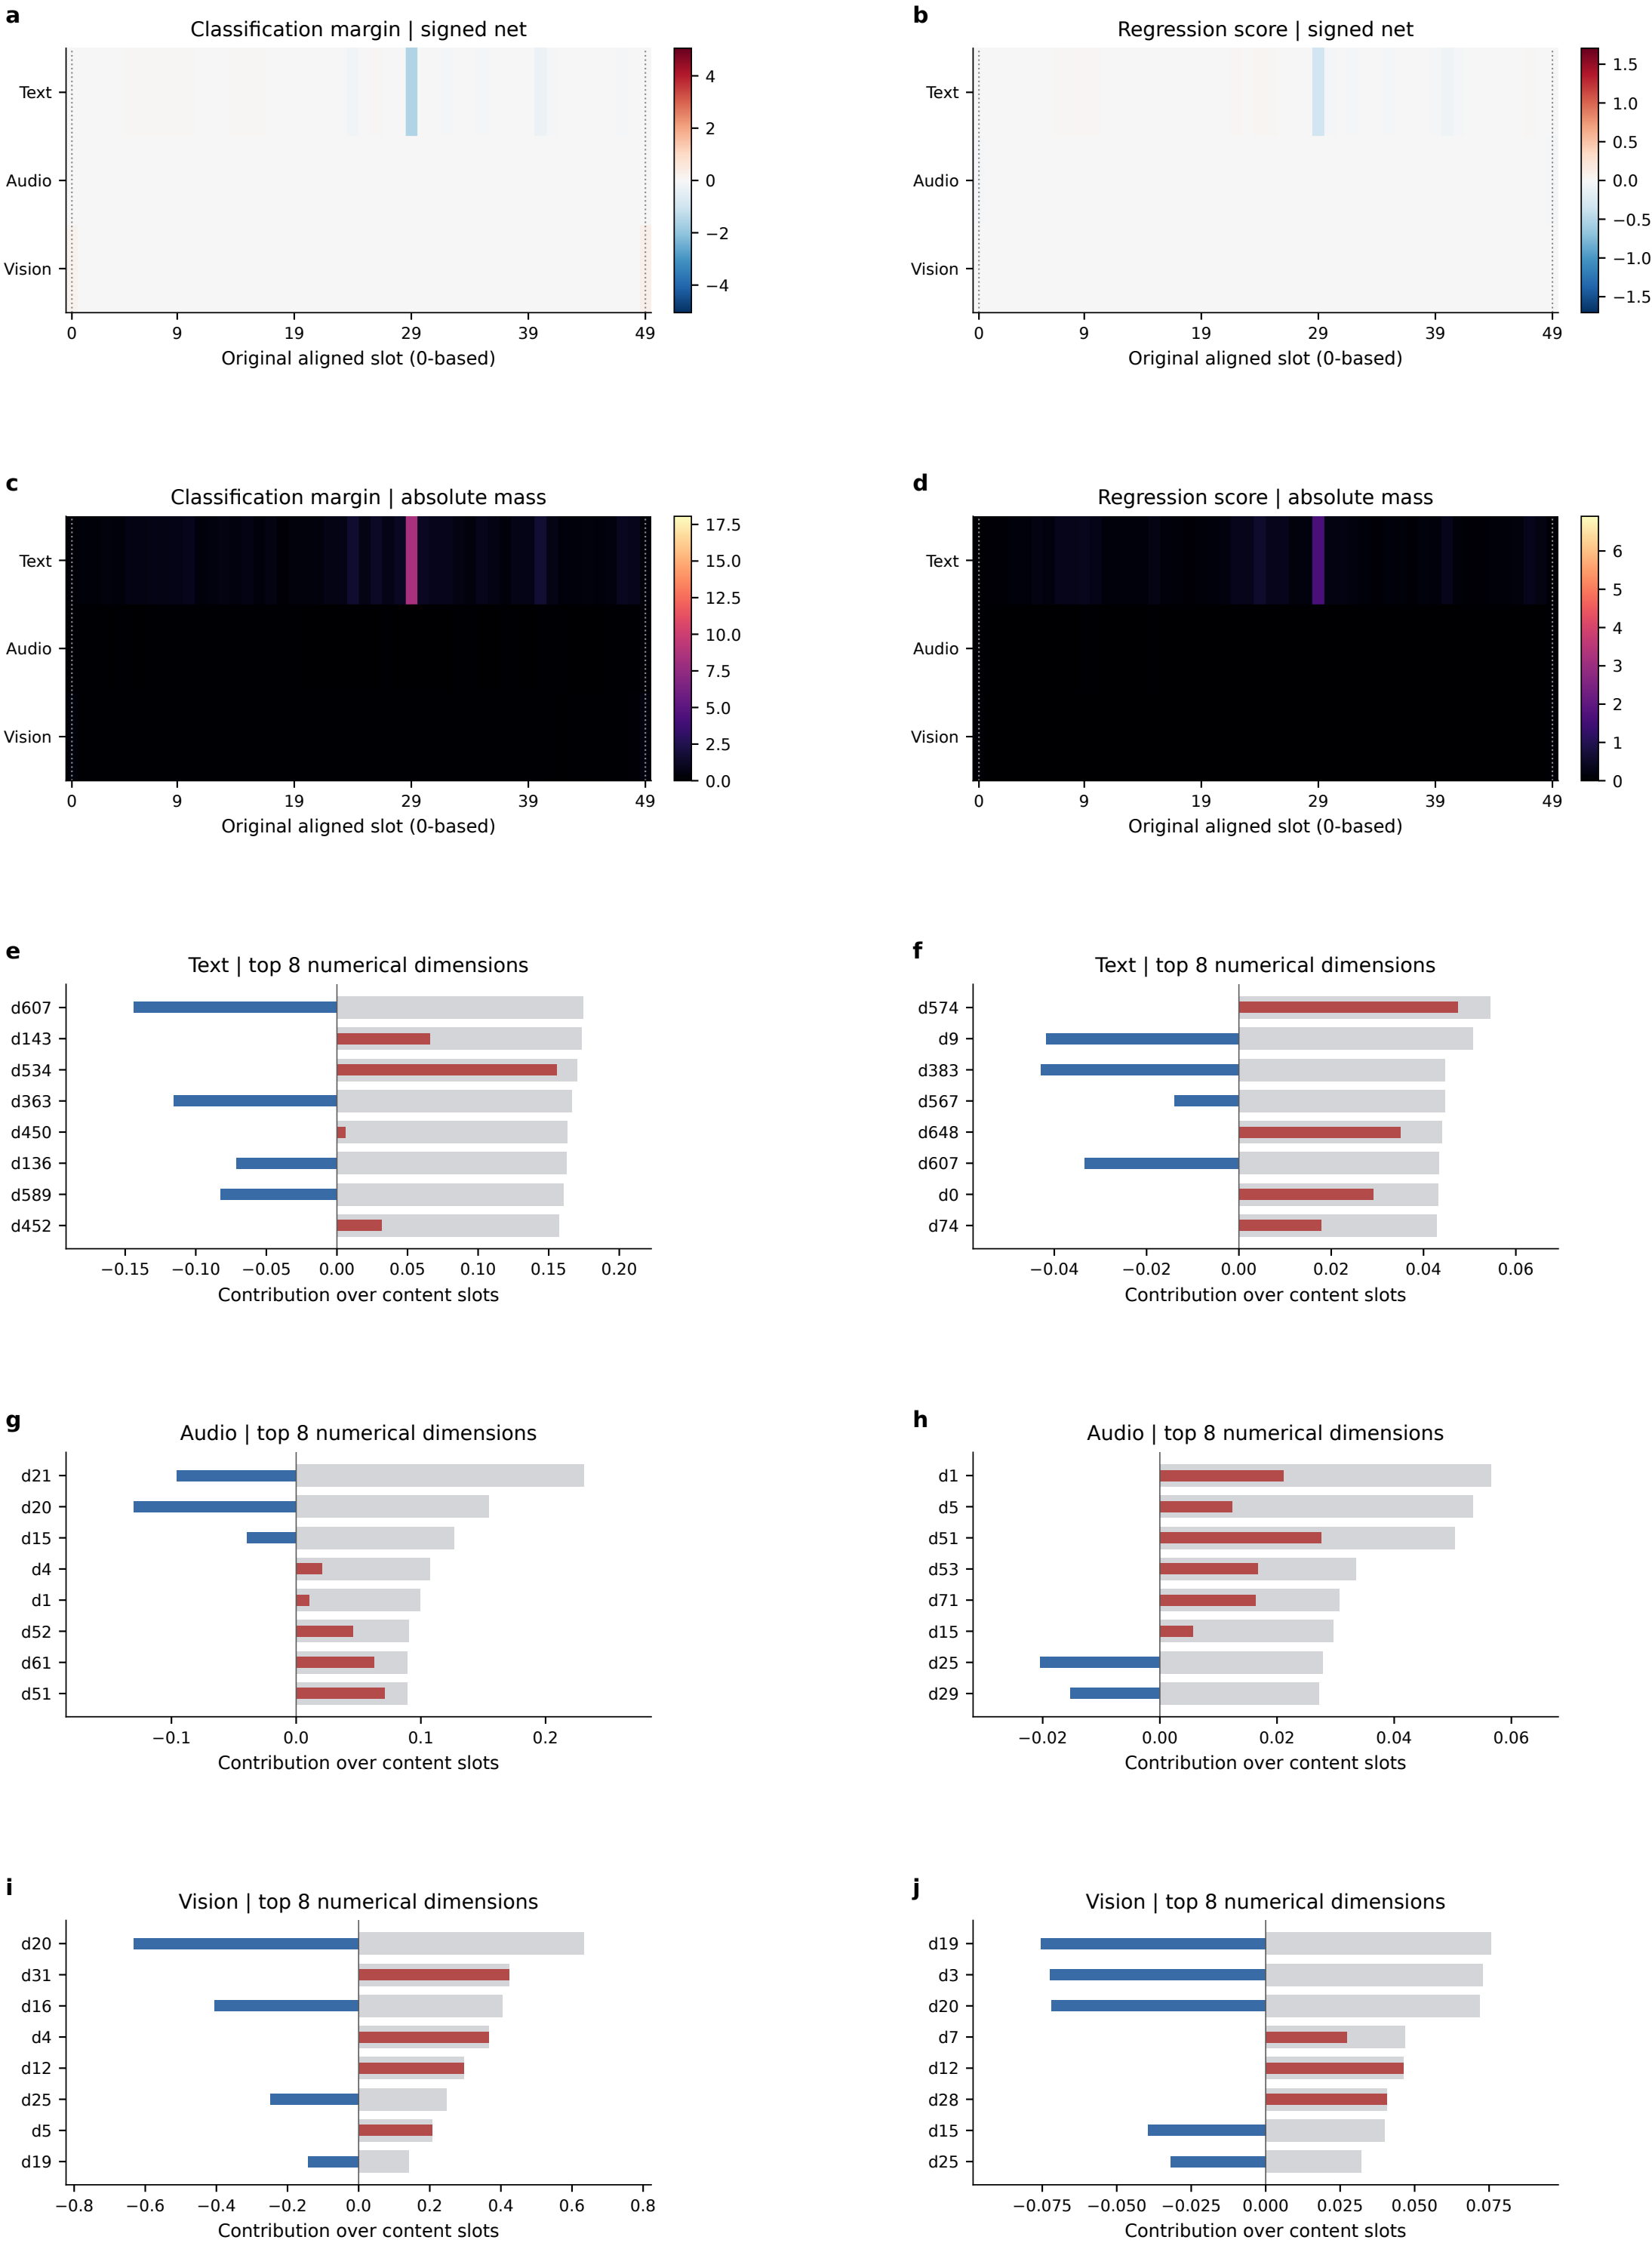
Absolute contribution mass Positive signed net Negative signed net

Train-mean reference. Heatmap scales are shared across all 20 samples for each target. Gray cells: padding; dotted lines: CLS/SEP. Absolute mass is not a modality share. Dimension rankings exclude specials and padding. Dimensions are numerical indices, not semantic labels or video timestamps.

Attachment 4 | Sample 18 | Time-step and feature evidence

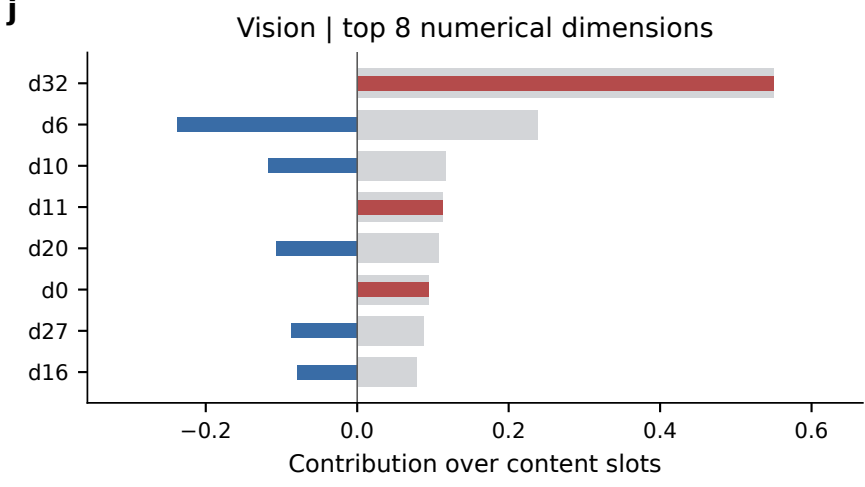
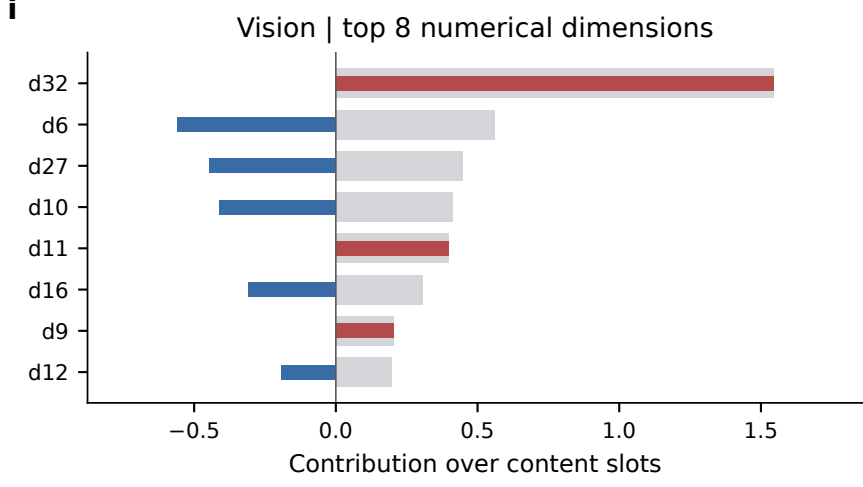
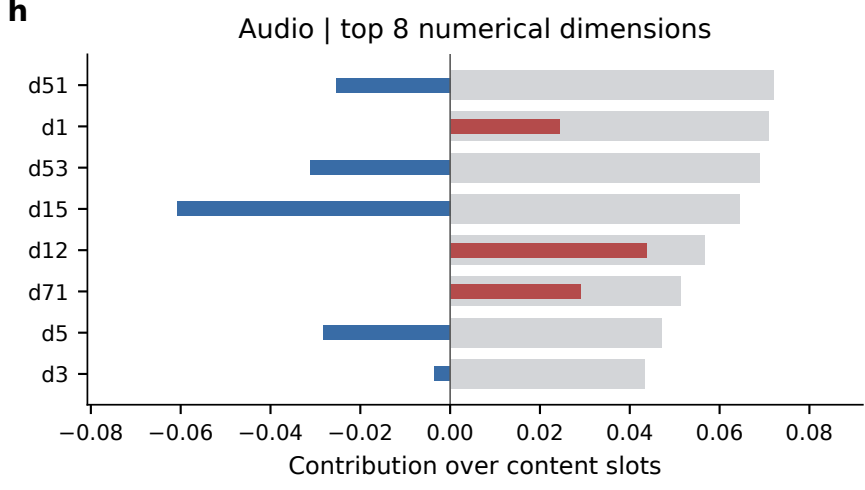
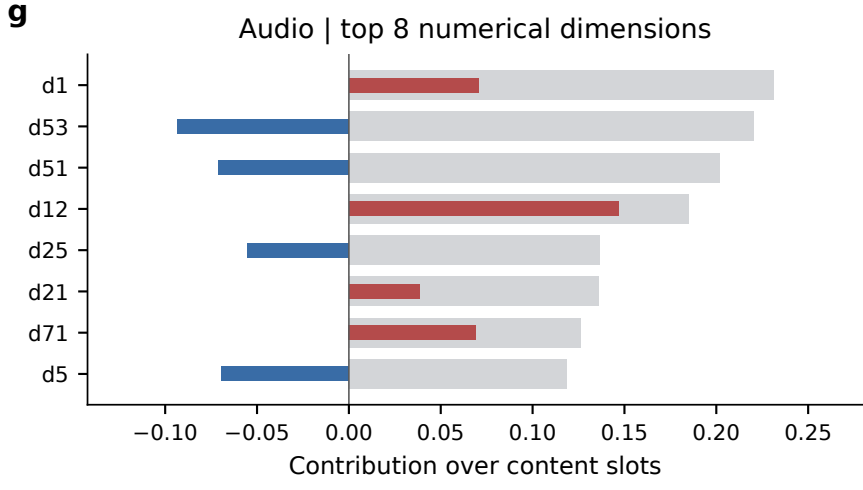
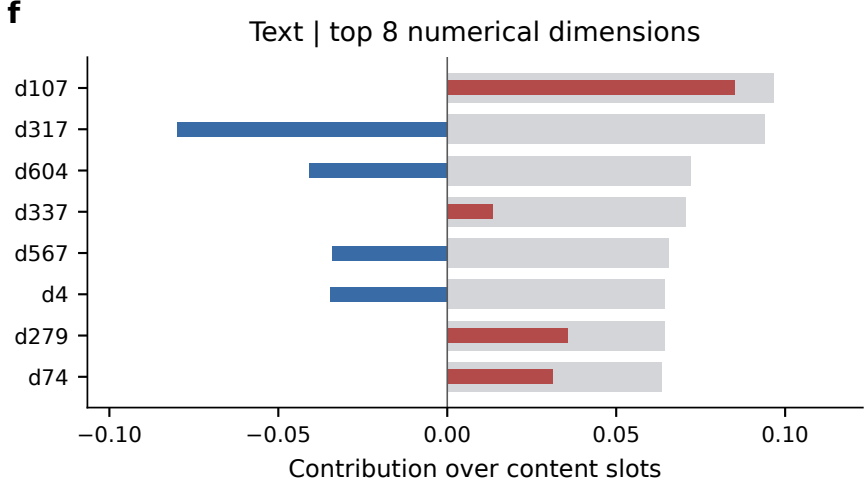
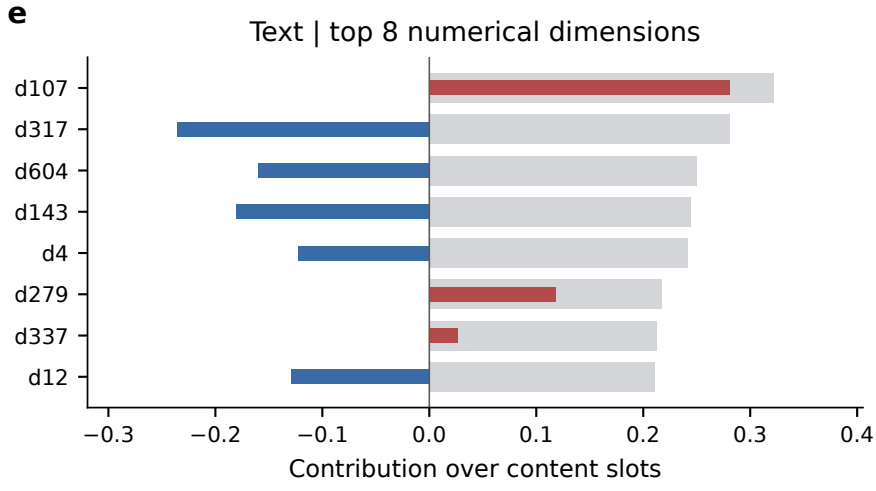
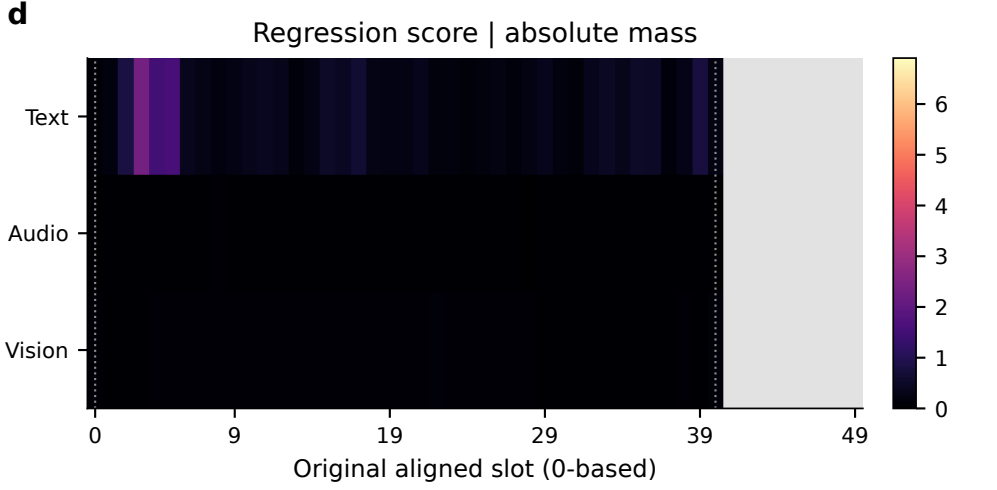
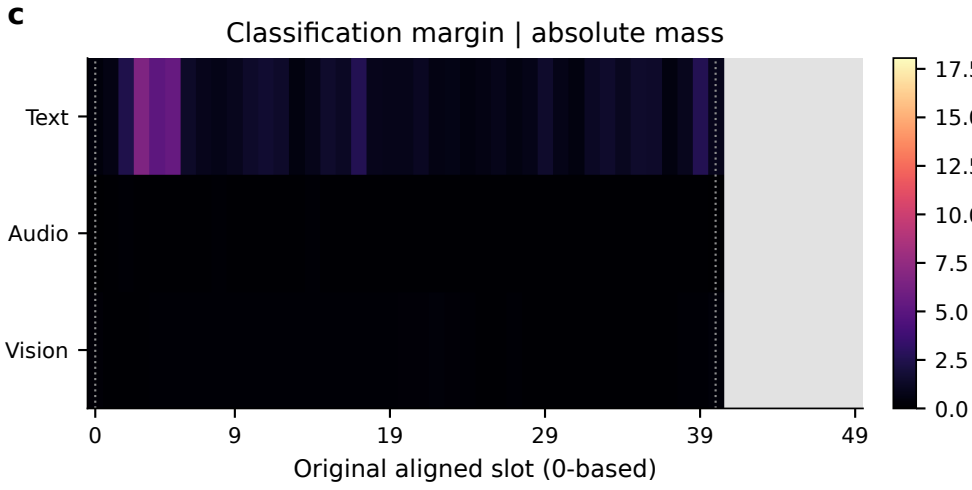
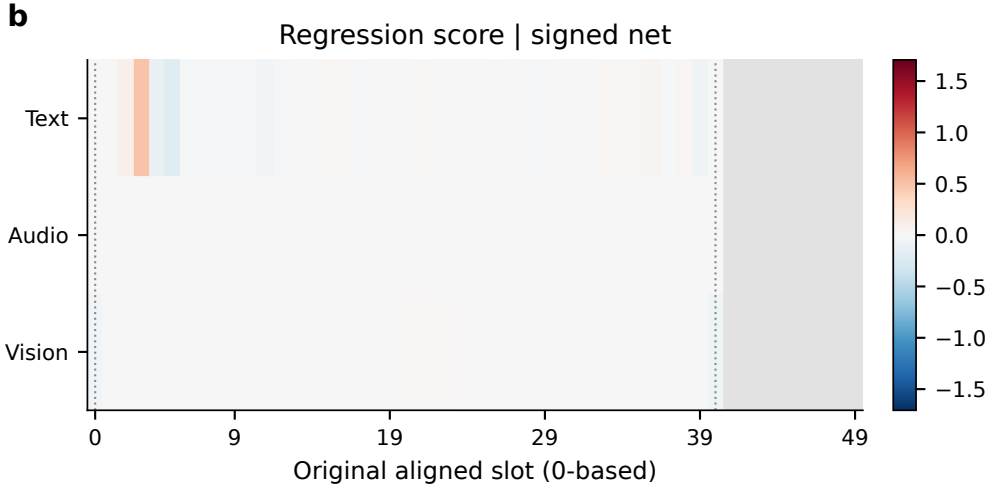
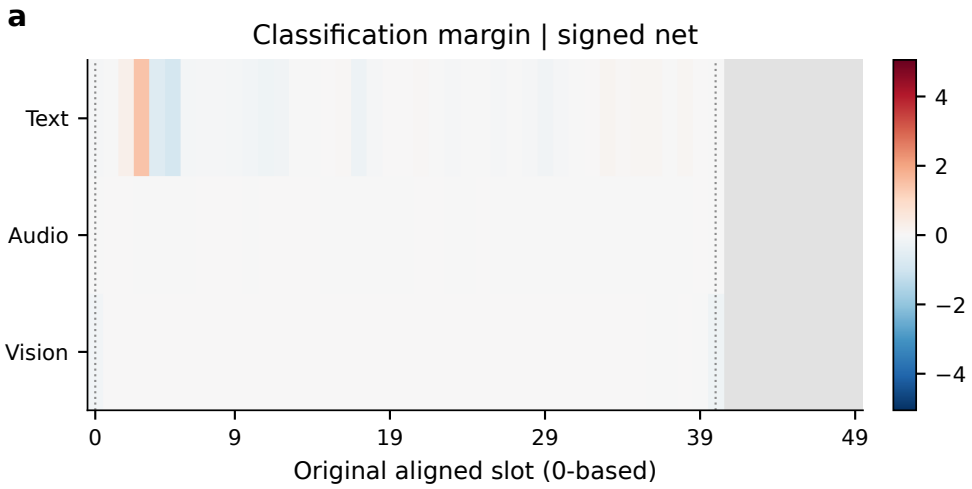
Predicted class 1 | Class probabilities [0: 0.316, 1: 0.458, 2: 0.226] | Regression -0.0311

Text truncated: 48 of 57 content tokens retained.



Attachment 4 | Sample 19 | Time-step and feature evidence

Predicted class 2 | Class probabilities [0: 0.166, 1: 0.055, 2: 0.778] | Regression 0.6792



— Absolute contribution mass — Positive signed net — Negative signed net

Train-mean reference. Heatmap scales are shared across all 20 samples for each target. Gray cells: padding; dotted lines: CLS/SEP. Absolute mass is not a modality share. Dimension rankings exclude specials and padding. Dimensions are numerical indices, not semantic labels or video timestamps.

Attachment 4 | Sample 20 | Time-step and feature evidence

Predicted class 2 | Class probabilities [0: 0.012, 1: 0.051, 2: 0.937] | Regression 1.2047

