









Attachment 4 | Sample 01 | Intervention and stability checks


Predicted class 1 | Class probabilities [0: 0.073, 1: 0.490, 2: 0.436] | Regression 0.1541



Replacement keeps masks fixed. Slot curves use the train-mean reference; random shading is +/- 1 SD over 20 draws. Word bars compare attribution with original-minus-reencoded output, so they need not agree. These checks assess model behavior, not correctness or real-world causality. N/A denotes an undefined rank correlation.


Attachment 4 | Sample 02 | Intervention and stability checks

Predicted class 0 | Class probabilities [0: 0.571, 1: 0.190, 2: 0.239] | Regression -0.2530



Attachment 4 | Sample 03 | Intervention and stability checks


Predicted class 1 | Class probabilities [0: 0.251, 1: 0.432, 2: 0.316] | Regression 0.0704



Replacement keeps masks fixed. Slot curves use the train-mean reference; random shading is +/- 1 SD over 20 draws. Word bars compare attribution with original-minus-reencoded output, so they need not agree. These checks assess model behavior, not correctness or real-world causality. N/A denotes an undefined rank correlation.









Attachment 4 | Sample 04 | Intervention and stability checks

Predicted class 2 | Class probabilities [0: 0.309, 1: 0.088, 2: 0.603] | Regression 0.4894



Attachment 4 | Sample 05 | Intervention and stability checks





Predicted class 2 | Class probabilities [0: 0.031, 1: 0.208, 2: 0.761] | Regression 0.6403



Replacement keeps masks fixed. Slot curves use the train-mean reference; random shading is +/- 1 SD over 20 draws. Word bars compare attribution with original-minus-reencoded output, so they need not agree. These checks assess model behavior, not correctness or real-world causality. N/A denotes an undefined rank correlation.

Predicted class 2 | Class probabilities [0: 0.027, 1: 0.114, 2: 0.859] | Regression 0.8748

Predicted class 2 | Class probabilities [0: 0.027, 1: 0.114, 2: 0.859] | Regression 0.8748











Replacement keeps masks fixed. Slot curves use the train-mean reference; random shading is ± 1 SD over 20 draws. Word bars compare attribution with original-minus-reencoded output, so they need not agree. These checks assess model behavior, not correctness or real-world causality. N/A denotes an undefined rank correlation.

Attachment 4 | Sample 07 | Intervention and stability checks

Predicted class 2 | Class probabilities [0: 0.014, 1: 0.053, 2: 0.933] | Regression 1.2134


Text truncated: 48 of 63 content tokens retained.



Replacement keeps masks fixed. Slot curves use the train-mean reference; random shading is +/- 1 SD over 20 draws. Word bars compare attribution with original-minus-reencoded output, so they need not agree. These checks assess model behavior, not correctness or real-world causality. N/A denotes an undefined rank correlation.


Attachment 4 | Sample 08 | Intervention and stability checks

Predicted class 2 | Class probabilities [0: 0.052, 1: 0.247, 2: 0.701] | Regression 0.3236




Attachment 4 | Sample 09 | Intervention and stability checks

Predicted class 0 | Class probabilities [0: 0.922, 1: 0.037, 2: 0.041] | Regression -2.0012



Attachment 4 | Sample 10 | Intervention and stability checks


Predicted class 0 | Class probabilities [0: 0.892, 1: 0.033, 2: 0.076] | Regression -1.4864



Replacement keeps masks fixed. Slot curves use the train-mean reference; random shading is +/- 1 SD over 20 draws. Word bars compare attribution with original-minus-reencoded output, so they need not agree. These checks assess model behavior, not correctness or real-world causality. N/A denotes an undefined rank correlation.

Attachment 4 | Sample 11 | Intervention and stability checks








Predicted class 0 | Class probabilities [0: 0.912, 1: 0.049, 2: 0.039] | Regression -1.9479



Replacement keeps masks fixed. Slot curves use the train-mean reference; random shading is +/- 1 SD over 20 draws. Word bars compare attribution with original-minus-reencoded output, so they need not agree. These checks assess model behavior, not correctness or real-world causality. N/A denotes an undefined rank correlation.

Attachment 4 | Sample 12 | Intervention and stability checks

Predicted class 0 | Class probabilities [0: 0.790, 1: 0.128, 2: 0.082] | Regression -0.6350











Replacement keeps masks fixed. Slot curves use the train-mean reference; random shading is ± 1 SD over 20 draws. Word bars compare attribution with original-minus-reencoded output, so they need not agree. These checks assess model behavior, not correctness or real-world causality. N/A denotes an undefined rank correlation.

Attachment 4 | Sample 13 | Intervention and stability checks

Predicted class 1 | Class probabilities [0: 0.087, 1: 0.613, 2: 0.300] | Regression 0.0550









Raw vision is all zero; visual attribution explains zero-input behavior.



Replacement keeps masks fixed. Slot curves use the train-mean reference; random shading is +/- 1 SD over 20 draws. Word bars compare attribution with original-minus-reencoded output, so they need not agree. These checks assess model behavior, not correctness or real-world causality. N/A denotes an undefined rank correlation.

Attachment 4 | Sample 14 | Intervention and stability checks









Predicted class 2 | Class probabilities [0: 0.052, 1: 0.368, 2: 0.579] | Regression 0.3282



Replacement keeps masks fixed. Slot curves use the train-mean reference; random shading is +/- 1 SD over 20 draws. Word bars compare attribution with original-minus-reencoded output, so they need not agree. These checks assess model behavior, not correctness or real-world causality. N/A denotes an undefined rank correlation.

Attachment 4 | Sample 15 | Intervention and stability checks









Predicted class 2 | Class probabilities [0: 0.022, 1: 0.089, 2: 0.889] | Regression 1.0632



Replacement keeps masks fixed. Slot curves use the train-mean reference; random shading is ± 1 SD over 20 draws. Word bars compare attribution with original-minus-reencoded output, so they need not agree. These checks assess model behavior, not correctness or real-world causality. N/A denotes an undefined rank correlation.

Predicted class 0 | Class probabilities [0: 0.949, 1: 0.023, 2: 0.028] | Regression -2.0530








Predicted class 0 | Class probabilities [0: 0.949, 1: 0.023, 2: 0.028] | Regression -2.0530



Replacement keeps masks fixed. Slot curves use the train-mean reference; random shading is ± 1 SD over 20 draws. Word bars compare attribution with original-minus-reencoded output, so they need not agree. These checks assess model behavior, not correctness or real-world causality. N/A denotes an undefined rank correlation.

Predicted class 2 | Class probabilities [0: 0.015, 1: 0.041, 2: 0.944] | Regression 1.3862

Predicted class 2 | Class probabilities [0: 0.015, 1: 0.041, 2: 0.944] | Regression 1.3862











Replacement keeps masks fixed. Slot curves use the train-mean reference; random shading is ± 1 SD over 20 draws. Word bars compare attribution with original-minus-reencoded output, so they need not agree. These checks assess model behavior, not correctness or real-world causality. N/A denotes an undefined rank correlation.

Attachment 4 | Sample 18 | Intervention and stability checks

Predicted class 1 | Class probabilities [0: 0.316, 1: 0.458, 2: 0.226] | Regression -0.0311









Text truncated: 48 of 57 content tokens retained.



Replacement keeps masks fixed. Slot curves use the train-mean reference; random shading is +/- 1 SD over 20 draws. Word bars compare attribution with original-minus-reencoded output, so they need not agree. These checks assess model behavior, not correctness or real-world causality. N/A denotes an undefined rank correlation.

Attachment 4 | Sample 19 | Intervention and stability checks





Predicted class 2 | Class probabilities [0: 0.166, 1: 0.055, 2: 0.778] | Regression 0.6792



Replacement keeps masks fixed. Slot curves use the train-mean reference; random shading is +/- 1 SD over 20 draws. Word bars compare attribution with original-minus-reencoded output, so they need not agree. These checks assess model behavior, not correctness or real-world causality. N/A denotes an undefined rank correlation.

Predicted class 2 | Class probabilities [0: 0.012, 1: 0.051, 2: 0.937] | Regression 1.2047

Predicted class 2 | Class probabilities [0: 0.012, 1: 0.051, 2: 0.937] | Regression 1.2047



Replacement keeps masks fixed. Slot curves use the train-mean reference; random shading is ± 1 SD over 20 draws. Word bars compare attribution with original-minus-reencoded output, so they need not agree. These checks assess model behavior, not correctness or real-world causality. N/A denotes an undefined rank correlation.